



Thank You for your purchase
Microsoft AI-102 Exam Question & Answers
Designing and Implementing a Microsoft Azure AI
Solution Exam

Product Questions: 337

Version: 26.2

Topic 1, Wide World Importers

Case study

Overview

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview

Existing Environment

A company named Wide World Importers is developing an e-commerce platform.

You are working with a solutions architect to design and implement the features of the e-commerce platform. The platform will use microservices and a serverless environment built on Azure.

Wide World Importers has a customer base that includes English, Spanish, and Portuguese speakers.

Applications

Wide World Importers has an App Service plan that contains the web apps shown in the following table.

Name	Description
Product Management	An app used by employees to create and manage products. The app and the expected inputs from the employees are in English.
Inventory Tracking	An app used by employees to manage inventory when dispatching orders, receiving refunds, and receiving consignments from suppliers.

Azure Resources

You have the following resources:

An Azure Active Directory (Azure AD) tenant

The tenant supports internal authentication.

All employees belong to a group named AllUsers.

Senior managers belong to a group named LeadershipTeam.

An Azure Functions resource

A function app posts to Azure Event Grid when stock levels of a product change between OK, Low Stock, and Out of Stock. The function app uses the Azure Cosmos DB change feed.

An Azure Cosmos DB account

The account uses the Core (SQL) API.

The account stores data for the Product Management app and the Inventory Tracking app.

An Azure Storage account

The account contains blob containers for assets related to products.

The assets include images, videos, and PDFs.

An Azure Cognitive Services resource named wwics

A Video Indexer resource named wwivi

Requirements

Business Goals

Wide World Importers wants to leverage AI technologies to differentiate itself from its competitors.

Planned Changes

Wide World Importers plans to start the following projects:

A product creation project: Help employees create accessible and multilingual product entries, while expediting product entry creation.

A smart e-commerce project: Implement an Azure Cognitive Search solution to display products for customers to browse.

A shopping on-the-go project: Build a chatbot that can be integrated into smart speakers to support customers.

Business Requirements

Wide World Importers identifies the following business requirements for all the projects:

Provide a multilingual customer experience that supports English, Spanish, and Portuguese.
Whenever possible, scale based on transaction volumes to ensure consistent performance.
Minimize costs.

Governance and Security Requirements

Wide World Importers identifies the following governance and security requirements:

Data storage and processing must occur in datacenters located in the United States.
Azure Cognitive Services must be inaccessible directly from the internet.

Accessibility Requirements

Wide World Importers identifies the following accessibility requirements:

All images must have relevant alt text.
All videos must have transcripts that are associated to the video and included in product descriptions.
Product descriptions, transcripts, and all text must be available in English, Spanish, and Portuguese.

Product Creation Requirements

Wide World Importers identifies the following requirements for improving the Product Management app:

Minimize how long it takes for employees to create products and add assets.
Remove the need for manual translations.

Smart E-Commerce Requirements

Wide World Importers identifies the following requirements for the smart e-commerce project:

Ensure that the Cognitive Search solution meets a Service Level Agreement (SLA) of 99.9% availability for searches and index writes.
Provide users with the ability to search insight gained from the images, manuals, and videos associated with the products.
Support autocomplete and autosuggestion based on all product name variants.
Store all raw insight data that was generated, so the data can be processed later.
Update the stock level field in the product index immediately upon changes.
Update the product index hourly.

Shopping On-the-Go Requirements

Wide World Importers identifies the following requirements for the shopping on-the-go chatbot:

Answer common questions.

Support interactions in English, Spanish, and Portuguese.

Replace an existing FAQ process so that all Q&A is managed from a central location.

Provide all employees with the ability to edit Q&As. Only senior managers must be able to publish updates.

Support purchases by providing information about relevant products to customers. Product displays must include images and warnings when stock levels are low or out of stock.

Product JSON Sample

You have the following JSON sample for a product.

```
{  
    "sku": "b1",  
    "name": {  
        "en": "Bicycle",  
        "es": "Bicicleta",  
        "pt": "Bicicleta"  
    },  
    "stocklevel": "Out of Stock",  
    "description": {  
        "en": "Bicycle",  
        "es": "Bicicleta",  
        "pt": "Bicicleta"  
    },  
    "image":  
    {"uri": "https://upload.worldwideimporters.org/bicycle.jpg",  
        "alttext": {  
            "en": "Bicycle",  
            "es": "Bicicleta",  
            "pt": "Bicicleta"  
        }  
    },  
    "createdUtc": "2020-02-14T06:08:39Z",  
    "language": "en"  
}
```

Question: 1

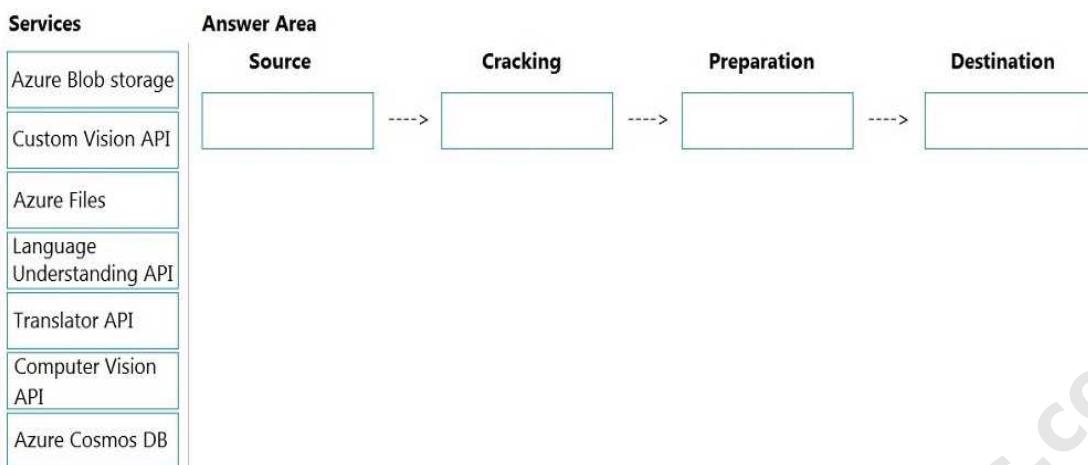
DRAG DROP

You are developing the smart e-commerce project.

You need to design the skillset to include the contents of PDFs in searches.

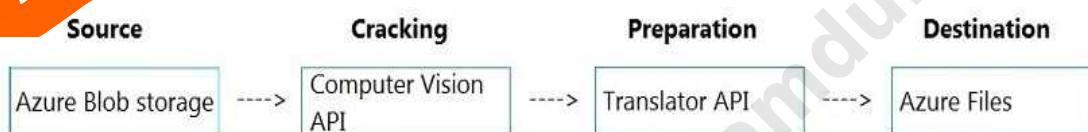
How should you complete the skillset design diagram? To answer, drag the appropriate services to the correct stages. Each service may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.



Answer:

Explanation:



Box 1: Azure Blob storage

At the start of the pipeline, you have unstructured text or non-text content (such as images, scanned documents, or JPEG files). Data must exist in an Azure data storage service that can be accessed by an indexer.

Box 2: Computer Vision API

Scenario: Provide users with the ability to search insight gained from the images, manuals, and videos associated with the products.

The Computer Vision Read API is Azure's latest OCR technology (learn what's new) that extracts printed text (in several languages), handwritten text (English only), digits, and currency symbols from images and multi-page PDF documents.

Box 3: Translator API

Scenario: Product descriptions, transcripts, and all text must be available in English, Spanish, and Portuguese.

Box 4: Azure Files

Scenario: Store all raw insight data that was generated, so the data can be processed later.

Incorrect Answers:

The custom vision API from Microsoft Azure learns to recognize specific content in imagery and becomes smarter with training and time.

Reference:

<https://docs.microsoft.com/en-us/azure/search/cognitive-search-concept-intro>

<https://docs.microsoft.com/en-us/azure/cognitive-services/computer-vision/overview-ocr>

Question: 2

DRAG DROP

You are planning the product creation project.

You need to recommend a process for analyzing videos.

Which four actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order. (Choose four.)

Actions	Answer Area
Index the video by using the Video Indexer API.	
Upload the video to blob storage.	
Analyze the video by using the Computer Vision API.	
Extract the transcript from Microsoft Stream.	
Send the transcript to the Language Understanding API as an utterance.	
Extract the transcript from the Video Indexer API.	
Translate the transcript by using the Translator API.	
Upload the video to file storage.	

Answer:

Explanation:

Actions	Answer Area
Index the video by using the Video Indexer API.	Upload the video to blob storage.
Upload the video to blob storage.	Index the video by using the Video Indexer API.
Analyze the video by using the Computer Vision API.	Extract the transcript from the Video Indexer API.
Extract the transcript from Microsoft Stream.	Translate the transcript by using the Translator API.
Send the transcript to the Language Understanding API as an utterance.	
Extract the transcript from the Video Indexer API.	
Translate the transcript by using the Translator API.	
Upload the video to file storage.	



Scenario: All videos must have transcripts that are associated to the video and included in product descriptions.

Product descriptions, transcripts, and all text must be available in English, Spanish, and Portuguese.

Step 1: Upload the video to blob storage

Given a video or audio file, the file is first dropped into a Blob Storage. T

Step 2: Index the video by using the Video Indexer API.

When a video is indexed, Video Indexer produces the JSON content that contains details of the specified video insights. The insights include: transcripts, OCRs, faces, topics, blocks, etc.

Step 3: Extract the transcript from the Video Indexer API.

Step 4: Translate the transcript by using the Translator API.

Reference:

<https://azure.microsoft.com/en-us/blog/get-video-insights-in-even-more-languages/>

<https://docs.microsoft.com/en-us/azure/media-services/video-indexer/video-indexer-output-json-v2>

Question: 3

HOTSPOT

You are planning the product creation project.

You need to build the REST endpoint to create the multilingual product descriptions.

How should you complete the URI? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

api.cognitive.microsofttranslator.com	?api-version=3.0&to=es&to=pt
api-nam.cognitive.microsofttranslator.com	/detect
westus.tts.speech.microsoft.com	/languages
wwics.cognitiveservices.azure.com/translator	/text-to-speech
	/translate

Answer:

Explanation: 1. **api.cognitive.microsofttranslator.com**(global recommended - it tooks nearest location)

Box 1: api-nam.cognitive.microsofttranslator.com

<https://docs.microsoft.com/en-us/azure/cognitive-services/translator/reference/v3-0-reference>

Box 2: /translate 2. **/translate**

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/translator/reference/v3-0-translate>

Question: 4

You are developing the smart e-commerce project.

You need to implement autocompletion as part of the Cognitive Search solution.

Which three actions should you perform? Each correct answer presents part of the solution. (Choose three.)

NOTE: Each correct selection is worth one point.

- A. Make API queries to the autocomplete endpoint and include suggesterName in the body.
- B. Add a suggester that has the three product name fields as source fields.
- C. Make API queries to the search endpoint and include the product name fields in the searchFields query parameter.
- D. Add a suggester for each of the three product name fields.
- E. Set the searchAnalyzer property for the three product name variants.
- F. Set the analyzer property for the three product name variants.

Answer: ABF

Explanation:

Scenario: Support autocomplete and autosuggestion based on all product name variants.

A: Call a suggester-enabled query, in the form of a Suggestion request or Autocomplete request, using an API. API usage is illustrated in the following call to the Autocomplete REST API.

```
POST /indexes/myboxgames/docs/autocomplete?search&api-version=2020-06-30
{
  "search": "minecraf",
  "suggesterName": "sg"
}
```

B: In Azure Cognitive Search, typeahead or "search-as-you-type" is enabled through a suggester. A suggester provides a list of fields that undergo additional tokenization, generating prefix sequences to support matches on partial terms. For example, a suggester that includes a City field with a value for "Seattle" will have prefix combinations of "sea", "seat", "seatt", and "seattl" to support typeahead.

F. Use the default standard Lucene analyzer ("analyzer": null) or a language analyzer (for example, "analyzer": "en.Microsoft") on the field.

Reference:

<https://docs.microsoft.com/en-us/azure/search/index-add-suggesters>

Question: 5

HOTSPOT

You are developing the shopping on-the-go project.

You are configuring access to the QnA Maker resources.

Which role should you assign to AllUsers and LeadershipTeam? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Cognitive Services QnA Maker Editor: (Leadership Team - if content editing is part of their role)
Can create, edit, and test changes in the knowledge base but cannot publish.

AllUsers: QnA Maker Reader (provides read-only access to the QnA knowledge base, allowing users to query and retrieve answers).

LeadershipTeam: Cognitive Services User (allows the leadership team to read, write, and publish changes to the QnA knowledge base)

Cognitive Services User role and the Cognitive Services QnA Maker Editor role, is the ability to publish changes to the QnA Maker knowledge base.

Answer Area

AllUsers:	<table border="1"><tr><td>Cognitive Service User</td></tr><tr><td>Contributor</td></tr><tr><td>Owner</td></tr><tr><td>QnA Maker Editor</td></tr><tr><td>QnA Maker Read</td></tr></table>	Cognitive Service User	Contributor	Owner	QnA Maker Editor	QnA Maker Read
Cognitive Service User						
Contributor						
Owner						
QnA Maker Editor						
QnA Maker Read						
LeadershipTeam:	<table border="1"><tr><td>Cognitive Service User</td></tr><tr><td>Contributor</td></tr><tr><td>Owner</td></tr><tr><td>QnA Maker Editor</td></tr><tr><td>QnA Maker Read</td></tr></table>	Cognitive Service User	Contributor	Owner	QnA Maker Editor	QnA Maker Read
Cognitive Service User						
Contributor						
Owner						
QnA Maker Editor						
QnA Maker Read						

Answer:

Explanation:

<https://learn.microsoft.com/en-us/azure/search/search-security-rbac>

Answer Area

AllUsers:

Cognitive Service User
Contributor
Owner
QnA Maker Editor
QnA Maker Read

Management-Accountant:CognitiveServiceUser

ConsultantAccountant:
CognitiveServicesQnAMakerEdit

Agent-CustomerServices:
CognitiveServicesQnAMakerRead

LeadershipTeam:



All User : QnA Maker Reader

LeadershipTeam: Cognitive Services User.

Cognitive Service User
Contributor
Owner
QnA Maker Editor
QnA Maker Read

Box 1: QnA Maker Editor

Scenario: Provide all employees with the ability to edit Q&As.

The QnA Maker Editor (read/write) has the following permissions:

Create KB API
Update KB API
Replace KB API
Replace Alterations
"Train API" [in
new service model v5]

Box 2: Contributor

Scenario: Only senior managers must be able to publish updates.

Contributor permission: All except ability to add new members to roles

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/reference-role-based-access-control>

Question: 6

HOTSPOT

You need to develop code to upload images for the product creation project. The solution must meet the accessibility requirements.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

```
public static async Task<string> SuggestAltText(ComputerVisionClient client,
{
    List<VisualFeatureTypes?> features = new List<VisualFeatureTypes?>()
    {
        VisualFeatureTypes.Description
        VisualFeatureTypes.ImageType
        VisualFeatureTypes.Objects
        VisualFeatureTypes.Tags
    };
    ImageAnalysis results = await client.AnalyzeImageAsync(image, features);
}

var c = results.Brands.DetectedBrands[0]
var c = results.Description.Captions[0]
var c = results.Metadata[0]
var c = results.Objects[0]

if(c.Confidence>0.5) return(c.Text);
}
```

Dictionary
stream
string

image)

Explanation:

```
public static async Task<string> SuggestAltText(ComputerVisionClient client,
{
    List<VisualFeatureTypes?> features = new List<VisualFeatureTypes?>()
    {
        VisualFeatureTypes.Description
        VisualFeatureTypes.ImageType
        VisualFeatureTypes.Objects
        VisualFeatureTypes.Tags
    };
    ImageAnalysis results = await client.AnalyzeImageAsync(image, features);

    var c = results.Brands.DetectedBrands[0]
    var c = results.Description.Captions[0]
    var c = results.Metadata[0]
    var c = results.Objects[0]

    if(c.Confidence>0.5) return(c.Text);
}
```

Answer:

Dictionary
stream
string

image)

Reference:

<https://github.com/Azure-Samples/cognitive-services-dotnet-sdk-samples/blob/master/documentation-samples/quickstarts/ComputerVision/Program.cs>

Question: 7

HOTSPOT

You are developing the shopping on-the-go project.

You need to build the Adaptive Card for the chatbot.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
version": "1.3",
"body": [
  {
    "type": "TextBlock",
    "size": "Medium",
    "weight": "Bolder",
    "text": "${if(language == 'en', 'en', name)}  
name  
name.en  
name[language]"
  },
  {
    "type": "TextBlock",
    "$when": "${stockLevel != 'OK'}"  
"$when": "${stockLevel == 'OK'}"  
"$when": "${stockLevel.OK}"
    color : Attention
  },
  {
    "type": "Image",
    "url": "${image.uri}",
    "size": "Medium",
    "altText": "${image.altText.en}  
image.altText.language  
image.altText["language"]  
image.altText[language]"
  }
]
```

Answer:

Explanation:

```
version": "1.3",
"body": [
  {
    "type": "TextBlock",
    "size": "Medium",
    "weight": "Bolder",
    "text": "${if(language == 'en', 'en', name)}  
name  
name.en  
name[language]"
  },
  {
    "type": "TextBlock",
    "text": "$when": "${stockLevel != 'OK'}"  
"$when": "${stockLevel == 'OK'}"  
"$when": "${stockLevel.OK}"
    color : Attention
  },
  {
    "type": "Image",
    "url": "${image.uri}",
    "size": "Medium",
    "altText": "${image.altText.en}  
image.altText.language  
image.altText["language"]  
image.altText[language]"
  }
]
```

Box 1: name.en

Box 2: "\$when": "\${stockLevel != 'OK'}"

Product displays must include images and warnings when stock levels are low or out of stock.

Box 3:image.altText.en

Topic 2, Contoso, Ltd.

Case Study

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other questions in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

General Overview

Contoso, Ltd. is an international accounting company that has offices in France, Portugal, and the United Kingdom. Contoso has a professional services department that contains the

Name	Position	Office
Accountant	Manager	United Kingdom, France, Portugal
Accountant	Consultant	United Kingdom, France, Portugal
Customer Service	Manager	United Kingdom
Customer Service	Agent	United Kingdom
Bookkeeper	Manager	United Kingdom, France, Portugal
Bookkeeper	Consultant	United Kingdom, France, Portugal

Infrastructure

Contoso has the following subscriptions:

- Azure
- Microsoft 365
- Microsoft Dynamics 365

Azure Active (Azure AD) Directory

Contoso has Azure Active Directory groups for securing role-based access. The company uses the following group naming conventions:

- ICountryJ-[Level]-[Role]
- [Level]-[Role]

Intellectual Property

Contoso has the intellectual property shown in the following table.

Content	Format	Language	Content store	Domain
Weekly webinars	Video	English	Azure Blob storage	Vid.contoso.com
Blogs	Text	English, French, Portuguese	WordPress	Pt-blog.contoso.com Blog.contoso.com Fr-blog.contoso.com
Wikis	Text	English, French, Portuguese	Azure Cosmos DB	Internal.contoso.com/wiki
Monthly conference recordings	Video	English	SharePoint Online	Contoso.sharepoint.com
Frequently asked questions (FAQs)	Text	English	SharePoint Online	Contoso.sharepoint.com

Text-based content is provided only in one language and is not translated.

Planned Projects

Contoso plans to develop the following:

- A document processing workflow to extract information automatically from PDFs and images of financial documents
- A customer-support chatbot that will answer questions by using FAQs
- A searchable knowledgebase of all the intellectual property

Technical Requirements

Contoso identifies the following technical requirements:

- All content must be approved before being published.
- All planned projects must support English, French, and Portuguese.
- All content must be secured by using role-based access control (RBAC).
- RBAC role assignments must use the principle of least privilege.
- RBAC roles must be assigned only to Azure Active Directory groups.
- AI solution responses must have a confidence score that is equal to or greater than 70 percent.
- When the response confidence score of an AI response is lower than 70 percent, the response must be improved by human input.

Chatbot Requirements

Contoso identifies the following requirements for the chatbot:

- Provide customers with answers to the FAQs.
- Ensure that the customers can chat to a customer service agent.
- Ensure that the members of a group named Management-Accountants can approve the FAQs.
- Ensure that the members of a group named Consultant-Accountants can create and amend the FAQs.
- Ensure that the members of a group named the Agent-CustomerServices can browse the FAQs.
- Ensure that access to the customer service agents is managed by using Omnichannel for Customer Service.
- When the response confidence score is low, ensure that the chatbot can provide other response options to the customers.

Document Processing Requirements

Contoso identifies the following requirements for document processing:

- The document processing solution must be able to process standardized financial documents that have the following characteristics:
 - Contain fewer than 20 pages.
 - Be formatted as PDF or JPEG files.
 - Have a distinct standard for each office.

- The document processing solution must be able to extract tables and text from the financial documents.
- The document processing solution must be able to extract information from receipt images.
- Members of a group named Management-Bookkeeper must define how to extract tables from the financial documents.
- Members of a group named Consultant-Bookkeeper must be able to process the financial documents.

Knowledgebase Requirements

Contoso identifies the following requirements for the knowledgebase:

- Supports searches for equivalent terms
- Can transcribe jargon with high accuracy
- Can search content in different formats, including video
- Provides relevant links to external resources for further research

Question: 8

You need to develop an extract solution for the receipt images. The solution must meet the document processing requirements and the technical requirements.

You upload the receipt images to the From Recognizer API for analysis, and the API returns the following JSON.

```
"documentResults": [
  {
    "docType": "prebuilt:receipt",
    "pageRange": [
      1,
      1
    ],
    "fields": {
      "ReceiptType": {
        "type": "string",
        "valueString": "Itemized",
        "confidence": 0.672
      },
      "MerchantName": {
        "type": "string",
        "valueString": "Tailwind",
        "text": "Tailwind",
        "boundingBox": [],
        "page": 1,
        "confidence": 0.913,
        "elements": [
          "#/readResults/0/lines/0/words/0"
        ]
      },
      ...
    }
  }
]
```

Which expression should you use to trigger a manual review of the extracted information by a member of the Consultant-Bookkeeper group?

- A. documentResults.docType == "prebuilt:receipt" This option checks the document type, not the confidence level of a specific field.
- B. documentResults.fields.".confidence < 0.7 JSON structure does not support using an asterisk (*) to represent all fields.
- C. documentResults.fields.ReceiptType.confidence > 0.7
- D. documentResults.fields.MerchantName.confidence < 0.7 is below 0.7, which indicates low confidence.

Answer: D

Explanation:

Need to specify the field name, and then use < 0.7 to handle trigger if confidence score is less than 70%.

Reference:

<https://docs.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/api-v2-0/reference-sdk-api-v2-0>

Question: 9

HOTSPOT

You are developing the knowledgebase by using Azure Cognitive Search.

You need to build a skill that will be used by indexers.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
{  
    "@odata.type": "#Microsoft.Skills.Text.EntityRecognitionSkill",  
    "categories": [],  
    "categories": [ "Email", "Persons", "Organizations"],  
    "categories": [ "Locations", "Persons", "Organizations"],  
  
    "minimumPrecision": 0.7,  
    "inputs": [  
        { "name": "text",  
          "source": "/document/content"}],  
    "outputs": [  
        { "name": "persons", "targetName": "people"},  
        { "name": "locations", "targetName": "locations"},  
        { "name": "organizations", "targetName": "organizations"},  
  
        1 { "name": "entities" }  
        { "name": "categories" }  
        { "name": "namedEntities" }]
```

category:Person,Location,Organization

value: "paris","john Doe"

Location,Person,Organizations

entities fields

category,value,offset,
confidence

namedEntities fields

category,subcategory
text,offset,length,
confidentscore

Answer:

Explanation:

Box 1: "categories": ["Locations", "Persons", "Organizations"],

Locations, Persons, Organizations are in the outputs.

Scenario: Contoso plans to develop a searchable knowledgebase of all the intellectual property

Note: The categories parameter is an array of categories that should be extracted. Possible category

types: "Person", "Location", "Organization", "Quantity", "Datetime", "URL", "Email". If no category is provided, all types are returned.

entities and namedEntities are often used interchangeably,

namedEntities is a specific name that can be given to the skill output.

Box 2: {"name": "entities"}

The include wikis, so should include entities in the outputs.

Note: entities is an array of complex types that contains rich information about the entities extracted from text, with the following fields

name (the actual entity name. This represents a "normalized" form)
wikipediaId
wikipediaLanguage
wikipediaUrl (a link to Wikipedia page for the entity)
etc.

Reference:

<https://docs.microsoft.com/en-us/azure/search/cognitive-search-skill-entity-recognition>

Question: 10

You are developing the chatbot.

You create the following components:

- A QnA Maker resource
- A chatbot by using the Azure Bot Framework SDK

You need to add an additional component to meet the technical requirements and the chatbot requirements. What should you add?

dispatch-You need a way to determine if the user's input should go to QnA Maker

- A. Dispatch for an answer, or if it should be handled by a different service or logic
(like a natural language understanding model for more complex intent).
B. chatdown
C. Language Understanding
D. Microsoft Translator

Chatdown is a tool used to create a conversation transcript for your chatbot but does not provide routing functionality

Answer: A

Explanation:

Scenario: All planned projects must support English, French, and Portuguese.

If a bot uses multiple LUIS models and QnA Maker knowledge bases (knowledge bases), you can use the Dispatch tool to determine which LUIS model or QnA Maker knowledge base best matches the user input. The dispatch tool does this by creating a single LUIS app to route user input to the correct model.

Reference:

<https://docs.microsoft.com/en-us/azure/bot-service/bot-builder-tutorial-dispatch>

Question: 11

You are developing the document processing workflow.

You need to identify which API endpoints to use to extract text from the financial documents. The solution must meet the document processing requirements.

Which two API endpoints should you identify? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. /vision/v3.2/read/analyzeResults
B. /formrecognizer/v2.0/prebuilt/receipt/analyze
C. /vision/v3.2/read/analyze

analyzing receipts using a prebuilt model.
While it's designed for receipts, it's important to consider using prebuilt models for financial documents.

- D. /vision/v3.2/describe
E. /formrecognizer/v2.0/custom/models{modelId}/ analyze

Answer: BC

Explanation:

C: Analyze Receipt - Get Analyze Receipt Result.

Query the status and retrieve the result of an Analyze Receipt operation.

Request URL: <https://{{endpoint}}/formrecognizer/v2.0-preview/prebuilt/receipt/analyzeResults/{{resultId}}>

E: POST {Endpoint}/vision/v3.2/read/analyze

Use this interface to get the result of a Read operation, employing the state-of-the-art Optical Character Recognition (OCR) algorithms optimized for text-heavy documents.

Scenario: Contoso plans to develop a document processing workflow to extract information automatically from PDFs and images of financial documents

The document processing solution must be able to process standardized financial documents that have the following characteristics:

- Contain fewer than 20 pages.
- Be formatted as PDF or JPEG files.
- Have a distinct standard for each office.

*The document processing solution must be able to extract tables and text from the financial documents.

The document processing solution must be able to extract information from receipt images.

Reference:

<https://westus2.dev.cognitive.microsoft.com/docs/services/form-recognizer-api-v2-preview/operations/GetAnalyzeReceiptResult>

<https://docs.microsoft.com/en-us/rest/api/computervision/3.1/read/read>

Question: 12

You are developing the chatbot.

You create the following components:

- * A QnA Maker resource
- * A chatbot by using the Azure Bot Framework SDK.

You need to integrate the components to meet the chatbot requirements.

Which property should you use?

- A. QnADialogResponseOptions.CardNoMatchText
- B. Qna MakerOptions-ScoreThreshold (**chatbot**)
- C. Qna Maker Options StrickFilters
- D. QnaMakerOptions.RankerType (**user**)

QnaMakerOptions.ScoreThreshold

Essential for ensuring that your chatbot only provides answers when QnA Maker is confident in its response. This improves the quality of answers

since we need to measure chatbot responsescorethreshold is correct

QnaMakerOptions.RankerType

Use QuestionOnly if you want to improve accuracy and make sure the user's query matches directly with a question in the KB.

This property determines how QnA Maker searches for a matching answer within the knowledge base.

This property is used to specify the type of ranker used to sort QnA Maker results but doesn't directly control the integration threshold.

Answer: D

Explanation:

Scenario: When the response confidence score is low, ensure that the chatbot can provide other response options to the customers.

When no good match is found by the ranker, the confidence score of 0.0 or "None" is returned and the default response is "No good match found in the KB". You can override this default response in the bot or application code calling the endpoint. Alternately, you can also set the override response in Azure and this changes the default for all knowledge bases deployed in a particular QnA Maker service.

Choosing Ranker type: By default, QnA Maker searches through questions and answers. If you want to search through questions only, to generate an answer, use the RankerType=QuestionOnly in the POST body of the GenerateAnswer request.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/concepts/best-practices>

Question: 13

HOTSPOT

You build a QnA Maker resource to meet the chatbot requirements.

Which RBAC role should you assign to each group? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Management-Accountants

▼
Owner
Contributor
Cognitive Services User
Cognitive Services QnA Maker Read
Cognitive Services QnA Maker Editor

Management Accountants:
Cognitive Services User

B.

Consultant-Accountants

▼
Owner
Contributor
Cognitive Services User
Cognitive Services QnA Maker Read
Cognitive Services QnA Maker Editor

Consultant Accountants:
Cognitive Services QnA Maker Edit

Agent-CustomerServices

▼
Owner
Contributor
Cognitive Services User
Cognitive Services QnA Maker Read
Cognitive Services QnA Maker Editor

Agent-Customer Services:
Cognitive Services User**Answer:**

Explanation:

Box 1: Cognitive Service User

Ensure that the members of a group named Management-Accountants can approve the FAQs.

Approve=publish.

Cognitive Service User (read/write/publish): API permissions: All access to Cognitive Services resource except for ability to:

1. Add new members to roles.
2. Create new resources.

Box 2: Cognitive Services QnA Maker Editor

Ensure that the members of a group named Consultant-Accountants can create and amend the FAQs.

QnA Maker Editor: API permissions:

1. Create KB API
2. Update KB API
3. Replace KB API
4. Replace Alterations
5. "Train API" [in new service model v5]

Box 3: Cognitive Services QnA Maker Read

Ensure that the members of a group named the Agent-CustomerServices can browse the FAQs.

QnA Maker Read: API Permissions:

1. Download KB API
 2. List KBs for user API
 3. Get Knowledge base details
 4. Download Alterations
- Generate Answer

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/concepts/role-based-access-control>

Question: 14

DRAG DROP

You are developing a solution for the Management-Bookkeepers group to meet the document processing requirements. The solution must contain the following components:

A From Recognizer resource

An Azure web app that hosts the Form Recognizer sample labeling tool

The Management-Bookkeepers group needs to create a custom table extractor by using the sample labeling tool.

Which three actions should the Management-Bookkeepers group perform in sequence? To answer, move the appropriate cmdlets from the list of cmdlets to the answer area and arrange them in the correct order.

Actions
Train a custom model
Label the sample documents
Create a new project and load sample documents
Create a composite model

Answer Area

- 1.Create a new project and load sample documents
- 2.Label the sample documents
- 3.Train a custom model



Answer:

Explanation:

Step 1: Create a new project and load sample documents

Create a new project. Projects store your configurations and settings.

Step 2: Label the sample documents

When you create or open a project, the main tag editor window opens.

Step 3: Train a custom model.

Finally, train a custom model.

composite model-
to analyze documents that can belong to one of several types or
have variations within a single type

Reference:

<https://docs.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/label-tool>

Question: 15

You are developing the knowledgebase.

You use Azure Video Analyzer for Media (previously Video indexer) to obtain transcripts of webinars.

You need to ensure that the solution meets the knowledgebase requirements.

What should you do?

B.This limits the insights you can extract from webinars.
you lose valuable visual insights like scene detection,
people recognition, etc., which can be helpful in a knowledgebase.

- A. Create a custom language model
 - B. Configure audio indexing for videos only
 - C. Enable multi-language detection for videos
 - D. Build a custom Person model for webinar presenters
- While useful for identifying presenters,
- C.Useful if webinars are multi-lingual, but doesn't improve accuracy for specific terminology within a single language.

Answer: B

Explanation:

Can search content in different formats, including video

Audio and video insights (multi-channels). When indexing by one channel, partial result for those models will be available.

Keywords extraction: Extracts keywords from speech and visual text.

Named entities extraction: Extracts brands, locations, and people from speech and visual text via natural language processing (NLP).

Topic inference: Makes inference of main topics from transcripts. The 2nd-level IPTC taxonomy is included.

Artifacts: Extracts rich set of "next level of details" artifacts for each of the models.

Sentiment analysis: Identifies positive, negative, and neutral sentiments from speech and visual text.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-video-analyzer/video-analyzer-for-media-docs/video-indexer-overview>

Question: 16

You are developing the knowledgebase by using Azure Cognitive Search.

You need to process wiki content to meet the technical requirements.

What should you include in the solution? This is the correct choice for ingesting content from Azure Blob Storage, where wiki content, often stored as files (e.g., HTML, Markdown), would likely be located.

- A. an indexer for Azure Blob storage attached to a skillset that contains the language detection skill

and the text translation skill

- B. an indexer for Azure Blob storage attached to a skillset that contains the language detection skill
- C. an indexer for Azure Cosmos DB attached to a skillset that contains the document extraction skill and the text translation skill
- D. an indexer for Azure Cosmos DB attached to a skillset that contains the language detection skill and the text translation skill

B. While including a language detection skill is important, it misses the requirement of translating the content, which is crucial for a multilingual knowledgebase.

Answer: C

Explanation:

The wiki contains text in English, French and Portuguese.

Scenario: All planned projects must support English, French, and Portuguese.

The Document Extraction skill extracts content from a file within the enrichment pipeline. This allows you to take advantage of the document extraction step that normally happens before the skillset execution with files that may be generated by other skills.

Note: The Translator Text API will be used to determine the from language. The Language detection skill is not required.

D. This option suggests using an Azure Cosmos DB indexer, which is not the right fit for wiki content stored as files. Also, the document extraction skill is not necessary here as the wiki files are already formatted in a readable manner.

Reference:

<https://docs.microsoft.com/en-us/azure/search/cognitive-search-skill-document-extraction>

<https://docs.microsoft.com/en-us/azure/search/cognitive-search-skill-text-translation>

Question: 17

You are developing the knowledgebase by using Azure Cognitive Search.

You need to meet the knowledgebase requirements for searching equivalent terms.

What should you include in the solution?

- A. a synonym map
- B. a suggester
- C. a custom analyzer
- D. a built-in key phrase extraction skill

Answer: A

Explanation:

Within a search service, synonym maps are a global resource that associate equivalent terms, expanding the scope of a query without the user having to actually provide the term. For example, assuming "dog", "canine", and "puppy" are mapped synonyms, a query on "canine" will match on a document containing "dog".

Create synonyms: A synonym map is an asset that can be created once and used by many indexes.

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-synonyms>

Topic 3, Misc. Questions

Question: 18

DRAG DROP

You have 100 chatbots that each has its own Language Understanding model.

Frequently, you must add the same phrases to each model.

You need to programmatically update the Language Understanding models to include the new phrases.

How should you complete the code? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Values
AddPhraseListAsync
Phraselist
PhraselistCreateObject
Phrases
SavePhraselistAsync
UploadPhraseListAsync

Answer Area

```
var phraselistId = await client.Features.  
(appId, versionId, new   
{  
    EnabledForAllModels = false,  
    IsExchangeable = true,  
    Name = "PL1",  
    Phrases = "item1,item2,item3,item4,item5"  
});
```

Explanation:

Values
AddPhraseListAsync
Phraselist
PhraselistCreateObject
Phrases
SavePhraselistAsync
UploadPhraseListAsync

Answer Area

```
var phraselistId = await client.Features. AddPhraseListAsync  
(appId, versionId, new  PhraselistCreateObject  
{  
    EnabledForAllModels = false,  
    IsExchangeable = true,  
    Name = "PL1",  
    Phrases = "item1,item2,item3,item4,item5"  
});
```

Box 1: AddPhraseListAsync
Example: Add phraselist feature

client.Features: This suggests you're using a LUIS authoring client, and interacting with its Features property.



```
var phraselistId = await client.Features.AddPhraseListAsync(appId, versionId, new  
PhraselistCreateObject  
{  
    EnabledForAllModels = false,  
    IsExchangeable = true,  
    Name = "QuantityPhraselist",  
    Phrases = "few,more,extra"  
});
```

AddPhraseListAsync

PhraselistCreateObject

This is where you need the correct class for creating a phrase list object.

Box 2: PhraselistCreateObject

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/client-libraries-rest-api>

Question: 19

DRAG DROP

You plan to use a Language Understanding application named app1 that is deployed to a container.

App1 was developed by using a Language Understanding authoring resource named lu1.

App1 has the versions shown in the following table.

Version	Trained date	Published date
V1.2	None	None
V1.1	2020-10-01	None
V1.0	2020-09-01	2020-09-15

You need to create a container that uses the latest deployable version of app1.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order. (Choose three.)

Actions	Answer Area
Run a container that has version set as an environment variable.	
Export the model by using the Export as JSON option.	
Select v1.1 of app1.	
Run a container and mount the model file.	
Select v1.0 of app1.	
Export the model by using the Export for containers (GZIP) option.	
Select v1.2 of app1.	

Answer:

Explanation:

v1.2: Trained but not published.

v1.1: Trained but not published.

v1.0: Trained and published (on 2020-09-15).

Therefore, the latest deployable version is v1.0

1. Select v1.0 of app1: You must first select the specific version that you want to deploy. In this case, that would be the published version v1.0.

2. Export model by using export contains (GZIP) option: The published model needs to be exported in a format that can be used by the containerized LUIS app. Gzip preferred for deploy in container

3. Run container and mount the model file: You should start your container

Actions	Answer Area
Run a container that has version set as an environment variable.	Export the model by using the Export for containers (GZIP) option.
Export the model by using the Export as JSON option.	Select v1.1 of app1.
Select v1.1 of app1.	Run a container and mount the model file.
Run a container and mount the model file.	
Select v1.0 of app1.	
Export the model by using the Export for containers (GZIP) option.	
Select v1.2 of app1.	

Step 1: Export the model using the Export for containers (GZIP) option.

Export versioned app's package from LUIS portal

The versioned app's package is available from the Versions list page.

Sign on to the LUIS portal.

Select the app in the list.

Select Manage in the app's navigation bar.

Select Versions in the left navigation bar.

Select the checkbox to the left of the version name in the list.

Select the Export item from the contextual toolbar above the list.

Select Export for container (GZIP).

The package is downloaded from the browser.

The screenshot shows the 'Versions' list page in the LUIS portal. At the top, there are buttons for 'Rename', 'Clone', and 'Export'. The 'Export' button is highlighted with a blue box. A dropdown menu from the 'Export' button lists 'Export as JSON' and 'Export for container (GZIP)'. Below the buttons, there is a search bar and filters for 'All' and 'Search for version(s)'. The main table lists versions with columns for 'Version name', 'Created', and 'Last modified'. One row is selected, showing '0.1 (Active & Production)' with creation and modification dates of '5/3/18' and '9/6/18' respectively.

Version name	Created	Last modified
0.1 (Active & Production)	5/3/18	9/6/18

Step 2: Select v1.1 of app1.

A trained or published app packaged as a mounted input to the container with its associated App ID.

Step 3: Run a contain and mount the model file.

Run the container, with the required input mount and billing settings.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-container-howto>

Question: 20

You need to build a chatbot that meets the following requirements:

Supports chit-chat, knowledge base, and multilingual models

Performs sentiment analysis on user messages

Selects the best language model automatically

What should you integrate into the chatbot?

- A. QnA Maker, Language Understanding, and Dispatch
- B. Translator, Speech, and Dispatch
- C. Language Understanding, Text Analytics, and QnA Maker
- D. Text Analytics, Translator, and Dispatch

the best model or service based on user input.

In your case, the dispatcher will determine whether to send a query to LUIS for intent recognition, to QnA Maker for knowledgebase lookup

text analytics can be done as add one but dispatch is important to route the query.

Answer: C

Explanation:

Language Understanding: An AI service that allows users to interact with your applications, bots, and IoT devices by using natural language.

QnA Maker is a cloud-based Natural Language Processing (NLP) service that allows you to create a natural conversational layer over your data.

a. It is used to find the most appropriate answer for any input from your custom knowledge base (KB) of information.

Text Analytics: Mine insights in unstructured text using natural language processing (NLP)—no machine learning expertise required. Gain a deeper understanding of customer opinions with sentiment analysis. The Language Detection feature of the Azure Text Analytics REST API evaluates text input

Incorrect Answers:

A, B, D: Dispatch uses sample utterances for each of your bot's different tasks (LUIS, QnA Maker, or custom), and builds a model that can be used to properly route your user's request to the right task, even across multiple bots.

Reference:

<https://azure.microsoft.com/en-us/services/cognitive-services/text-analytics/>

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/overview/overview>

Question: 21

Your company wants to reduce how long it takes for employees to log receipts in expense reports. All the receipts are in English.

You need to extract top-level information from the receipts, such as the vendor and the transaction total. The solution must minimize development effort.

Which Azure Cognitive Services service should you use?

- A. Custom Vision
- B. Personalizer
- C. Form Recognizer
- D. Computer Vision **OCP**

Answer: C

Explanation:

Azure Form Recognizer is a cognitive service that lets you build automated data processing software using machine learning technology. Identify and extract text, key/value pairs, selection marks, tables, and structure from your documents—the service outputs structured data that includes the relationships in the original file, bounding boxes, confidence and more.

Form Recognizer is composed of custom document processing models, prebuilt models for invoices, receipts, IDs and business cards, and the layout model.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/form-recognizer>

Question: 22

HOTSPOT

You need to create a new resource that will be used to perform sentiment analysis and optical character recognition (OCR). The solution must meet the following requirements:

- Use a single key and endpoint to access multiple services.
- Consolidate billing for future services that you might use.
- Support the use of Computer Vision in the future.

How should you complete the HTTP request to create the new resource? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

▼ https://management.azure.com/subscriptions/xxxxxxxx-xxxx-
PATCH
POST
PUT

xxxx-xxxx-
xxxxxxxxxxxx/resourceGroups/RG1/providers/Microsoft.CognitiveServices/
accounts/CS1?api-version=2017-04-18

{

 "location": "West US",
 "kind": "▼",
 CognitiveServices
 ComputerVision
 TextAnalytics

 "sku": {
 "name": "S0"

 },
 "properties": {},
 "identity": {
 "type": "SystemAssigned"

 }
}

Answer:

Explanation:

Answer Area

▼ https://management.azure.com/subscriptions/xxxxxxxx-xxxx-
PATCH
POST
PUT

POST

xxxx-xxxx-
xxxxxxxxxxxx/resourceGroups/RG1/providers/Microsoft.CognitiveServices/
accounts/CS1?api-version=2017-04-18

{

 "location": "West US",
 "kind": "▼",
 CognitiveServices
 ComputerVision
 TextAnalytics

 "sku": {
 "name": "S0"

 },
 "properties": {},
 "identity": {
 "type": "SystemAssigned"

 }

CognitiveServices

Box 1: PUT

Sample Request: PUT <https://management.azure.com/subscriptions/00000000-0000-0000-0000-000000000000/resourceGroups/test-rg/providers/Microsoft.DeviceUpdate/accounts/contoso?api-version=2020-03-01-preview>

Incorrect Answers:

PATCH is for updates.

Box 2: CognitiveServices

Microsoft Azure Cognitive Services provide us to use its pre-trained models for various Business Problems related to Machine Learning.

List of Different Services are:

Decision

Language (includes sentiment analysis)

Speech

Vision (includes OCR)

Web Search

Reference:

<https://docs.microsoft.com/en-us/rest/api/deviceupdate/resourcemanager/accounts/create>

<https://www.analyticsvidhya.com/blog/2020/12/microsoft-azure-cognitive-services-api-for-ai-development/>

Question: 23

You are developing a new sales system that will process the video and text from a public-facing website.

You plan to monitor the sales system to ensure that it provides equitable results regardless of the user's location or background.

Which two responsible AI principles provide guidance to meet the monitoring requirements? Each correct answer presents part of the solution. (Choose two.)

NOTE: Each correct selection is worth one point.

- A. transparency
- B. fairness
- C. inclusiveness
- D. reliability and safety
- E. privacy and security

Fairness and inclusiveness

Answer: BC

Explanation:

<https://docs.microsoft.com/en-us/learn/modules/get-started-ai-fundamentals/8-understand->

[responsible-ai](#)

Question: 24

DRAG DROP

You plan to use containerized versions of the Anomaly Detector API on local devices for testing and in on-premises datacenters.

You need to ensure that the containerized deployments meet the following requirements:

Prevent billing and API information from being stored in the command-line histories of the devices that run the container.

Control access to the container images by using Azure role-based access control (Azure RBAC).

Which four actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order. (Choose four.)

NOTE: More than one order of answer choices is correct. You will receive credit for any of the correct orders you select.

Actions	Answer Area
Create a custom Dockerfile.	1.Pull Anomaly Detector container image
Pull the Anomaly Detector container image.	2.create a custom docker file
Distribute a docker run script.	3.Build the image
Push the image to an Azure container registry.	4.Publish image on Azure Container Registry
Build the image.	
Push the image to Docker Hub.	

Answer:

Explanation:

Step 1: Pull the Anomaly Detector container image.

Step 2: Create a custom Dockerfile

Step 3: Build the image

Step 4: Push the image to an Azure container registry.

<https://docs.microsoft.com/en-us/azure/cognitive-services/containers/container-reuse-recipe>

Question: 25

HOTSPOT

You plan to deploy a containerized version of an Azure Cognitive Services service that will be used for text analysis.

You configure https://contoso.cognitiveservices.azure.com as the endpoint URI for the service, and you pull the latest version of the Text Analytics Sentiment Analysis container.

You need to run the container on an Azure virtual machine by using Docker.

How should you complete the command? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

```
docker run --rm -it -p 5000:5000 --memory 8g --cpus 1 \
```

```
http://contoso.blob.core.windows.net
https://contoso.cognitiveservices.azure.com
mcr.microsoft.com/azure-cognitive-services/textanalytics/keyphrase
mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment
```

Eula=accept \

```
Billing=
```

```
http://contoso.blob.core.windows.net
https://contoso.cognitiveservices.azure.com
mcr.microsoft.com/azure-cognitive-services/textanalytics/keyphrase
mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment
```

ApiKey=xxxxxxxxxxxxxxxxxxxxxx

**docker run -rm -it -p 5000:5000 **

Answer:

Explanation: **mcr.microsoft.com/azure-cognitive-services/text-analytics/sentiment \ Eula=accept \ Billing=https://contoso.cognitiveservices.azure.com**

```
docker run --rm -it -p 5000:5000 --memory 8g --cpus 1 \
```

```
http://contoso.blob.core.windows.net
https://contoso.cognitiveservices.azure.com
mcr.microsoft.com/azure-cognitive-services/textanalytics/keyphrase
mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment
```

Eula=accept \ **container image name**

```
Billing=
```

```
http://contoso.blob.core.windows.net
https://contoso.cognitiveservices.azure.com
mcr.microsoft.com/azure-cognitive-services/textanalytics/keyphrase
mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment
```

ApiKey=xxxxxxxxxxxxxxxxxxxxxx

Specifies the billing endpoint for the container

Box 1: **mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment**

To run the Sentiment Analysis v3 container, execute the following docker run command.

```
docker run --rm -it -p 5000:5000 --memory 8g --cpus 1 \
mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment \
Eula=accept \
Billing={ENDPOINT_URI} \
ApiKey={API_KEY} is the endpoint for accessing the Text Analytics API. https://<your-custom-
subdomain>.cognitiveservices.azure.com
```

Box 2: **https://contoso.cognitiveservices.azure.com**

{ENDPOINT_URI} is the endpoint for accessing the Text Analytics API: https://<your-custom-

subdomain>.cognitiveservices.a The endpoint for accessing the Text Analytics API. zure.com

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/text-analytics/how-tos/text-analytics-how-to-install-containers?tabs=sentiment>

Question: 26

You have the following C# method for creating Azure Cognitive Services resources programmatically.

```
static void create_resource(CognitiveServicesManagementClient client, string
resource_name, string kind, string account_tier, string location)
{
    CognitiveServicesAccount parameters =
        new CognitiveServicesAccount(null, null, kind, location, resource_name,
new CognitiveServicesAccountProperties(), new Sku(account_tier));
    var result = client.Accounts.Create(resource_group_name, account_tier,
parameters);
}
```

You need to call the method to create a free Azure resource in the West US Azure region. The resource will be used to generate captions of images automatically.

Which code should you use?

F0 - Free

- A. create_resource(client, "res1", "ComputerVision", "F0", "westus")
- B. create_resource(client, "res1", "CustomVision.Prediction", "F0", "westus")
- C. create_resource(client, "res1", "ComputerVision", "S0", "westus") standard
- D. create_resource(client, "res1", "CustomVision.Prediction", "S0", "westus")

Answer: B

Explanation:

<https://azure.microsoft.com/en-us/pricing/details/cognitive-services/computer-vision/>

Question: 27

You successfully run the following HTTP request.

```
POST https://management.azure.com/subscriptions/18c51a87-3a69-47a8-aedc-
a54745f708a1/resourceGroups/RG1/providers/Microsoft.CognitiveServices/accounts/contosol/rege-
nerateKey?api-version=2017-04-18
Body>{"keyName": "Key2"}
```

What is the result of the request?

- A. A key for Azure Cognitive Services was generated in Azure Key Vault.
- B. A new query key was generated.
- C. The primary subscription key and the secondary subscription key were rotated.
- D. The secondary subscription key was reset.

Answer: B

Explanation:

<https://docs.microsoft.com/en-us/rest/api/searchmanagement/2021-04-01-preview/query-keys/create>

Question: 28

You build a custom Form Recognizer model.

You receive sample files to use for training the model as shown in the following table.

Name	Type	Size
File1	PDF	20 MB
File2	MP4	100 MB
File3	JPG	20 MB
File4	PDF	100 MB
File5	GIF	1 MB
File6	JPG	40 MB

Which three files can you use to train the model? Each correct answer presents a complete solution.
(Choose three.)

NOTE: Each correct selection is worth one point.

- A. File1
- B. File2
- C. File3
- D. File4
- E. File5
- F. File6

Answer: A, B, F

Explanation:

Input requirements

Form Recognizer works on input documents that meet these requirements:

Format must be JPG, PNG, PDF (text or scanned), or TIFF. Text-embedded PDFs are best because there's no possibility of error in character extraction and location.

File size must be less than 50 MB.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/form-recognizer/overview>

Question: 29

You have a Video Indexer service that is used to provide a search interface over company videos on your company's website.

You need to be able to search for videos based on who is present in the video. What should you do?

- A. Create a person model and associate the model to the videos.
- B. Create person objects and provide face images for each object.
- C. Invite the entire staff of the company to Video Indexer.
- D. Edit the faces in the videos.
- E. Upload names to a language model.

Answer: A

Explanation:

Video Indexer supports multiple Person models per account. Once a model is created, you can use it by providing the model ID of a specific Person model when uploading/indexing or reindexing a video. Training a new face for a video updates the specific custom model that the video was associated with.

Note: Video Indexer supports face detection and celebrity recognition for video content. The celebrity recognition feature covers about one million faces based on commonly requested data source such as IMDB, Wikipedia, and top LinkedIn influencers. Faces that aren't recognized by the celebrity recognition feature are detected but left unnamed. Once you label a face with a name, the face and name get added to your account's Person model. Video Indexer will then recognize this face in your future videos and past videos.

Reference:

<https://docs.microsoft.com/en-us/azure/media-services/video-indexer/customize-person-model-with-api>

Question: 30

You use the Custom Vision service to build a classifier.

After training is complete, you need to evaluate the classifier.

Which two metrics are available for review? Each correct answer presents a complete solution.
(Choose two.)

NOTE: Each correct selection is worth one point.

- A. recall **it's a single value that combines precision and recall, but it's not directly available in the Custom Vision evaluation interface as two separate metrics. F-score is calculated based on precision and recall.**
- B. F-score **Weighted accuracy is a metric that accounts for class imbalances**
- C. weighted accuracy **Weighted accuracy is a metric that accounts for class imbalances**
- D. precision
- E. area under the curve (AUC) **AUC is a metric used to evaluate binary classifiers. Custom Vision supports multi-class classification**

Answer: AD

Explanation:

Custom Vision provides three metrics regarding the performance of your model: precision, recall, and AP.

Reference:

<https://www.tallan.com/blog/2020/05/19/azure-custom-vision/>

Question: 31

DRAG DROP

You are developing a call to the Face API. The call must find similar faces from an existing list named employefaces. The employefaces list contains 60,000 images.

How should you complete the body of the HTTP request? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Values
"faceListId"
"LargeFaceListId"
"matchFace"
"matchPerson"

Answer Area

```
{  
  "faceId": "18c51a87-3a69-47a8-aedc-a54745f708a1",  
  [REDACTED]: "employefaces",  
  "maxNumOfCandidatesReturned": 1,  
  "mode": [REDACTED]  
}
```

Answer:

Explanation:

Values
"faceListId"
"LargeFaceListId"
"matchFace"
"matchPerson"

Answer Area

```
{  
  "faceId": "18c51a87-3a69-47a8-aedc-a54745f708a1",  
  "LargeFaceListId": "employefaces",  
  "maxNumOfCandidatesReturned": 1,  
  "mode": "matchFace"  
}
```



Box 1: LargeFaceListID

LargeFaceList: Add a face to a specified large face list, up to 1,000,000 faces.

Note: Given query face's faceId, to search the similar-looking faces from a faceId array, a face list or a

large face list. A "faceListId" is created by FaceList - Create containing persistedFacIds that will not expire. And a "largeFaceListId" is created by LargeFaceList - Create containing persistedFacIds that will also not expire.

Incorrect Answers:

Not "faceListId": Add a face to a specified face list, up to 1,000 faces.

Box 2: matchFace

Find similar has two working modes, "matchPerson" and "matchFace". "matchPerson" is the default mode that it tries to find faces of the same person as possible by using internal same-person thresholds. It is useful to find a known person's other photos. Note that an empty list will be returned if no faces pass the internal thresholds. "matchFace" mode ignores same-person thresholds and returns ranked similar faces anyway, even the similarity is low. It can be used in the cases like searching celebrity-looking faces.

Reference:

<https://docs.microsoft.com/en-us/rest/api/faceapi/face/findsimilar>

Question: 32

DRAG DROP

You are developing a photo application that will find photos of a person based on a sample image by using the Face API.

You need to create a POST request to find the photos.

How should you complete the request? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Values	Answer Area
detect	POST {Endpoint}/face/v1.0/ <input type="text"/>
findsimilar	Request Body
group	
identify	
matchFace	
matchPerson	
verify	

```
{  
    "faceId": "c5c24a82-6845-4031-9d5d-978df9175426",  
    "largeFaceListId": "sample_list",  
    "largeFaceListId": "sample_list",  
    "maxNumOfCandidatesReturned": 10,  
    "mode": ""  
}
```

Explanation:

Box 1: findsimilar

1. /findsimilar
endpoint is used to find faces similar to a target face within a face list.

2.matchFace -It must be set to matchFace when searching a large face list.

Answer:

<https://docs.microsoft.com/en-us/rest/api/faceapi/face/find-similar>

Box 2: matchPerson

Find similar has two working modes, "matchPerson" and "matchFace". "matchPerson" is the default mode that it tries to find faces of the same person as possible by using internal same-person thresholds. It is useful to find a known person's other photos. Note that an empty list will be returned if no faces pass the internal thresholds. "matchFace" mode ignores same-person thresholds and returns ranked similar faces anyway, even the similarity is low. It can be used in the cases like searching celebrity-looking faces.

Reference:

<https://docs.microsoft.com/en-us/rest/api/faceapi/face/detectwithurl>

<https://docs.microsoft.com/en-us/rest/api/faceapi/face/findsimilar>

Question: 33

HOTSPOT

You develop a test method to verify the results retrieved from a call to the Computer Vision API. The call is used to analyze the existence of company logos in images. The call returns a collection of brands named brands.

You have the following code segment.

```
foreach (var brand in brands)
{
    if (brand.Confidence >= .75)
        Console.WriteLine($"Logo of {brand.Name} between {brand.Rectangle.X},
{brand.Rectangle.Y} and {brand.Rectangle.W}, {brand.Rectangle.H}");
}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Statements	Yes	No
The code will return the name of each detected brand with a confidence equal to or higher than 75 percent.	<input checked="" type="radio"/>	<input type="radio"/>
The code will return coordinates for the bottom-left corner of the rectangle that contains the brand logo of the displayed brands.	<input checked="" type="radio"/>	<input type="radio"/>
The code will return coordinates for the bottom-right corner of the rectangle that contains the brand logo of the displayed brands.	<input type="radio"/>	<input checked="" type="radio"/>

Answer:

Explanation:

Box 1: Yes

Box 2: No

Box 3: No

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/computer-vision/concept-detecting-faces>

Question: 34

HOTSPOT

You develop an application that uses the Face API.


You need to add multiple images to a person group.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
Parallel.For(0, PersonCount, async i =>
{
    Guid personId = persons[i].PersonId;
    string personImageDir = $"{path/to/person/{i}/images}";
    foreach (string imagePath in Directory.GetFiles(personImageDir, "*.jpg"))
    {
        using (File t = File.OpenRead(imagePath))
        {
            await faceClient.PersonGroupPerson.
                (personGroupId, personId, t);
        }
    }
});
```



Answer:

Explanation:

Box 1: Stream

The File.OpenRead(String) method opens an existing file for reading.

Example: Open the stream and read it back.

```
using (FileStream fs = File.OpenRead(path))
```

Box 2: AddFaceFromStreamAsync

Step 5 on <https://docs.microsoft.com/en-us/azure/cognitive-services/face/face-api-how-to-topics/how-to-add-faces>

Question: 35**HOTSPOT**

You are developing an application that will use the Computer Vision client library. The application has the following code.

```
public async Task>AnalyzeImage(ComputerVisionClient client, string localImage)
{
    List<VisualFeatureTypes> features = new List<VisualFeatureTypes>()
    {
        VisualFeatureTypes.Description,
        VisualFeatureTypes.Tags,
    };
    using (Stream imageStream = File.OpenRead(localImage))
    {
        try
        {
            ImageAnalysis results = await client.AnalyzeImageInStreamAsync(imageStream, features);

            foreach (var caption in results.Description.Captions)
            {
                Console.WriteLine($"{caption.Text} with confidence {caption.Confidence}");
            }

            foreach (var tag in results.Tags)
            {
                Console.WriteLine($"{tag.Name} {tag.Confidence}");
            }
        }
        catch (Exception ex)
        {
            Console.WriteLine(ex.Message);
        }
    }
}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
The code will perform face recognition.	<input type="radio"/>	<input checked="" type="radio"/>
The code will list tags and their associated confidence.	<input checked="" type="radio"/>	<input type="radio"/>
The code will read a file from the local file system.	<input checked="" type="radio"/>	<input type="radio"/>

Answer:

Explanation:

Box 1: No

Box 2: Yes

Box 3: No

Type text here

Question: 36

You are developing a method that uses the Computer Vision client library. The method will perform optical character recognition (OCR) in images. The method has the following code.

```
public static async Task ReadImageUrl(ComputerVisionClient client, string urlFile)
{
    const int numberOfCharsInOperationId = 36;

    var txtHeaders = await client.ReadAsync(urlFile, language: "en");

    string opLocation = txtHeaders.OperationLocation;
    string operationId = opLocation.Substring(opLocation.Length -
    numberOfCharsInOperationId);

    ReadOperationResult results;

    results = await client.GetReadResultAsync(Guid.Parse(operationId));

    var textUrlFileResults = results.AnalyzeResult.ReadResults;
    foreach (ReadResult page in textUrlFileResults)
    {
        foreach (Line line in page.Lines)
        {
            Console.WriteLine(line.Text);
        }
    }
}
```

During testing, you discover that the call to the GetReadResultAsync method occurs before the read operation is complete.

You need to prevent the GetReadResultAsync method from proceeding until the read operation is complete.

Which two actions should you perform? Each correct answer presents part of the solution. (Choose two.)

NOTE: Each correct selection is worth one point.

- A. Remove the Guid.Parse(operationId) parameter.
- B. Add code to verify the results.Status value.
- C. Add code to verify the status of the txtHeaders.Status value.
- D. Wrap the call to GetReadResultAsync within a loop that contains a delay.

Answer: BD

Explanation:

Example code :

```
do
{
    results = await client.GetReadResultAsync(Guid.Parse(operationId));
}
while ((results.Status == OperationStatusCodes.Running ||
    results.Status == OperationStatusCodes.NotStarted));
```

Reference:

<https://github.com/Azure-Samples/cognitive-services-quickstart-code/blob/master/dotnet/ComputerVision/ComputerVisionQuickstart.cs>

Question: 37

You are building a Language Understanding model for an e-commerce platform. You need to construct an entity to capture billing addresses.

Which entity type should you use for the billing address?

- A. machine learned
- B. Regex
- C. geographyV2
- D. Pattern.any
- E. list

Answer: A

Explanation:

An ML entity can be composed of smaller sub-entities, each of which can have its own properties. For example, Address could have the following structure:

Address: 4567 Main Street, NY, 98052, USA

Building Number: 4567

Street Name: Main Street

State: NY

Zip Code: 98052

Country: USA

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-concept-entity-types>

Question: 38

You need to upload speech samples to a Speech Studio project. How should you upload the samples?

- A. Combine the speech samples into a single audio file in the .wma format and upload the file.
- B. Upload a .zip file that contains a collection of audio files in the .wav format and a corresponding text transcript file.
- C. Upload individual audio files in the FLAC format and manually upload a corresponding transcript in Microsoft Word format.
- D. Upload individual audio files in the .wma format.

Answer: B

Explanation:

To upload your data, navigate to the Speech Studio . From the portal, click Upload data to launch the wizard and create your first dataset. You'll be asked to select a speech data type for your dataset, before allowing you to upload your data.

The default audio streaming format is WAV

Use this table to ensure that your audio files are formatted correctly for use with Custom Speech:

Property	Value
File format	RIFF (WAV)
Sample rate	8,000 Hz or 16,000 Hz
Channels	1 (mono)
Maximum length per audio	2 hours
Sample format	PCM, 16-bit
Archive format	.zip
Maximum archive size	2 GB

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/speech-service/how-to-custom-speech-test-and-train>

Question: 39

You are developing a method for an application that uses the Translator API.

The method will receive the content of a webpage, and then translate the content into Greek (el).

The result will also contain a transliteration that uses the Roman alphabet.

You need to create the URI for the call to the Translator API. You have the following URI.

<https://api.cognitive.microsofttranslator.com/translate?api-version=3.0>

Which three additional query parameters should you include in the URI? Each correct answer presents part of the solution. (Choose three.)

NOTE: Each correct selection is worth one point.

- A. toScript=Cyril
- B. from=el
- C. textType=html
- D. to=el
- E. textType=plain
- F. toScript=Latn

Answer: A, D, F

Explanation:

C: textType is an optional parameter. It defines whether the text being translated is plain text or HTML text (used for web pages).

D: to is a required parameter. It specifies the language of the output text. The target language must be one of the supported languages included in the translation scope.

F: toScript is an optional parameter. It specifies the script of the translated text. We use Latin (Roman alphabet) script.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/translator/reference/v3-0-translate>

Question: 40

You have a chatbot that was built by using the Microsoft Bot Framework. You need to debug the chatbot endpoint remotely.

Which two tools should you install on a local computer? Each correct answer presents part of the solution. (Choose two.)

NOTE: Each correct selection is worth one point.

- A. Fiddler
- B. Bot Framework Composer
- C. Bot Framework Emulator
- D. Bot Framework CLI
- E. ngrok
- F. nginx

Answer: CE

Explanation:

Bot Framework Emulator is a desktop application that allows bot developers to test and debug bots, either locally or remotely.

ngrok is a cross-platform application that "allows you to expose a web server running on your local machine to the internet." Essentially, what we'll be doing is using ngrok to forward messages from external channels on the web directly to our local machine to allow debugging, as opposed to the

standard messaging endpoint configured in the Azure portal.

Reference:

<https://docs.microsoft.com/en-us/azure/bot-service/bot-service-debug-emulator>

Question: 41

DRAG DROP

You are building a retail chatbot that will use a QnA Maker service.

You upload an internal support document to train the model. The document contains the following question "What is your warranty period?"

Users report that the chatbot returns the default QnA Maker answer when they ask the following question "How long is the warranty coverage?"

The chatbot returns the correct answer when the users ask the following question 'What is your warranty period?"

Both questions should return the same answer.

You need to increase the accuracy of the chatbot responses.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order. (Choose three.)

Actions	Answer Area
Add a new question and answer (QnA) pair.	
Retrain the model.	
Add additional questions to the document.	
Republish the model.	
Add alternative phrasing to the question and answer (QnA) pair.	

Answer:

Explanation:

Add alternative phrasing to the question and answer (QnA) pair.



Retrain the model.



Republish the model.



Step 1: Add alternative phrasing to the question and answer (QnA) pair.

Add alternate questions to an existing QnA pair to improve the likelihood of a match to a user query.

Step 2: Retrain the model.

Periodically select Save and train after making edits to avoid losing changes.

Step 3: Republish the model

Note: A knowledge base consists of question and answer (QnA) pairs. Each pair has one answer and a pair contains all the information associated with that answer.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/how-to/edit-knowledge-base>

Question: 42

You need to measure the public perception of your brand on social media messages. Which Azure Cognitive Services service should you use?


- A. Text Analytics
- B. Content Moderator
- C. Computer Vision
- D. Form Recognizer

Answer: A

Explanation:

Text Analytics Cognitive Service could be used to quickly determine the public perception for a specific topic, event or brand.

Example: A NodeJS app which pulls Tweets from Twitter using the Twitter API based on a specified search term. Then pass these onto Text Analytics for sentiment scoring before storing the data and building a visualisation in PowerBI. The Architecture looked something like this:



Reference:

<https://www.linkedin.com/pulse/measuring-public-perception-azure-cognitive-services-steve-dalai>

Question: 43

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You build a language model by using a Language Understanding service. The language model is used to search for information on a contact list by using an intent named FindContact.

A conversational expert provides you with the following list of phrases to use for training.

Find contacts in London. Who do I know in Seattle?

Search for contacts in Ukraine.

You need to implement the phrase list in Language Understanding.

Solution: You create a new intent for location.

Does this meet the goal?

- A. Yes
- B. No

Answer: A

Explanation:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-concept-intent>

Question: 44

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You build a language model by using a Language Understanding service. The language model is used to search for information on a contact list by using an intent named FindContact.

A conversational expert provides you with the following list of phrases to use for training.

Find contacts in London.

Who do I know in Seattle? Search for contacts in Ukraine.

You need to implement the phrase list in Language Understanding.

Solution: You create a new entity for the domain.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation:

Instead use a new intent for location.

Note: An intent represents a task or action the user wants to perform. It is a purpose or goal expressed in a user's utterance.

Define a set of intents that corresponds to actions users want to take in your application. Reference:
<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-concept-intent>

Question: 45

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You build a language model by using a Language Understanding service. The language model is used to search for information on a contact list by using an intent named FindContact.

A conversational expert provides you with the following list of phrases to use for training.

Find contacts in London. Who do I know in Seattle?
Search for contacts in Ukraine.

You need to implement the phrase list in Language Understanding.

Solution: You create a new pattern in the FindContact intent.

Does this meet the goal?

- A. Yes
- B. No**

Answer: B

Explanation:

Instead use a new intent for location.

Note: An intent represents a task or action the user wants to perform. It is a purpose or goal expressed in a user's utterance.

Define a set of intents that corresponds to actions users want to take in your application. Reference: <https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-concept-intent>

Question: 46

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You develop an application to identify species of flowers by training a Custom Vision model. You receive images of new flower species.

You need to add the new images to the classifier.

Solution: You add the new images, and then use the Smart Labeler tool.

Does this meet the goal?

- A. Yes
- B. No**

Answer: B

Explanation:

The model need to be extended and retrained.

Note: Smart Labeler to generate suggested tags for images. This lets you label a large number of images more quickly when training a Custom Vision model.

Question: 47

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You develop an application to identify species of flowers by training a Custom Vision model. You receive images of new flower species.

You need to add the new images to the classifier.

Solution: You add the new images and labels to the existing model. You retrain the model, and then publish the model.

Does this meet the goal?

- A. Yes
- B. No

Answer: A

Explanation:

The model needs to be extended and retrained.

Question: 48

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You develop an application to identify species of flowers by training a Custom Vision model. You receive images of new flower species.

You need to add the new images to the classifier.

Solution: You create a new model, and then upload the new images and labels.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation:

The model needs to be extended and retrained.

Question: 49

HOTSPOT

You are developing a service that records lectures given in English (United Kingdom).

You have a method named AppendToTranscriptFile that takes translated text and a language identifier.

You need to develop code that will provide transcripts of the lectures to attendees in their respective language. The supported languages are English, French, Spanish, and German.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
static async Task TranslateSpeechAsync()
{
    var config = SpeechTranslationConfig.FromSubscription("69cad5cc-0ab3-4704-bdff-afbf4aa07d85", "uksouth");

    var lang = new List<string>
    {
        {"en-GB"}
        {"fr", "de", "es"}
        {"French", "Spanish", "German"}
        {languages}
    }

    config.SpeechRecognitionLanguage = "en-GB";
    lang.ForEach(config.AddTargetLanguage);

    using var audioConfig = AudioConfig.FromDefaultMicrophoneInput();
    using var recognizer = new IntentRecognizer (config, audioConfig);

    var result = await recognizer.RecognizeOnceAsync();
    if (result.Reason == ResultReason.TranslatedSpeech)
```

Answer:**Explanation:****Answer Area**

```
static async Task TranslateSpeechAsync()
{
    var config = SpeechTranslationConfig.FromSubscription("69cad5cc-0ab3-4704-bdff-afbf4aa07d85", "uksouth");

    var lang = new List<string>
    {
        {"en-GB"}
        {"fr", "de", "es"} Box 1: {"fr", "de", "es"}
        {"French", "Spanish", "German"}
        {languages}
    }

    config.SpeechRecognitionLanguage = "en-GB";
    lang.ForEach(config.AddTargetLanguage);

    using var audioConfig = AudioConfig.FromDefaultMicrophoneInput();
    using var recognizer = new TranslationRecognizer (config, audioConfig);

    var result = await recognizer.RecognizeOnceAsync();
    if (result.Reason == ResultReason.TranslatedSpeech)
```

Box 1: {"fr", "de", "es"}

A common task of speech translation is to specify target translation languages, at least one is

required but multiples are supported. The following code snippet sets both French and German as translation language targets.

```
static async Task TranslateSpeechAsync()
{
    var translationConfig =
        SpeechTranslationConfig.FromSubscription(SPEECH__SUBSCRIPTION__KEY,
SPEECH__SERVICE__REGION);

    translationConfig.SpeechRecognitionLanguage = "it-IT";

    // Translate to languages. See, https://aka.ms/speech/sttt-languages
    translationConfig.AddTargetLanguage("fr");
    translationConfig.AddTargetLanguage("de");
}
```

Box 2: TranslationRecognizer

After you've created a SpeechTranslationConfig, the next step is to initialize a TranslationRecognizer.

Example code:

```
static async Task TranslateSpeechAsync()
{
    var translationConfig =
        SpeechTranslationConfig.FromSubscription(SPEECH__SUBSCRIPTION__KEY,
SPEECH__SERVICE__REGION);

    var fromLanguage = "en-US";
    var toLanguages = new List<string> { "it", "fr", "de" };
    translationConfig.SpeechRecognitionLanguage = fromLanguage;
    toLanguages.ForEach(translationConfig.AddTargetLanguage);

    using var recognizer = new TranslationRecognizer(translationConfig);
}
```

Question: 50

HOTSPOT

You are developing a text processing solution.

You develop the following method.

```
static void GetKeyPhrases(TextAnalyticsClient textAnalyticsClient, string text)
{
    var response = textAnalyticsClient.ExtractKeyPhrases(text);
    Console.WriteLine("Key phrases:");

    foreach (string keyphrase in response.Value)
    {
        Console.WriteLine($"\\t{keyphrase}");
    }
}
```

You call the method by using the following code.

```
GetKeyPhrases(textAnalyticsClient, "the cat sat on the mat");
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
The call will output key phrases from the input string to the console.	<input checked="" type="radio"/>	<input type="radio"/>
The output will contain the following words: the, cat, sat, on, and mat.	<input type="radio"/>	<input checked="" type="radio"/>
The output will contain the confidence level for key phrases.	<input type="radio"/>	<input checked="" type="radio"/>

Answer:

Explanation:

Answer Area

Statements	Yes	No
The call will output key phrases from the input string to the console.	<input checked="" type="radio"/>	<input type="radio"/>
The output will contain the following words: the, cat, sat, on, and mat.	<input type="radio"/>	<input checked="" type="radio"/>
The output will contain the confidence level for key phrases.	<input type="radio"/>	<input checked="" type="radio"/>

Box 1: Yes

The Key Phrase Extraction API evaluates unstructured text, and for each JSON document, returns a list of key phrases.

Box 2: No

'the' is not a key phrase.

This capability is useful if you need to quickly identify the main points in a collection of documents. For example, given input text "The food was delicious and there were wonderful staff", the service returns the main talking points: "food" and "wonderful staff".

Box 3: No

Key phrase extraction does not have confidence levels.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/text-analytics/how-tos/text-analytics-how-to-keyword-extraction>

Question: 51

You deploy a web app that is used as a management portal for indexing in Azure Cognitive Search. The app is configured to use the primary admin key.

During a security review, you discover unauthorized changes to the search index. You suspect that the primary access key is compromised.

You need to prevent unauthorized access to the index management endpoint. The solution must minimize downtime.

What should you do next?

- A. Regenerate the primary admin key, change the app to use the secondary admin key, and then regenerate the secondary admin key.
- B. Change the app to use a query key, and then regenerate the primary admin key and the secondary admin key.
- C. Regenerate the secondary admin key, change the app to use the secondary admin key, and then regenerate the primary key.
- D. Add a new query key, change the app to use the new query key, and then delete all the unused query keys.

Answer: C

Explanation:

Question: 52

You have an existing Azure Cognitive Search service.

You have an Azure Blob storage account that contains millions of scanned documents stored as images and PDFs.

You need to make the scanned documents available to search as quickly as possible. What should you do?

- A. Split the data into multiple blob containers. Create a Cognitive Search service for each container. Within each indexer definition, schedule the same runtime execution pattern.
- B. Split the data into multiple blob containers. Create an indexer for each container. Increase the search units. Within each indexer definition, schedule a sequential execution pattern.
- C. Create a Cognitive Search service for each type of document.
- D. Split the data into multiple virtual folders. Create an indexer for each folder. Increase the search

units. Within each indexer definition, schedule the same runtime execution pattern.

Answer: D

Explanation:

Incorrect Answers:

A: Need more search units to process the data in parallel. B: Run them in parallel, not sequentially.

C: Need a blob indexer.

Note: A blob indexer is used for ingesting content from Azure Blob storage into a Cognitive Search index. Index large datasets

Indexing blobs can be a time-consuming process. In cases where you have millions of blobs to index, you can speed up indexing by partitioning your data and using multiple indexers to process the data in parallel. Here's how you can set this up:

Partition your data into multiple blob containers or virtual folders Set up several data sources, one per container or folder.

Create a corresponding indexer for each data source. All of the indexers should point to the same target search index.

One search unit in your service can run one indexer at any given time. Creating multiple indexers as described above is only useful if they actually run in parallel.

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-howto-indexing-azure-blob-storage>

Question: 53

You need to implement a table projection to generate a physical expression of an Azure Cognitive Search index.

Which three properties should you specify in the skillset definition JSON configuration table node?

Each correct answer presents part of the solution. (Choose three.)

NOTE: Each correct selection is worth one point.

- A. tableName
- B. generatedKeyName
- C. dataSource
- D. dataSourceConnection
- E. source

Answer: ABE

Explanation:

Defining a table projection.

Each table requires three properties:

tableName: The name of the table in Azure Storage.

generatedKeyName: The column name for the key that uniquely identifies this row.

source: The node from the enrichment tree you are sourcing your enrichments from. This node is

usually the output of a shaper, but could be the output of any of the skills.

Reference:

<https://docs.microsoft.com/en-us/azure/search/knowledge-store-projection-overview>

Question: 54

HOTSPOT

You are creating an enrichment pipeline that will use Azure Cognitive Search. The knowledge store contains unstructured JSON data and scanned PDF documents that contain text.

Which projection type should you use for each data type? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area


JSON data:	<ul style="list-style-type: none">File projectionObject projectionTable projection
Scanned data:	<ul style="list-style-type: none">File projectionObject projectionTable projection

Answer:


Explanation:

Answer Area

JSON data:



Scanned data:



Box 1: Object projection

Object projections are JSON representations of the enrichment tree that can be sourced from any node.

Box 2: File projection

File projections are similar to object projections and only act on the normalized_images collection.

Reference:

<https://docs.microsoft.com/en-us/azure/search/knowledge-store-projection-overview>

Question: 55

HOTSPOT

You are building an Azure Cognitive Search custom skill.

You have the following custom skill schema definition.

```
{  
    "@odata.type": "#Microsoft.Skills.Custom.WebApiSkill",  
    "description": "My custom skill description",  
    "uri": "https://contoso-webskill.azurewebsites.net/api/process",  
    "context": "/document/organizations/*",  
    "inputs": [  
        {  
            "name": "companyName",  
            "source": "/document/organizations/*"  
        }  
    ],  
    "outputs": [  
        {  
            "name": "companyDescription",  
        }  
    ]  
}
```

For each of the following statements, select Yes if the statement. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
CompanyDescription is available for indexing.	<input checked="" type="radio"/>	<input type="radio"/>
The definition calls a web API as part of the enrichment process.	<input checked="" type="radio"/>	<input type="radio"/>
The enrichment step is called only for the first organization under "/document/organizations".	<input type="radio"/>	<input checked="" type="radio"/>

Answer:

Explanation:

Answer Area

Statements	Yes	No
CompanyDescription is available for indexing.	<input type="radio"/>	<input checked="" type="radio"/>
The definition calls a web API as part of the enrichment process.	<input type="radio"/>	<input checked="" type="radio"/>
The enrichment step is called only for the first organization under "/document/organizations".	<input checked="" type="radio"/>	<input type="radio"/>

Box 1: Yes

Once you have defined a skillset, you must map the output fields of any skill that directly contributes values to a given field in your search index.

Box 2: Yes

The definition is a custom skill that calls a web API as part of the enrichment process.

Box 3: No

For each organization identified by entity recognition, this skill calls a web API to find the description of that organization.

Reference:

<https://docs.microsoft.com/en-us/azure/search/cognitive-search-output-field-mapping>

Question: 56

You have the following data sources:

Finance: On-premises Microsoft SQL Server database

Sales: Azure Cosmos DB using the Core (SQL) API

Logs: Azure Table storage

HR: Azure SQL database

You need to ensure that you can search all the data by using the Azure Cognitive Search REST API.

What should you do?

- A. Configure multiple read replicas for the data in Sales.
- B. Mirror Finance to an Azure SQL database.
- C. Migrate the data in Sales to the MongoDB API.
- D. Ingest the data in Logs into Azure Sentinel.

Answer: B

Explanation:

On-premises Microsoft SQL Server database cannot be used as an index data source.

Note: Indexer in Azure Cognitive Search: : Automate aspects of an indexing operation by configuring a data source and an indexer that you can schedule or run on demand. This feature is supported for a limited number of data source types on Azure.

Indexers crawl data stores on Azure.

Azure Blob Storage

Azure Data Lake Storage Gen2 (in preview)

Azure Table Storage

Azure Cosmos DB

Azure SQL Database

SQL Managed Instance

SQL Server on Azure Virtual Machines

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-indexer-overview#supported-data-sources>

Question: 57

You are building a multilingual chatbot.

You need to send a different answer for positive and negative messages.

Which two Text Analytics APIs should you use? Each correct answer presents part of the solution.
(Choose two.)

NOTE: Each correct selection is worth one point.

- A. Linked entities from a well-known knowledge base
- B. Sentiment Analysis
- C. Key Phrases
- D. Detect Language
- E. Named Entity Recognition

Answer: BD

Explanation:

B: The Text Analytics API's Sentiment Analysis feature provides two ways for detecting positive and negative sentiment. If you send a Sentiment Analysis request, the API will return sentiment labels (such as "negative", "neutral" and "positive") and confidence scores at the sentence and document-level.

D: The Language Detection feature of the Azure Text Analytics REST API evaluates text input for each document and returns language identifiers with a score that indicates the strength of the analysis. This capability is useful for content stores that collect arbitrary text, where language is unknown.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/text-analytics/how-tos/text-analytics-how-to-sentiment-analysis?tabs=version-3-1>

<https://docs.microsoft.com/en-us/azure/cognitive-services/text-analytics/how-tos/text-analytics-how-to-language-detection>

Question: 58

DRAG DROP

You plan to build a chatbot to support task tracking.

You create a Language Understanding service named lu1.

You need to build a Language Understanding model to integrate into the chatbot. The solution must minimize development time to build the model.

Which four actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order. (Choose four.)

Actions	Answer Area
Train the application.	3
Publish the application.	4
Add a new application.	1
Add example utterances.	2
Add the prebuilt domain ToDo.	

Answer:

Explanation:

1. Add a new application
2. Add a prebuilt domain intent ToDo (it has already utterances so we can skip this step)
3. Train
4. Publish

Also Correct

Question: 59

You are building a bot on a local computer by using the Microsoft Bot Framework. The bot will use an existing Language Understanding model.

You need to translate the Language Understanding model locally by using the Bot Framework CLI.

What should you do first?

- A. From the Language Understanding portal, clone the model.
- B. Export the model as an .lu file.**
- C. Create a new Speech service.
- D. Create a new Language Understanding service.

Answer: B

Explanation:

You might want to manage the translation and localization for the language understanding content for your bot independently.

Translate command in the @microsoft/bf-lu library takes advantage of the Microsoft text translation API to automatically machine translate .lu files to one or more than 60+ languages supported by the Microsoft text translation cognitive service.

What is translated?

- An .lu file and optionally translate Comments in the lu file
- LU reference link texts
- List of .lu files under a specific path.

Reference:

<https://github.com/microsoft/botframework-cli/blob/main/packages/luis/docs/translate-command.md>

Question: 60

DRAG DROP

You are using a Language Understanding service to handle natural language input from the users of a web-based customer agent.

The users report that the agent frequently responds with the following generic response: "Sorry, I don't understand that."

You need to improve the ability of the agent to respond to requests.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order. (Choose three.)

Actions	Answer Area
Add prebuilt domain models as required.	
Validate the utterances logged for review and modify the model.	
Migrate authoring to an Azure resource authoring key.	
. Enable active learning.	
Enable log collection by using Log Analytics.	
. Train and republish the Language Understanding model.	

Answer:

Explanation:

- enable active learning
- validate the utterances
- train and republish

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-how-to-review-endpoint-utterances#log-user-queries-to-enable-active-learning>

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-concept-prebuilt-model>

Question: 61

You build a conversational bot named bot1.

You need to configure the bot to use a QnA Maker application.

From the Azure Portal, where can you find the information required by bot1 to connect to the QnA Maker application?

- A. Access control (IAM)
- B. Properties
- C. Keys and Endpoint
- D. Identity

Answer: C

Explanation:

Obtain values to connect your bot to the knowledge base

1. In the QnA Maker site, select your knowledge base.
2. With your knowledge base open, select the SETTINGS tab. Record the value shown for service name. This value is useful for finding your knowledge base of interest when using the QnA Maker portal interface. It's not used to connect your bot app to this knowledge base.
3. Scroll down to find Deployment details and record the following values from the Postman sample HTTP request:
4. POST /knowledgebases/<knowledge-base-id>/generateAnswer
5. Host: <your-host-url>
6. Authorization: EndpointKey <your-endpoint-key>

Reference:

<https://docs.microsoft.com/en-us/azure/bot-service/bot-builder-howto-qna>

Question: 62

HOTSPOT

You are building a chatbot by using the Microsoft Bot Framework Composer.

You have the dialog design shown in the following exhibit.

AskForName > BeginDialog > Text Show code

Prompt for text
Text input
Collection information - Ask for a word or sentence.

[Learn more](#)

Bot Asks **User Input** **Other**

Property string
user.name

Output Format string
ex. =toUpperCase(this.value), \${toUpperCase(this.value)}

Value expression
fx =coalesce(@user.Name,@personName)

Expected responses (intent:
#TextInput_Response_GH5FTe)

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area

Statements

Yes No

user.name is an entity.

The dialog asks for a user name and a user age and assigns appropriate values to the user.name and user.age properties.

The chatbot attempts to take the first non-null entity value for userName or personName and assigns the value to user.name.

Answer:

Explanation:

Answer Area

Statements	Yes	No
------------	-----	----

user.name is an entity.

<input type="radio"/>	<input checked="" type="radio"/>
-----------------------	----------------------------------

The dialog asks for a user name and a user age and assigns appropriate values to the user.name and user.age properties.

<input checked="" type="radio"/>	<input type="radio"/>
----------------------------------	-----------------------

The chatbot attempts to take the first non-null entity value for userName or personName and assigns the value to user.name.

<input checked="" type="radio"/>	<input type="radio"/>
----------------------------------	-----------------------

Box 1: No

User.name is a property.

Box 2: Yes

Box 3: Yes

The coalesce() function evaluates a list of expressions and returns the first non-null (or non-empty for string) expression.

Reference:

<https://docs.microsoft.com/en-us/composer/concept-language-generation>

<https://docs.microsoft.com/en-us/azure/data-explorer/kusto/query/coalescefuction>

Question: 63

HOTSPOT

You are building a chatbot for a Microsoft Teams channel by using the Microsoft Bot Framework SDK. The chatbot will use the following code.

```
protected override async Task OnMembersAddedAsync(IList<ChannelAccount>
membersAdded, ITurnContext<IConversationUpdateActivity> turnContext,
CancellationToken cancellationToken)
{
    foreach (var member in membersAdded)
        if (member.Id != turnContext.Activity.Recipient.Id)
            await turnContext.SendActivityAsync($"Hi there - {member.Name}.
{WelcomeMessage}", cancellationToken: cancellationToken);
}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
OnMembersAddedAsync will be triggered when a user joins the conversation.	<input checked="" type="radio"/>	<input type="radio"/>
When a new user joins the conversation, the existing users in the conversation will see the chatbot greeting.	<input type="radio"/>	<input checked="" type="radio"/>
OnMembersAddedAsync will be initialized when a user sends a message.	<input type="radio"/>	<input checked="" type="radio"/>

Verified**Answer:**

Explanation:

Answer Area

Statements	Yes	No
OnMembersAddedAsync will be triggered when a user joins the conversation.	<input checked="" type="radio"/>	<input type="radio"/>
When a new user joins the conversation, the existing users in the conversation will see the chatbot greeting.	<input checked="" type="radio"/>	<input type="radio"/>
OnMembersAddedAsync will be initialized when a user sends a message.	<input type="radio"/>	<input checked="" type="radio"/>

Box 1: Yes

The ActivityHandler.OnMembersAddedAsync method overrides this in a derived class to provide logic for when members other than the bot join the conversation, such as your bot's welcome logic.

Box 2: Yes

membersAdded is a list of all the members added to the conversation, as described by the conversation update activity.

Box 3: No

Reference:

<https://docs.microsoft.com/en-us/dotnet/api/microsoft.bot.builder.activityhandler.onmembersaddedasync?view=botbuilder-dotnet-stable>

Question: 64

HOTSPOT

You are building a chatbot by using the Microsoft Bot Framework SDK.

You use an object named UserProfile to store user profile information and an object named ConversationData to store information related to a conversation.

You create the following state accessors to store both objects in state.

```
var userStateAccessors = _userState.CreateProperty<UserProfile>(nameof(UserProfile));
```

```
var conversationStateAccessors =  
_conversationState.CreateProperty<ConversationData>(nameof(ConversationData));
```

The state storage mechanism is set to Memory Storage.

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Statements	Yes	No
The code will create and maintain the <code>UserProfile</code> object in the underlying storage layer.	<input checked="" type="radio"/>	<input type="radio"/>
The code will create and maintain the <code>ConversationData</code> object in the underlying storage layer.	<input checked="" type="radio"/>	<input type="radio"/>
The <code>UserProfile</code> and <code>ConversationData</code> objects will persist when the Bot Framework runtime terminates.	<input type="radio"/>	<input checked="" type="radio"/>

Answer:

Explanation:

Statements	Yes	No
The code will create and maintain the <code>UserProfile</code> object in the underlying storage layer.	<input type="radio"/>	<input checked="" type="radio"/>
The code will create and maintain the <code>ConversationData</code> object in the underlying storage layer.	<input type="radio"/>	<input checked="" type="radio"/>
The <code>UserProfile</code> and <code>ConversationData</code> objects will persist when the Bot Framework runtime terminates.	<input checked="" type="radio"/>	<input type="radio"/>

Box 1: Yes

You create property accessors using the `CreateProperty` method that provides a handle to the `BotState` object. Each state property accessor allows you to get or set the value of the associated state property.

Box 2: Yes

Box 3: No

Before you exit the turn handler, you use the state management objects' `SaveChangesAsync()` method to write all state changes back to storage.

Reference:

<https://docs.microsoft.com/en-us/azure/bot-service/bot-builder-howto-v4-state>

Question: 65

HOTSPOT

You are building a chatbot that will provide information to users as shown in the following exhibit.

Passengers

Sarah Hum

Jeremy Goldberg

Evan Litvak

2 Stops**Tue, May 30, 2017 10:25 PM**

San Francisco

Amsterdam



San Francisco

Amsterdam

SFO

AMS

SFO

AMS

Non-Stop**Fri, Jun 2, 2017 11:55 PM**

San Francisco

Amsterdam



San Francisco

Amsterdam

SFO

AMS

SFO

AMS

Total

\$4,032.54

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

NOTE: Each correct selection is worth one point.

Answer Area

The chatbot is showing [answer choice].



an Adaptive Card
a Hero Card
a Thumbnail Card



The card includes [answer choice].



an action set
an image
an image group
media



Answer:

Explanation:

Box 1: Adaptive card

Box 2: an image

Reference:

<https://docs.microsoft.com/en-us/microsoftteams/platform/task-modules-and-cards/cards/cards-reference>

<https://docs.microsoft.com/en-us/composer/how-to-send-cards?tabs=v1x>

Question: 66

HOTSPOT

You are reviewing the design of a chatbot. The chatbot includes a language generation file that contains the following fragment.

```
# Greet(user)
- ${Greeting()}, ${user.name}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Statements	Yes	No
<code> \${user.name} </code> retrieves the user name by using a prompt.	<input type="radio"/>	<input checked="" type="radio"/>
Greet () is the name of the language generation template.	<input checked="" type="radio"/>	<input type="radio"/>
<code> \${Greeting() } </code> is a reference to a template in the language generation file.	<input checked="" type="radio"/>	<input type="radio"/>

Answer:

Explanation:

Statements	Yes	No
<code> \${user.name} </code> retrieves the user name by using a prompt.	<input type="radio"/>	<input type="radio"/>
Greet () is the name of the language generation template.	<input type="radio"/>	<input checked="" type="radio"/>
<code> \${Greeting() } </code> is a reference to a template in the language generation file.	<input type="radio"/>	<input checked="" type="radio"/>

Box 1: No

Example: Greet a user whose name is stored in `user.name`

- `${welcomeUser(user.name) }`

Verified this is correct

Example: Greet a user whose name you don't know:

- `${welcomeUser() }`

Box 2: No

Greet(User) is a Send a response action.

Box 3: Yes

Reference:

<https://docs.microsoft.com/en-us/composer/how-to-ask-for-user-input>

Question: 67

A customer uses Azure Cognitive Search.

The customer plans to enable a server-side encryption and use customer-managed keys (CMK) stored in Azure.

What are three implications of the planned change? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. The index size will increase.
- B. Query times will increase.
- C. A self-signed X.509 certificate is required.
- D. The index size will decrease.
- E. Query times will decrease.
- F. Azure Key Vault is required.

Answer: ABF

Explanation:

"Customer-managed keys require an additional billable service, Azure Key Vault, which can be in a different region, but under the same subscription, as Azure Cognitive Search. Enabling CMK encryption will increase index size and degrade query performance." same document also lists Azure Key Vault as a requirement:

<https://docs.microsoft.com/en-us/azure/search/search-security-overview#data-protection>

Question: 68

You are developing a new sales system that will process the video and text from a public-facing website.

You plan to notify users that their data has been processed by the sales system.

Which responsible AI principle does this help meet?

- A. transparency [verified-https://learn.microsoft.com/en-us/training/modules/preparing-for-ai-engineering/](https://learn.microsoft.com/en-us/training/modules/preparing-for-ai-engineering/)
- B. fairness
- C. inclusiveness
- D. reliability and safety

Answer: D

Explanation:

"When an AI application relies on personal data, such as a facial recognition system that takes images of people to recognize them; you should make it clear to the user how their data is used and retained, and who has access to it." from: <https://docs.microsoft.com/en-us/learn/paths/prepare-for-ai-engineering/>

Question: 69

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You create a web app named app1 that runs on an Azure virtual machine named vm1. Vm1 is on an Azure virtual network named vnet1.

You plan to create a new Azure Cognitive Search service named service1.

You need to ensure that app1 can connect directly to service1 without routing traffic over the public internet.

Solution: You deploy service1 and a public endpoint to a new virtual network, and you configure Azure Private Link.

Does this meet the goal?

A. Yes

B. No

Answer: B

Explanation:

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/cognitive-services-virtual-networks?tabs=portal#use-private-endpoints>

Question: 70

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You create a web app named app1 that runs on an Azure virtual machine named vm1. Vm1 is on an Azure virtual network named vnet1.

You plan to create a new Azure Cognitive Search service named service1.

You need to ensure that app1 can connect directly to service1 without routing traffic over the public internet.

Solution: You deploy service1 and a public endpoint, and you configure an IP firewall rule.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation:

Reference:

<https://docs.microsoft.com/en-us/azure/private-link/private-link-overview>

Question: 71

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You create a web app named app1 that runs on an Azure virtual machine named vm1. Vm1 is on an Azure virtual network named vnet1.

You plan to create a new Azure Cognitive Search service named service1.

You need to ensure that app1 can connect directly to service1 without routing traffic over the public internet.

Solution: You deploy service1 and a public endpoint, and you configure a network security group (NSG) for vnet1.

Does this meet the goal? **need private endpoint for the Cognitive Search service**

- A. Yes
- B. No

<https://learn.microsoft.com/en-us/azure/private-link/private-link-overview#network-security-groups>

Answer: A

Explanation:

<https://docs.microsoft.com/en-us/azure/virtual-network/network-security-groups-overview#network-security-groups>

<https://docs.microsoft.com/en-us/azure/virtual-network/virtual-networks-overview>

Question: 72

You plan to perform predictive maintenance.

You collect IoT sensor data from 100 industrial machines for a year. Each machine has 50 different sensors that generate data at one-minute intervals. In total, you have 5,000 time series datasets.

You need to identify unusual values in each time series to help predict machinery failures.

Which Azure Cognitive Services service should you use?

- A. Anomaly Detector
- B. Cognitive Search
- C. Form Recognizer
- D. Custom Vision

Answer: A

Explanation:

Question: 73

HOTSPOT

You are developing a streaming Speech to Text solution that will use the Speech SDK and MP3 encoding.

You need to develop a method to convert speech to text for streaming MP3 data.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

```
var audioFormat =  (AudioStreamContainerFormat.MP3);
 AudioConfig SetProperty
 AudioStreamFormat GetCompressedFormat
 AudioStreamFormat GetWaveFormatPCM
 PullAudioInputStream

var speechConfig = SpeechConfig.FromSubscription("18c51a87-3a69-47a8-aedc-a54745f708a1", "westus");

var audioConfig = AudioConfig.FromStreamInput(pushStream, audioFormat);

using (var recognizer = new  (speechConfig, audioConfig))
{
     KeywordRecognizer
     SpeakerRecognizer
     SpeechRecognizer
     SpeechSynthesizer

    var result = await recognizer.RecognizeOnceAsync();
    var text = result.Text;
}
```

Answer:

Explanation:

```
var audioFormat =  (AudioStreamContainerFormat.MP3);  
AudioConfig SetProperty  
AudioStreamFormat GetCompressedFormat  
AudioStreamFormat GetWaveFormatPCM  
PullAudioInputStream  
  
var speechConfig = SpeechConfig.FromSubscription("18c51a87-3a69-47a8-aedc-a54745f708a1", "westus");  
var audioConfig = AudioConfig.FromStreamInput(pushStream, audioFormat);  
using (var recognizer = new  (speechConfig, audioConfig))  
    KeywordRecognizer  
    SpeakerRecognizer  
    SpeechRecognizer  
    SpeechSynthesizer  
  
{  
    var result = await recognizer.RecognizeOnceAsync();  
    var text = result.Text;  
}  
}
```

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/speech-service/how-to-use-codec-compressed-audio-input-streams?tabs=debian&pivots=programming-language-csharp>

Question: 74

HOTSPOT

You are developing an internet-based training solution for remote learners.

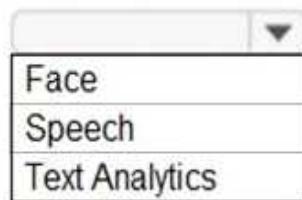
Your company identifies that during the training, some learners leave their desk for long periods or become distracted.

You need to use a video and audio feed from each learner's computer to detect whether the learner is present and paying attention. The solution must minimize development effort and identify each learner.

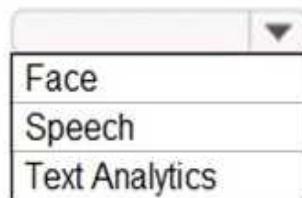
Which Azure Cognitive Services service should you use for each requirement? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

From a learner's video feed, verify whether the learner is present:



From a learner's facial expression in the video feed, verify whether the learner is paying attention:



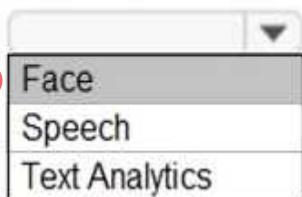
From a learner's audio feed, detect whether the learner is talking:



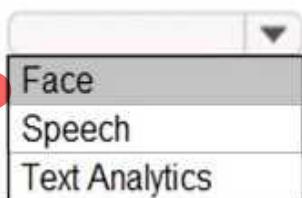
Answer:

Explanation:


From a learner's video feed, verify whether the learner is present:



From a learner's facial expression in the video feed, verify whether the learner is paying attention:



From a learner's audio feed, detect whether the learner is talking:



Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/what-are-cognitive-services>

Question: 75

You plan to provision a QnA Maker service in a new resource group named RG1.

In RG1, you create an App Service plan named AP1.

Which two Azure resources are automatically created in RG1 when you provision the QnA Maker service? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Language Understanding
- B. Azure SQL Database
- C. Azure Storage
- D. Azure Cognitive Search
- E. Azure App Service

Answer: DE

Explanation:

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/how-to/set-up-qnamaker-service-azure?tabs=v1#delete-azure-resources>

"When you create a QnAMaker resource, you host the data in your own Azure subscription. Azure Search is used to index your data." & "When you create a QnAMaker resource, you host the runtime in your own Azure subscription. App Service is the compute engine that runs the QnA Maker queries for you."

Question: 76

You are building a language model by using a Language Understanding service.

You create a new Language Understanding resource.

You need to add more contributors.

What should you use?

- A. a conditional access policy in Azure Active Directory (Azure AD)
- B. the Access control (IAM) page for the authoring resources in the Azure portal
- C. the Access control (IAM) page for the prediction resources in the Azure portal

Answer: B

Explanation:

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-how-to-collaborate>

Question: 77

HOTSPOT

You have a Computer Vision resource named contoso1 that is hosted in the West US Azure region.

You need to use contoso1 to make a different size of a product photo by using the smart cropping feature.

How should you complete the API URL? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Verified - it should be Supported Cognitive Services endpoints.

```
curl -H "Ocp-Apim-Subscription-Key: xxx" /  
-o "sample.png" -H "Content-Type: application/json" /  
"https://api.projectoxford.ai  
"https://contoso1.cognitiveservices.azure.com  
"https://westus.api.cognitive.microsoft.com" /vision/v3.1/  
areaOfInterest  
detect  
generateThumbnail  
?width=100&height=100&smartCropping=true" /  
-d "{\"url\": \"https://upload.litwareinc.org/litware/bicycle.jpg\"}"
```

so it is contoso1.cognitiveservices.azure.com

Answer:

Explanation:

westus.api.cognitive.microsoft.com

generateThumbnail

<https://docs.microsoft.com/en-us/rest/api/computervision/3.1/generate-thumbnail/generate-thumbnail#examples>

POST

https://westus.api.cognitive.microsoft.com/vision/v3.1/generateThumbnail?width=500&height=500
&smartCropping=True

Ocp-Apim-Subscription-Key: {API key}

Question: 78

DRAG DROP

You are developing a webpage that will use the Video Indexer service to display videos of internal company meetings.

You embed the Player widget and the Cognitive Insights widget into the page.

You need to configure the widgets to meet the following requirements:

Ensure that users can search for keywords.

Display the names and faces of people in the video.

Show captions in the video in English (United States).

How should you complete the URL for each widget? To answer, drag the appropriate values to the

correct targets. Each value may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Values	Answer Area
en-US	
false	
people,keywords	
people,search	
search	
true	

Cognitive Insights Widget
`https://www.videoindexer.ai/embed/insights/<accountId>/<videoId>/?widgets=` Value `controls=` Value

Player Widget
`https://www.videoindexer.ai/embed/player/<accountId>/<videoId>/? showcaptions=` Value `captions=` Value

Answer:

Explanation:

Cognitive Insights Widget

`https://www.videoindexer.ai/embed/insights/<accountId>/<videoId>/?widgets=` people,keywords `controls=` search

Player Widget

`https://www.videoindexer.ai/embed/player/<accountId>/<videoId>/? showcaptions=` true `captions=` en-US

Question: 79

DRAG DROP

You train a Custom Vision model to identify a company's products by using the Retail domain.

You plan to deploy the model as part of an app for Android phones.

You need to prepare the model for deployment.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions	Answer Area
Change the model domain.	
Retrain the model.	
Test the model.	
Export the model.	

Answer:

Explanation:

In user want to change to deploy offline model

1. Change model domain to compact model
2. Retrain compact model
3. Export model

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/export-your-model>

Question: 80

HOTSPOT

You are developing an application to recognize employees' faces by using the Face Recognition API. Images of the faces will be accessible from a URI endpoint.

The application has the following code.

```
static async void AddFace(string subscription_key, string personGroupId, string personId, string imageURI)
{
    var client = new HttpClient();
    client.DefaultRequestHeaders.Add("Ocp-Apim-Subscription-Key", subscription_key);
    var endpointURI = $"https://westus.api.cognitive.microsoft.com/face/v1.0/persongroups/{personGroupId}/persons/{personId}/persistedFaces";
    HttpResponseMessage response;
    var body = "{ \"url\": \"" + imageURI + "\"}";
    var content = new StringContent(body, Encoding.UTF8, "application/json");
    var response = await client.PutAsync(endpointURI, content);
}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Statements	Yes	No
The code will add a face image to a person object in a person group.	<input checked="" type="radio"/>	<input type="radio"/>
The code will work for a group of 10,000 people.	<input type="radio"/>	<input checked="" type="radio"/> true if s0 tier else free Tier F0 false
AddFace can be called multiple times to add multiple face images to a person object.	<input checked="" type="radio"/>	<input type="radio"/>

Answer:

Explanation:

- A. True
- B. True
- C. True

Verified all true

B: see this example code from documentation that uses PersonGroup of size 10,000 :

<https://docs.microsoft.com/en-us/azure/cognitive-services/face/face-api-how-to-topics/how-to-add-faces>

the question wants to trick you into thinking you need to use a LargePersonGroup for a size of 10,000 - but the documentation for it doesn't include this limitation or criteria:

<https://docs.microsoft.com/en-us/azure/cognitive-services/face/face-api-how-to-topics/how-to-use-large-scale>

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/face/face-api-how-to-topics/use-persondirectory>

Question: 81

DRAG DROP

You have a Custom Vision resource named acvdev in a development environment.

You have a Custom Vision resource named acvprod in a production environment.

In acvdev, you build an object detection model named obj1 in a project named proj1.

You need to move obj1 to acvprod.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions	Answer Area
Use the ExportProject endpoint on acvdev.	
Use the GetProjects endpoint on acvdev.	
Use the ImportProject endpoint on acvprod.	
Use the ExportIteration endpoint on acvdev.	
Use the GetIterations endpoint on acvdev.	
Use the UpdateProject endpoint on acvprod.	

Answer:

Explanation:

-  Use the GetProjects endpoint on acvdev.
-  Use the ExportProject endpoint on acvdev.
-  Use the ImportProject endpoint on acvprod.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-tutorial-pattern#what-did-this-tutorial-accomplish>

Question: 82

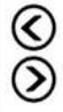
DRAG DROP

You train a Custom Vision model used in a mobile app.

You receive 1,000 new images that do not have any associated data.

You need to use the images to retrain the model. The solution must minimize how long it takes to retrain the model.

Which three actions should you perform in the Custom Vision portal? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions	Answer Area
Upload the images by category.	
Get suggested tags.	
Upload all the images.	
Group the images locally into category folders.	
Review the suggestions and confirm the tags.	
Tag the images manually.	

Answer:

Explanation:

- 1.) upload all the images
- 2.) Get suggested tags
- 3.) Review the suggestions and confirm the tags

reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/suggested-tags>

Question: 83

You are building a Language Understanding model for an e-commerce chatbot. Users can speak or type their billing address when prompted by the chatbot.

You need to construct an entity to capture billing addresses.

Which entity type should you use?

- A. machine learned **verified it is machine learning entity with structure**
- B. Regex
- C. list
- D. Pattern.any

Answer: B

Explanation:

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-concept-entity-types>

ML Entity with Structure

An ML entity can be composed of smaller sub-entities, each of which can have its own properties.

For example, Address could have the following structure:

Address: 4567 Main Street, NY, 98052, USA

Building Number: 4567
Street Name: Main Street
State: NY
Zip Code: 98052
Country: USA

Question: 84

You are building an Azure Webblob that will create knowledge bases from an array of URLs.

You instantiate a QnAMakerClient object that has the relevant API keys and assign the object to a variable named client.

You need to develop a method to create the knowledge bases.

Which two actions should you include in the method? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Create a list of FileDTO objects that represents data from the WebJob.
- B. Call the client. Knowledgebase. CreateAsync method.
- C. Create a list of QnADTO objects that represents data from the WebJob.
- D. Create a CreaceKbDTO object.

verified knowledgebase.createAsync,createkbDTO Answer: AC

Explanation:

Reference:

<https://docs.microsoft.com/en-us/rest/api/cognitiveservices-qnamaker/qnamaker4.0/knowledgebase/create>

Question: 85

HOTSPOT

You are developing an application that includes language translation.

The application will translate text retrieved by using a function named getTextToBeTranslated. The text can be in one of many languages. The content of the text must remain within the Americas Azure geography.

You need to develop code to translate the text to a single language.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

```
 . . .
var endpoint =
    "https://api.cognitive.microsofttranslator.com/translate";
    "https://api.cognitive.microsofttranslator.com/transliterate";
    "https://api-apc.cognitive.microsofttranslator.com/detect";
    "https://api-nam.cognitive.microsofttranslator.com/detect";
    "https://api-nam.cognitive.microsofttranslator.com/translate"; ;  
  
var apiKey = "FF956C68B83B21B38691ABD200A4C606";
var text = getTextToBeTranslated();
var body = '[{"Text":"' + text + '"}]';
var client = new HttpClient();
client.DefaultRequestHeaders.Add("Ocp-Apim-Subscription-Key", apiKey);  
  
var uri = endpoint + "?from=en";
var uri = endpoint + "?suggestedFrom=en";
var uri = endpoint + "?to=en";  
  
 HttpResponseMessage response;
var content = new StringContent(body, Encoding.UTF8, "application/json");
var response = await client.PutAsync(uri, content);
. . .
```

Answer:

Explanation:

```
 . . .
var endpoint =
    "https://api.cognitive.microsofttranslator.com/translate";
    "https://api.cognitive.microsofttranslator.com/transliterate";
    "https://api-apc.cognitive.microsofttranslator.com/detect";
    "https://api-nam.cognitive.microsofttranslator.com/detect";
    "https://api-nam.cognitive.microsofttranslator.com/translate"; ;  
  
var apiKey = "FF956C68B83B21B38691ABD200A4C606";
var text = getTextToBeTranslated();
var body = '[{"Text":"' + text + '"}]';
var client = new HttpClient();
client.DefaultRequestHeaders.Add("Ocp-Apim-Subscription-Key", apiKey);  
  
var uri = endpoint + "?from=en";
var uri = endpoint + "?suggestedFrom=en";
var uri = endpoint + "?to=en";  
  
 HttpResponseMessage response;
var content = new StringContent(body, Encoding.UTF8, "application/json");
var response = await client.PutAsync(uri, content);
. . .
```

Question: 86

You are building a natural language model.
You need to enable active learning.
What should you do?

- A. Add show-all-intents=true to the prediction endpoint query.
- B. Enable speech priming.
- C. Add log=true to the prediction endpoint query.**
- D. Enable sentiment analysis.

Answer: C

Explanation:

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-how-to-review-endpoint-utterances#log-user-queries-to-enable-active-learning>

Question: 87

You are developing a solution to generate a word cloud based on the reviews of a company's products.

Which Text Analytics REST API endpoint should you use?

- A. keyPhrases**
- B. sentiment
- C. languages
- D. entities/recognition/general

Answer: A

Explanation:

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/text-analytics/overview>

Question: 88

DRAG DROP

You have a web app that uses Azure Cognitive Search.

When reviewing billing for the app, you discover much higher than expected charges. You suspect that the query key is compromised.

You need to prevent unauthorized access to the search endpoint and ensure that users only have read only access to the documents collection. The solution must minimize app downtime.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions	Answer Area
Add a new query key.	(Red circle)
Regenerate the secondary admin key.	(Left arrow)
Change the app to use the secondary admin key.	(Up arrow)
Change the app to use the new key.	(Right arrow)
Regenerate the primary admin key.	(Down arrow)
Delete the compromised key.	(Red circle)

Answer:

Explanation:

Add a new query key.

Change the app to use the new key.

Delete the compromised key.

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-security-api-keys>

Question: 89

HOTSPOT

You are building a bot and that will use Language Understanding.

You have a LUDown file that contains the following content.

```
## Confirm
- confirm
- ok
- yes

## ExtractName
- call me steve !
- i am anna
- (i'm|i am) {@PersonName.Any}[]
- my name is {@PersonName.Any}[]

## Logout
- forget me
- log out

## SelectItem
- choose last
- choose the {@DirectionalReference=bottom left}
- choose {@DirectionalReference=top right}
- i like {@DirectionalReference=left} one

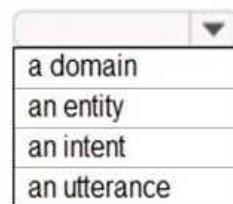
## SelectNone
- none

@ ml DirectionalReference
@ prebuilt personName
```

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

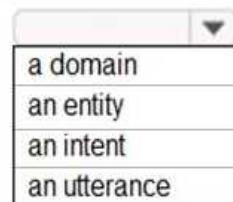
NOTE: Each correct selection is worth one point.

SelectItem is [answer choice].



a domain
an entity
an intent
an utterance

Choose {@DirectionalReference=top right} is [answer choice].



a domain
an entity
an intent
an utterance

Answer:

Explanation:

SelectItem is [answer choice].

a domain
an entity
an intent
an utterance

Choose {@DirectionalReference=top right} is [answer choice].

a domain
an entity
an intent
an utterance

Reference:

<https://github.com/solliancenet/tech-immersion-data-ai/blob/master/ai-exp1/README.md>

Question: 90

HOTSPOT

You are designing a conversation flow to be used in a chatbot.

You need to test the conversation flow by using the Microsoft Bot Framework Emulator.

How should you complete the .chat file? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

```
user=User1  
bot=watchbot  
user: I want a new watch.
```

```
bot: [ ] [Delay=3000]  
Attachment  
ConversationUpdate  
Typing
```

```
bot: I can help you with that! Let me see what I can find.
```

```
bot: Here's what I found.
```

```
bot:
```

```
[AttachmentLayout=  
adaptivecard  
carousel  
thumbnail]
```

```
[Attachment=https://contoso.blob.core.windows.net/watch01.jpg]
```

```
[Attachment=https://contoso.blob.core.windows.net/watch02.jpg]
```

```
user: I like the first one.
```

```
bot: Sure, pulling up more information.
```

```
bot: [Attachment=cards\watchProfileCard.json
```

```
user: That's nice! Thank you.
```

```
bot: Sure, you are most welcome!
```

```
adaptivecard  
carousel  
list
```

Answer:

Explanation:

```
user=User1  
bot=watchbot  
user: I want a new watch.
```

```
bot: [ ] [Delay=3000]
```

- Attachment
- ConversationUpdate
- Typing

```
bot: I can help you with that! Let me see what I can find.
```

```
bot: Here's what I found.
```

```
bot:
```

```
[AttachmentLayout=
```

- adaptivecard
- carousel
- thumbnail

```
[Attachment=https://contoso.blob.core.windows.net/watch01.jpg]
```

```
[Attachment=https://contoso.blob.core.windows.net/watch02.jpg]
```

```
user: I like the first one.
```

```
bot: Sure, pulling up more information.
```

```
bot: [Attachment=cards\watchProfileCard.json
```

```
user: That's nice! Thank you.
```

```
bot: Sure, you are most welcome!
```

```
[ ]
```

- adaptivecard
- carousel
- list

Reference:

<https://docs.microsoft.com/en-us/azure/bot-service/bot-builder-howto-add-media-attachments?view=azure-bot-service-4.0&tabs=csharp>

Question: 91

You build a bot by using the Microsoft Bot Framework SDK and the Azure Bot Service.

You plan to deploy the bot to Azure.

You register the bot by using the Bot Channels Registration service.

Which two values are required to complete the deployment? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. botId
- B. tenantId
- C. appId
- D. objectId
- E. appSecret

Answer: CE

Explanation:

Reference:

<https://github.com/MicrosoftDocs/bot-docs/blob/live/articles/bot-service-quickstart-registration.md>

Question: 92

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an Azure Cognitive Search service.

During the past 12 months, query volume steadily increased.

You discover that some search query requests to the Cognitive Search service are being throttled.

You need to reduce the likelihood that search query requests are throttled.

Solution: You migrate to a Cognitive Search service that uses a higher tier.

Does this meet the goal?

- A. Yes
- B. No

Answer: A

Explanation:

A simple fix to most throttling issues is to throw more resources at the search service (typically replicas for query-based throttling, or partitions for indexing-based throttling). However, increasing replicas or partitions adds cost, which is why it is important to know the reason why throttling is occurring at all.

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-performance-analysis>

Question: 93

DRAG DROP

You need to develop an automated call handling system that can respond to callers in their own

language. The system will support only French and English.

Which Azure Cognitive Services service should you use to meet each requirement? To answer, drag the appropriate services to the correct requirements. Each service may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Services

Answer Area

Speaker Recognition

Speech to Text

Text Analytics

Text to Speech

Translator

Detect the incoming language:

Respond in the callers' own language:

Answer:

Explanation:

Detect the incoming language:

verified Text Analytics

Respond in the callers' own language:

Translator

Box 1: Text Analytics

The Language Detection feature of the Azure Text Analytics REST API evaluates text input for each document and returns language identifiers with a score that indicates the strength of the analysis.

Box 2: Translator

Translator is a cloud-based neural machine translation service that is part of the Azure Cognitive Services family of REST APIs. Translator can be used with any operating system and powers many Microsoft products and services used by thousands of businesses worldwide to perform language translation and other language-related operations.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/text-analytics/how-tos/text-analytics-how-to-language-detection>

<https://docs.microsoft.com/en-us/azure/cognitive-services/translator/translator-overview>

Question: 94

You have receipts that are accessible from a URL.

You need to extract data from the receipts by using Form Recognizer and the SDK. The solution must use a prebuilt model.

Which client and method should you use?

- A. the FormRecognizerClient client and the StartRecognizeContentFromUri method
- B. the FormTrainingClient client and the StartRecognizeContentFromUri method
- C. the FormRecognizerClient client and the StartRecognizeReceiptsFromUri method
- D. the FormTrainingClient client and the StartRecognizeReceiptsFromUri method

Answer: D

Explanation:

To analyze receipts from a URL, use the StartRecognizeReceiptsFromUri method

Example code:

```
private static async Task AnalyzeReceipt(  
    FormRecognizerClient recognizerClient, string receiptUri)  
{  
    RecognizedFormCollection receipts = await recognizerClient.StartRecognizeReceiptsFromUri(new  
        Uri(receiptUri)).WaitForCompletionAsync();
```

Reference:

<https://docs.microsoft.com/en-us/azure/applied-ai-services/form-recognizer/quickstarts/client-library>

Question: 95

You have a collection of 50,000 scanned documents that contain text.

You plan to make the text available through Azure Cognitive Search.

You need to configure an enrichment pipeline to perform optical character recognition (OCR) and text analytics. The solution must minimize costs.

What should you attach to the skillset?

- A. a new Computer Vision resource
- B. a free (Limited enrichments) Cognitive Services resource
- C. an Azure Machine Learning pipeline

- D. a new Cognitive Services resource that uses the SO pricing tier

Answer: A

Explanation:

The Computer Vision API uses text recognition APIs to extract and recognize text information from images. Read uses the latest recognition models, and is optimized for large, text-heavy documents and noisy images.

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/solution-ideas/articles/cognitive-search-with-skillsets>

Question: 96

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an Azure Cognitive Search service.

During the past 12 months, query volume steadily increased.

You discover that some search query requests to the Cognitive Search service are being throttled.

You need to reduce the likelihood that search query requests are throttled.

Solution: You add indexes.

Does this meet the goal?

A. Yes

B. No

Answer: B

Explanation:

Instead, you could migrate to a Cognitive Search service that uses a higher tier.

Note: A simple fix to most throttling issues is to throw more resources at the search service (typically replicas for query-based throttling, or partitions for indexing-based throttling). However, increasing replicas or partitions adds cost, which is why it is important to know the reason why throttling is occurring at all.

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-performance-analysis>

Question: 97

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an Azure Cognitive Search service.

During the past 12 months, query volume steadily increased.

You discover that some search query requests to the Cognitive Search service are being throttled.

You need to reduce the likelihood that search query requests are throttled.

Solution: You enable customer-managed key (CMK) encryption.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation:

Customer-managed key (CMK) encryption does not affect throttling.

Instead, you could migrate to a Cognitive Search service that uses a higher tier.

Note: A simple fix to most throttling issues is to throw more resources at the search service (typically replicas for query-based throttling, or partitions for indexing-based throttling). However, increasing replicas or partitions adds cost, which is why it is important to know the reason why throttling is occurring at all.

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-performance-analysis>

Question: 98

DRAG DROP

You are developing an application that will recognize faults in components produced on a factory production line. The components are specific to your business.

You need to use the Custom Vision API to help detect common faults.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions**Answer Area**

Train the classifier model.

Upload and tag images.

Initialize the training dataset.

Train the object detection model.

Create a project.



Answer:

Explanation:

Create a project.

Upload and tag images.

Train the classifier model.

Step 1: Create a project

Create a new project.

Step 2: Upload and tag the images

Choose training images. Then upload and tag the images.

Step 3: Train the classifier model.

Train the classifier

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/custom-vision-service/getting-started-build-a-classifier>

Question: 99

HOTSPOT

You are building a model that will be used in an iOS app.

You have images of cats and dogs. Each image contains either a cat or a dog.

You need to use the Custom Vision service to detect whether the images is of a cat or a dog.

How should you configure the project in the Custom Vision portal? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Project Types:

Classification
Object Detection

Classification Types:

Multiclass (Single tag per image)
Multilabel (Multiple tags per image)

Domains:

Audit
Food
General
General (compact)
Landmarks
Landmarks (compact)
Retail
Retail (compact)

Answer:

Explanation:

Project Types:

Classification

Object Detection

Classification Types:

Multiclass (Single tag per image)

Multilabel (Multiple tags per image)

Domains:

Audit

Food

General

General (compact)

Landmarks

Landmarks (compact)

Retail

Retail (compact)

Box 1: Classification

multiclass- image assigned one tag only
multilabel - image can assigned multiple tags

Box 2: Multiclass

A multiclass classification project is for classifying images into a set of tags, or target labels. An image can be assigned to one tag only.

Box 3: General

General: Optimized for a broad range of image classification tasks. If none of the other specific domains are appropriate, or if you're unsure of which domain to choose, select one of the General domains.

Reference:

<https://cran.r-project.org/web/packages/AzureVision/vignettes/customvision.html>

Question: 100

HOTSPOT

You run the following command.

```
docker run --rm -it -p 5000:5000 --memory 10g --cpus 2 \
mcr.microsoft.com/azure-cognitive-services/textanalytics/sentiment\
Eula=accept \
Billing={ENDPOINT_URI} \
ApiKey={API_KEY}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

Statements	Yes	No
Going to http://localhost:5000/status will query the Azure endpoint to verify whether the API key used to start the container is valid.	<input type="radio"/>	<input type="radio"/>
The container logging provider will write log data.	<input type="radio"/>	<input checked="" type="radio"/>
Going to http://localhost:5000/swagger will provide the details to access the documentation for the available endpoints.	<input type="radio"/>	<input type="radio"/>

Answer:

Explanation:

Statements	Yes	No
Going to http://localhost:5000/status will query the Azure endpoint to verify whether the API key used to start the container is valid.	<input type="radio"/>	<input type="radio"/>
The container logging provider will write log data.	<input type="radio"/>	<input type="radio"/>
Going to http://localhost:5000/swagger will provide the details to access the documentation for the available endpoints.	<input type="radio"/>	<input type="radio"/>

Box 1: Yes

<http://localhost:5000/status> : Also requested with GET, this verifies if the api-key used to start the container is valid without causing an endpoint query.

Box 2: Yes

The command saves container and LUIS logs to output mount at C:\output, located on container host
verified The container includes a Fluentd logging provider, which allows your container to write logs

Box 3: Yes

<http://localhost:5000/swagger> : The container provides a full set of documentation for the endpoints and a Try it out feature. With this feature, you can enter your settings into a web-based HTML form and make the query without having to write any code. After the query returns, an example CURL command is provided to demonstrate the HTTP headers and body format that's required.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-container-howto>

Question: 101

You are training a Language Understanding model for a user support system.

You create the first intent named GetContactDetails and add 200 examples.

You need to decrease the likelihood of a false positive.

What should you do?

- A. Enable active learning.
- B. Add a machine learned entity.
- C. Add additional examples to the GetContactDetails intent.
- D. Add examples to the None intent.

Answer: A

Explanation:

Active learning is a technique of machine learning in which the machine learned model is used to identify informative new examples to label. In LUIS, active learning refers to adding utterances from the endpoint traffic whose current predictions are unclear to improve your model.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-glossary>

Question: 102

DRAG DROP

You are building a Language Understanding model for purchasing tickets.

You have the following utterance for an intent named PurchaseAndSendTickets.

Purchase [2 audit business] tickets to [Paris] [next Monday] and send tickets to [email@domain.com]

You need to select the entity types. The solution must use built-in entity types to minimize training data whenever possible.

Which entity type should you use for each label? To answer, drag the appropriate entity types to the correct labels. Each entity type may be used once, more than once, or not at all.

You may need to drag the split bar between panes or scroll to view content.

Entity Types	Answer Area
Email	
List	Paris:
Regex	email@domain.com:
GeographyV2	2 audit business:
Machine learned	

Answer:

Explanation:

Paris:	GeographyV2
email@domain.com:	Email
2 audit business:	Machine learned

Box 1: GeographyV2

The prebuilt geographyV2 entity detects places. Because this entity is already trained, you do not need to add example utterances containing GeographyV2 to the application intents.

Box 2: Email

Email prebuilt entity for a LUIS app: Email extraction includes the entire email address from an utterance. Because this entity is already trained, you do not need to add example utterances containing email to the application intents.

Box 3: Machine learned

The machine-learning entity is the preferred entity for building LUIS applications.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-reference-prebuilt-geographyv2>

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/luis-reference-prebuilt-email>

<https://docs.microsoft.com/en-us/azure/cognitive-services/luis/reference-entity-machine-learned-entity>

Question: 103

You are developing an application that will use Azure Cognitive Search for internal documents.

You need to implement document-level filtering for Azure Cognitive Search.

Which three actions should you include in the solution? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Send Azure AD access tokens with the search request.
- B. Retrieve all the groups.
- C. Retrieve the group memberships of the user.
- D. Add allowed groups to each index entry.
- E. Create one index per group.
- F. Supply the groups as a filter for the search requests.

Answer: CDF

Explanation:

Your documents must include a field specifying which groups have access. This information becomes the filter criteria against which documents are selected or rejected from the result set returned to the issuer.

D: A query request targets the documents collection of a single index on a search service.

CF: In order to trim documents based on group_ids access, you should issue a search query with a group_ids/any(g:search.in(g, 'group_id1, group_id2,...')) filter, where 'group_id1, group_id2,...' are the groups to which the search request issuer belongs.

Reference:

<https://docs.microsoft.com/en-us/azure/search/search-security-trimming-for-azure-search>

Question: 104

You are building a chatbot by using the Microsoft Bot Framework Composer as shown in the exhibit.
(Click the Exhibit tab.)

The screenshot shows a dialog flow in the Microsoft Bot Framework designer. A 'BeginDialog' event triggers a 'Bot Asks (Text)' step, which asks 'What is your name?'. This leads to a 'User input (Text)' step where the expression '(SCOPE).name = Input(Text)' is set. The 'User input' step has a dashed arrow pointing back to the 'BeginDialog' event, indicating it's a disposable property.

Prompt for text

Text Input
Collection information - Ask for a word or sentence.

Learn more

Bot response User input Other

Property: string

(SCOPE).name

Output format: string

Value: string

Expected responses (intent: #TextInput_Response_FuuyF4)

The chatbot contains a dialog named Get UserDetails. Get UserDetails contains a TextInput control that prompts users for their name.

The user input will be stored in a property named name.

You need to ensure that you can dispose of the property when the last active dialog ends.

Which scope should you assign to name?

- A. dialog
- B. user
- C. curr
- D. conversation

Answer: A

Explanation:

The dialog scope associates properties with the active dialog. Properties in the dialog scope are retained until the dialog ends.

Incorrect Answers:

A: The conversation scope associates properties with the current conversation. Properties in the conversation scope have a lifetime of the conversation itself. These properties are in scope while the bot is processing an activity associated with the conversation (for example, multiple users together in a Microsoft Teams channel).

B: The user scope associates properties with the current user. Properties in the user scope do not expire. These properties are in scope while the bot is processing an activity associated with the user.

C: The turn scope associates properties with the current turn. Properties in the turn scope expire at the end of the turn.

Reference:

<https://docs.microsoft.com/en-us/composer/concept-memory?tabs=v2x>

Question: 105

DRAG DROP

You have a chatbot that uses a QnA Maker application.

You enable active learning for the knowledge base used by the QnA Maker application.

You need to integrate user input into the model.

Which four actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions

Answer Area

Add a task to the Azure resource.

Approve and reject suggestions.

Publish the knowledge base.

Modify the automation task logic app to run an Azure Resource Manager template that creates the Azure Cognitive Services resource.

For the knowledge base, select Show active learning suggestions.

Save and train the knowledge base.

Select the properties of the Azure Cognitive Services resource.

2

4

1

3



Answer:

Explanation:

For the knowledge base, select Show active learning suggestions.

Approve and reject suggestions.

Save and train the knowledge base.

Publish the knowledge base.

Step 1: For the knowledge base, select Show active learning suggestions.

In order to see the suggested questions, on the Edit knowledge base page, select View Options, then select Show active learning suggestions.

Step 2: Approve and reject suggestions.

Each QnA pair suggests the new question alternatives with a check mark, , to accept the question or an x to reject the suggestions. Select the check mark to add the question.

Step 3: Save and train the knowledge base.

Select Save and Train to save the changes to the knowledge base.

Step 4: Publish the knowledge base.

Select Publish to allow the changes to be available from the GenerateAnswer API.

When 5 or more similar queries are clustered, every 30 minutes, QnA Maker suggests the alternate questions for you to accept or reject.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/qnamaker/how-to/improve-knowledge-base>

Question: 106

You need to enable speech capabilities for a chatbot.

Which three actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Enable WebSockets for the chatbot app.
- B. Create a Speech service.
- C. Register a Direct Line Speech channel.
- D. Register a Cortana channel.
- E. Enable CORS for the chatbot app.
- F. Create a Language Understanding service.

Answer: ABC

Explanation:

You can use the Speech service to voice-enable a chat bot.

The Direct Line Speech channel uses the text-to-speech service, which has neural and standard voices.

You'll need to make a small configuration change so that your bot can communicate with the Direct Line Speech channel using web sockets.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/speech-service/tutorial-voice-enable-your-bot-speech-sdk>

Question: 107

You are examining the Text Analytics output of an application.

The text analyzed is: "Our tour guide took us up the Space Needle during our trip to Seattle last week."

The response contains the data shown in the following table.

Text	Category	ConfidenceScore
Tour guide	PersonType	0.45
Space Needle	Location	0.38
Trip	Event	0.78
Seattle	Location	0.78
Last week	DateTime	0.80

Which Text Analytics API is used to analyze the text?

- A. Sentiment Analysis
- B. Named Entity Recognition
- C. Entity Linking
- D. Key Phrase Extraction

Answer: B

Explanation:

Question: 108**HOTSPOT**

You are developing an application that includes language translation.

The application will translate text retrieved by using a function named `get_text_to_be_translated`.

The text can be in one of many languages. The content of the text must remain within the Americas Azure geography.

You need to develop code to translate the text to a single language.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
    ...
    api_key = "FF956C6BB83B21B38691ABD200A4C606"
    text = get_text_to_be_translated()
    headers = {
        'Content-Type': 'application/json',
        'Ocp-Apim-Subscription-Key': api_key
    }
    body = {
        'Text': text
    }
    conn = httplib.HTTPSConnection
    conn.request("POST",
    response = conn.getresponse()
    conn.request("POST",
    response = conn.getresponse()
    response_data = response.r
    ...

```

<http://api-nam.cognitive.microsofttranslator.com/detect>

**Asia Pacific:**

api-apc.cognitive.microsofttranslator.com

Europe (except Switzerland):

api-eur.cognitive.microsofttranslator.com

confirmed from link

Explanation:

Question: 109**HOTSPOT**

You create a knowledge store for Azure Cognitive Search by using the following JSON.

```
"knowledgeStore": {
    "storageConnectionString": "DefaultEndpointsProtocol=https;AccountName=<Acct Name>;AccountKey=<Acct Key>;",
    "projections": [
        {
            "tables": [
                {
                    "tableName": "unrelatedDocument",
                    "generatedKeyName": "Documentid",
                    "source": "/document/pbiShape"
                },
                {
                    "tableName": "unrelatedKeyPhrases",
                    "generatedKeyName": "KeyPhraseid",
                    "source": "/document/pbiShape/keyPhrases"
                }
            ],
            "objects": [
                {
                    "storageContainer": "unrelatedocrlayout",
                    "source": null,
                    "sourceContext": "/document/normalized_images/*/layoutText",
                    "inputs": [
                        {
                            "name": "ocrLayoutText",
                            "source": "/document/normalized_images/*/layoutText"
                        }
                    ]
                }
            ],
            "files": []
        }
    ]
}
```

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic. NOTE Each correct selection is worth one point.

Answer:

Answer Area

1. Azure Blob Storage
2. Azure cosmos Db
3. Azure Data Lake Storage
4. Azure sql Database
5. Azure Table stage

2 projection groups

Images will [answer choice]

There will be [answer choice].

no projection groups
one projection group
two projection groups
four projection groups

not be saved.
be saved to a blob container.
be saved to file storage.
be saved to an Azure Data lake.

Answer:

be projected to azure Blob storage

Explanation:

Answer Area

6. Azure Sql Managed Instance
7. Sql Server on Azure VM
8. Azure files,Mysql
- 9.sharepoint, Azure cosmos for (mongodb or Apache Gemini),10. oneLake

Question: 110

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You create a web app named app1 that runs on an Azure virtual machine named vm1. Vm1 is on an Azure virtual network named vnet1.

You plan to create a new Azure Cognitive Search service named service1.

You need to ensure that app1 can connect directly to service1 without routing traffic over the public internet.

Solution: You deploy service1 and a private endpoint to vnet1.

Does this meet the goal?

- A. Yes
- B. No

Answer: A

Explanation:

A private endpoint is a network interface that uses a private IP address from your virtual network. This network interface connects you privately and securely to a service powered by Azure Private Link. By enabling a private endpoint, you're bringing the service into your virtual network.

The service could be an Azure service such as:

- Azure Storage
- Azure Cosmos DB
- Azure SQL Database
- Your own service using a Private Link Service.

Reference:

<https://docs.microsoft.com/en-us/azure/private-link/private-endpoint-overview>

Question: 111

You have a Language Understanding resource named lu1.

You build and deploy an Azure bot named bot1 that uses lu1.

You need to ensure that bot1 adheres to the Microsoft responsible AI principle of inclusiveness. How should you extend bot1?

- A. Implement authentication for bot1.
- B. Enable active learning for lu1.
- C. Host lu1 in a container.
- D. Add Direct Line Speech to bot1.

Answer: D

Explanation:

Inclusiveness: AI systems should empower everyone and engage people.

Direct Line Speech is a robust, end-to-end solution for creating a flexible, extensible voice assistant. It is powered by the Bot Framework and its Direct Line Speech channel, that is optimized for voice-in, voice-out interaction with bots.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/speech-service/direct-line-speech>

Question: 112

Your company uses an Azure Cognitive Services solution to detect faces in uploaded images. The method to detect the faces uses the following code.

```
static async Task DetectFaces(string imagePath)
{
    HttpClient client = new HttpClient();
    DefaultRequestHeaders.Add("Ocp-Apim-Subscription-Key", subscriptionKey);
    string requestParameters = "detectionModel=detection_01&returnFaceId=true&returnFaceLandmarks=false";
    string uri = endpoint + "/face/v1.0/detect?" + requestParameters;
    HttpResponseMessage response;
    byte[] byteData = GetImagesAsByteArray(imagePath);
    using (ByteArrayContent content = new ByteArrayContent(byteData))
    {
        Headers.ContentType = new MediaTypeHeaderValue("application/octet-stream");
        response = await PostAsync(uri, content);
        string contentString = await Content.ReadAsStringAsync();
        ProcessDetection(contentString);
    }
}
```

You discover that the solution frequently fails to detect faces in blurred images and in images that contain sideways faces.

- A. Use a different version of the Face API.
- B. Use the Computer Vision service instead of the Face service.
- C. Use the Identify method instead of the Detect method.
- D. Change the detection model.

You need to increase the likelihood that the solution can detect faces in blurred images and images that contain sideways faces.

What should you do?

Answer: D

Explanation:

Evaluate different models.

The best way to compare the performances of the detection models is to use them on a sample dataset. We recommend calling the Face - Detect API on a variety of images, especially images of many faces or of faces that are difficult to see, using each detection model. Pay attention to the number of faces that each model returns.

Reference:

<https://docs.microsoft.com/en-us/azure/cognitive-services/face/face-api-how-to-topics/specify-detection-model>

Question: 113

You have the following C# method.

```
static void create_resource(string resource_name, string kind, string account_tier, string location)
{
    CognitiveServicesAccount parameters =
        new CognitiveServicesAccount(null, null, kind, location, resource_name, new CognitiveServicesAccountProperties(), new Sku(account_tier));
    var result = cog_svc_client.Accounts.Create(resource_group_name, account_tier, parameters);
}
```

You need to deploy an Azure resource to the East US Azure region. The resource will be used to perform sentiment analysis.

How should you call the method?

- A. create_resource("res1", "ContentModerator", "S0", "eastus")
- B. create_resource("res1", "TextAnalytics", "S0", "eastus")
- C. create_resource("res1", "ContentModerator", "Standard", "East US")
- D. create_resource("res1", "TextAnalytics", "Standard", "East US")

Answer: B

Explanation:

To perform sentiment analysis, we specify TextAnalytics, not ContentModerator.

Possible SKU names include: 'F0','F1','S0','S1','S2','S3','S4','S5','S6','S7','S8'

Possible location names include: westus, eastus

Reference:

<https://docs.microsoft.com/en-us/powershell/module/az.cognitiveservices/new-azcognitiveservicesaccount>

Question: 114

You build a Language Understanding model by using the Language Understanding portal.

You export the model as a JSON file as shown in the following sample.

```
{  
    "text": "average amount of rain by month at chicago last year",  
    "intent": "Weather.CheckWeatherValue",  
    "entities": [  
        {  
            "entity": "Weather.WeatherRange",  
            "startPos": 0,  
            "endPos": 6,  
            "children": []  
        },  
        {  
            "entity": "Weather.WeatherCondition",  
            "startPos": 18,  
            "endPos": 21,  
            "children": []  
        },  
        {  
            "entity": "Weather.Historic",  
            "startPos": 23,  
            "endPos": 30,  
            "children": []  
        }  
    ]  
}
```

To what does the Weather.Historic entity correspond in the utterance?

- A. by month
- B. chicago
- C. rain
- D. location

Answer: A

Explanation:

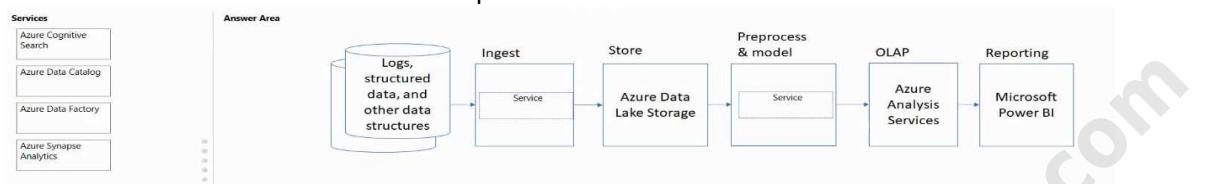
Question: 115

DRAG DROP

Match the Azure services to the appropriate locations in the architecture.

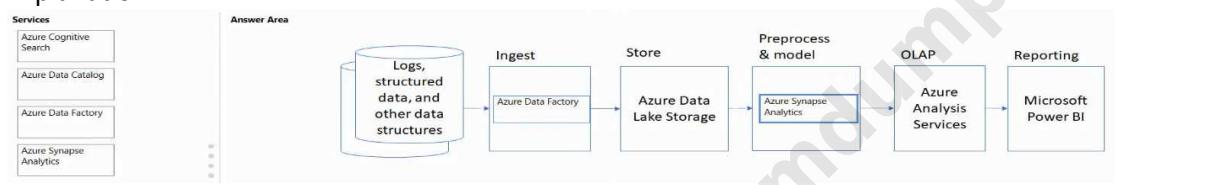
To answer, drag the appropriate service from the column on the left to its location on the right. Each service may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.



Answer:

Explanation:



Question: 116

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Statements		Yes	No
Platform as a service (PaaS) database offerings in Azure require less setup and configuration effort than infrastructure as a service (IaaS) database offerings.	<input checked="" type="radio"/>	<input type="radio"/>	
Platform as a service (PaaS) database offerings in Azure provide end users with the ability to control and update the operating system version.	<input type="radio"/>	<input checked="" type="radio"/>	
All relational and non-relational platform as a service (PaaS) database offerings in Azure can be paused to reduce costs.	<input type="radio"/>	<input checked="" type="radio"/>	

Not all PAAS offerings support pausing

Answer:

Explanation:

Statements		Yes	No
Platform as a service (PaaS) database offerings in Azure require less setup and configuration effort than infrastructure as a service (IaaS) database offerings.	<input checked="" type="radio"/>	<input type="radio"/>	
Platform as a service (PaaS) database offerings in Azure provide end users with the ability to control and update the operating system version.	<input type="radio"/>	<input checked="" type="radio"/>	
All relational and non-relational platform as a service (PaaS) database offerings in Azure can be paused to reduce costs.	<input type="radio"/>	<input checked="" type="radio"/>	

Question: 117

Which statement is an example of Data Manipulation Language (DML)?

- A. Revoke
- B. UPDATE
- C. DROP
- D. CREATE

Answer: B

Explanation:

Question: 118

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
Normalization involves eliminating relationships between database tables.	<input type="radio"/>	<input checked="" type="radio"/>
Normalizing a database reduces data redundancy.	<input checked="" type="radio"/>	<input type="radio"/>
Normalization improves data integrity.	<input type="radio"/>	<input checked="" type="radio"/>

Explanation:

Answer Area

Statements	Yes	No
Normalization involves eliminating relationships between database tables.	<input type="radio"/>	<input checked="" type="radio"/>
Normalizing a database reduces data redundancy.	<input checked="" type="radio"/>	<input type="radio"/>
Normalization improves data integrity.	<input checked="" type="radio"/>	<input type="radio"/>

Question: 119

HOTSPOT

Select the answer that correctly completes the sentence.

Answer Area

A block of code that runs in a database is called

a stored procedure.
 a table.
 a view.
 an index.

Explanation:

Answer Area

A block of code that runs in a database is called

Answer selection to complete the sentence A block of code that runs in a database is called

Question: 120

HOTSPOT

Select the answer that correctly completes the sentence.

Answer Area

The

clause can be used in Data Manipulation Language (DML) statements to specify the criteria that rows must match.

Explanation:

Answer Area

The clause can be used in Data Manipulation Language (DML) statements to specify the criteria that rows must match.

Answer:

Question: 121

Your company needs to implement a relational database in Azure. The solution must minimize ongoing maintenance. Which Azure service should you use?

- A. SQL Server on Azure Virtual Machines
- B. Azure SQL Database
- C. Azure HDInsight
- D. Azure Cosmos DB

Answer: B

Explanation:

Question: 122

You have a SQL query that combines customer data and order data.

- a. The query includes calculated columns. You need to create a database object that would allow other users to rerun the same SQL query. What should you create?
- A. an Index
- B. a view
- C. a scalar function
- D. a table

Answer: B

Explanation:

Question: 123

What are two benefits of platform as a service (PaaS) relational database offerings in Azure, such as Azure SQL Database? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. reduced administrative effort for managing the server infrastructure
- B. complete control over backup and restore processes
- C. in-database machine learning services S3
- D. access to the latest features

Answer: A, D

Explanation:

Question: 124

You have data saved in the following format.

```
FirstName,LastName,Age,LeisureHobby,SportsHobby  
John,Smith,23,Reading,Basketball  
Ben,Smith,21,Guitar,Curling
```

Which format was used?

- A. CSV
- B. JSON
- C. HTML
- D. YAML

Answer: A

Explanation:

Question: 125

What is a primary characteristic of a relational database?

- A. data is queried and manipulated by using a variant of the SQL language
- B. a lack of dependencies between tables
- C. a flexible data structure
- D. a large amount of duplicate data

Answer: C

Explanation:

Question: 126

HOTSPOT

Select the answer that correctly completes the sentence.

Answer Area

 A data analyst
 A data engineer
 A data scientist

is responsible for identifying which business rules must be applied to the data of a company.

data Analyst identify business rules

Answer:

Explanation:

Answer Area

 A data scientist

is responsible for identifying which business rules must be applied to the data of a company.

Question: 127

HOTSPOT

Select the answer that correctly completes the sentence.

Answer Area

 A data analyst
 A data engineer
 A data scientist
 A database administrator

is responsible for creating visuals and charts that help a company make informed decisions.

data Engineer for visuals

Answer:

Explanation:

Answer Area

A data scientist is responsible for identifying which business rules must be applied to the data of a company.

Question: 128

DRAG DROP

Match the types of workloads to the appropriate scenarios.

To answer, drag the appropriate workload type from the column on the left to its scenario on the right. Each workload type may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

Workload Types

Batch
Streaming

Answer Area

Workload type	Data for a product catalog will be loaded every 12 hours to a data warehouse.
Workload type	Thousands of data sets per second for online purchases will be loaded into a data warehouse in real time.
Workload type	Updates to inventory data will be loaded to a data warehouse every 1 million transactions.

Batch**Stream****Batch**

Answer:

Explanation:**Answer Area**

Batch	Data for a product catalog will be loaded every 12 hours to a data warehouse.
Streaming	Thousands of data sets per second for online purchases will be loaded into a data warehouse in real time.
Batch	Updates to inventory data will be loaded to a data warehouse every 1 million transactions.

Question: 129

What are two uses of data visualization? Each correct answer presents a complete solution. NOTE:

Each correct selection is worth one point.

- A. Communicate the significance of data.
- B. Represent trends and patterns over time.
- C. Implement machine learning to predict future values.
- D. Enforce business logic across reports.

Answer: A, B

Explanation:

Question: 130

What should you use to build a Microsoft Power BI paginated report?

- A. Power BI Report Builder
- B. Charciculator
- C. Power BI Desktop
- D. the Power BI service

Answer: A

Explanation:

Question: 131

Type text here

HOTSPOT

Select the answer that correctly completes the sentence.

Answer Area

The massively parallel processing (MPP) engine
of Azure Synapse Analytics

distributes processing across compute nodes.
distributes processing across control nodes.
redirects client connections across compute nodes.
redirects client connections across control nodes.

compute nodes to operate in parallel

Answer:

Explanation:

Answer Area

The massively parallel processing (MPP) engine
of Azure Synapse Analytics distributes processing across compute nodes.

Question: 132

Which scenario is an example of a streaming workload?

- A. sending transactions daily from point of sale (POS) devices
- B. sending cloud infrastructure metadata every 30 minutes
- C. sending transactions that are older than a month to an archive
- D. sending telemetry data from edge devices

Answer: D

Explanation:

Question: 133

What is the primary purpose of a data warehouse?

- A. to provide storage for transactional line-of-business (LOB) applications
- B. to provide transformation services between source and target data stores
- C. to provide read only storage of relational and non relational historical data
- D. to provide answers to complex queries that rely on data from multiple sources

confusing D looks C is correct but most say D

Answer: C

Explanation:

Question: 134**HOTSPOT**

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area

Statements

Azure Databricks is an Apache Spark-based analytics platform.
 Yes No
Azure Analysis Services is used for transactional workloads.
 Yes No
Azure Data Factory orchestrates data integration workflows.
 Yes No

Answer:

Explanation:

Answer Area

Statements

Azure Databricks is an Apache Spark-based analytics platform.
 Yes No
Azure Analysis Services is used for transactional workloads.
 Yes No
Azure Data Factory orchestrates data integration workflows.
 Yes No

Question: 135

You need to develop a solution to provide data to executives. The solution must provide an interactive graphical interface, depict various key performance indicators, and support data exploration by using drill down. What should you use in Microsoft Power BI?

- A. a report
- B. Microsoft Power Apps
- C. a view
- D. a dataflow

Answer: C

Explanation:

Question: 136

Your company has a reporting solution that has paginated reports. The reports query a dimensional model in a data warehouse. Which type of processing does the reporting solution use?

- A. Online Transaction Processing (OLTP)
- B. Online Analytical Processing (OLAP)
- C. batch processing
- D. stream processing

Answer: B

Explanation:

Question: 137

You need to create a visualization of running sales totals per quarter as shown in the following exhibit.



What should you create in Power BI Desktop;1

confirmed

- A. a waterfall chart **how an initial value is affected by a series of +ve and -ve values**
- B. a ribbon chart **-A combines ribbons for multiple categories into a single view.**
- C. a bar chart **-see data in ways, such as comparing numeric & Nominal discrete values**
- D. a decomposition tree

automatically aggregates data and enables drilling down into your dimensions in any order.

Answer: C

Explanation:

Question: 138

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No. NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
Stream processing has access to the most recent data received or data within a rolling time window.	<input checked="" type="radio"/>	<input type="radio"/>
Batch processing must occur immediately and have latency in the order of seconds or milliseconds.	<input type="radio"/>	<input checked="" type="radio"/>
Stream processing is used for simple response functions, aggregates, or calculations such as rolling averages.	<input checked="" type="radio"/>	<input type="radio"/>

Answer:

Explanation:

Answer Area

Statements	Yes	No
Stream processing has access to the most recent data received or data within a rolling time window.	<input checked="" type="radio"/>	<input type="radio"/>
Batch processing must occur immediately and have latency in the order of seconds or milliseconds.	<input checked="" type="radio"/>	<input type="radio"/>
Stream processing is used for simple response functions, aggregates, or calculations such as rolling averages.	<input type="radio"/>	<input checked="" type="radio"/>

Question: 139

Which Azure Storage service implements the key/value model?

- A. Azure Files

- B. Azure Blob
- C. Azure Table
- D. Azure Queue

Answer: C

Explanation:

Question: 140

HOTSPOT

Select the answer that correctly completes the sentence.

Answer Area

Varying fields for each entity in a JSON document is an example of

- relational data.
- semi-structured data.
- structured data.
- unstructured data.

semi structured data -mongodb

Answer:

Explanation:

Answer Area

Varying fields for each entity in a JSON document is an example of

- structured data.

Question: 141

What is used to define a query in a stream processing jobs in Azure Stream Analytics?

- A. SQL Go from zero to production in minutes using the
- B. XML no-code editor or SQL—which is easily extensible with custom code
- C. YAML and built-in machine learning capabilities for more
- D. KQL advanced scenarios

Answer: A

Explanation:

Question: 142

HOTSPOT

Select the answer that correctly completes the sentence.

Answer Area

In an infrastructure as a service (IaaS) instance of Microsoft SQL Server on Azure, you manage the

- elastic pool
- MySQL server
- PostgreSQL server
- virtual machine

that hosts SQL Server.



Answer:

Explanation:

Answer Area

In an infrastructure as a service (IaaS) instance of Microsoft SQL Server on Azure, you manage the virtual machine that hosts SQL Server.

Question: 143

HOTSPOT

Select the answer that correctly completes the sentence.

Answer Area

A relational database is appropriate for scenarios that involve a high volume of

changes to relationships between entities.
geographically distributed writes.
transactional writes.
writes that have varying data structures.

Explanation:

A relational database is appropriate for scenarios that involve a high volume of

 transactional writes.**Answer:****Question: 144**

Which property of a transactional workload guarantees that each transaction is treated as a single unit that either succeeds completely or fails completely?

- A. isolation
- B. atomicity
- C. consistency
- D. durability

Answer: B**Explanation:****Question: 145**

Which database transaction property ensures that transactional changes to a database are preserved during unexpected operating system restarts?

- A. durability
- B. atomicity
- C. consistency
- D. isolation

Answer: A**Explanation:****Question: 146****HOTSPOT**

Select the answer that correctly completes the sentence.

Answer Area

A JSON document is an example of

graph data.
relational data.
semi-structured data.
unstructured data.

semi-structured data**Answer:****Explanation:**

Answer Area

A JSON document is an example of relational data.

Question: 147**HOTSPOT**

Select the answer that correctly completes the sentence.

Answer Area

Relational data is stored in

a tabular form of rows & Columns

a file system as unstructured data.
a hierarchical folder structure.
a tabular form of rows and columns.
comma-separated value (CSV) files.

Answer:

Explanation:

Answer Area

Relational data is stored in a hierarchical folder structure.

Question: 148

Which database transaction property ensures that individual transactions are executed only once and either succeed in their entirety or roll back?

- A. consistency
- B. isolation
- C. atomicity
- D. durability

Answer: A

Explanation:

Question: 149

You manage an application that stores data in a shared folder on a Windows server. You need to move the shared folder to Azure Storage. Which type of Azure Storage should you use?

- A. table
- B. queue
- C. file **use file sync**
- D. blob

Answer: C

Explanation:

Question: 150

You need to recommend a non-relational data store that is optimized for storing and retrieving text

files, videos, audio streams, and virtual disk images. The data store must store data, some metadata, and a unique ID for each file. Which type of data store should you recommend?

- A. columnar
- B. key/value
- C. document
- D. object

Answer: D

Explanation:

Question: 151

HOTSPOT

Select the answer that correctly completes the sentence.

Answer Area

When using the Azure Cosmos DB Gremlin API, the container resource type is projected as a

graph.
table.
partition key.
document.

Answer:

Explanation:

Answer Area

When using the Azure Cosmos DB Gremlin API, the container resource type is projected as a

graph.

Question: 152

DRAG DROP

Match the Azure Cosmos DB APIs to the appropriate data structures.

To answer, drag the appropriate API from the column on the left to its data structure on the right.

Each API may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

APIs

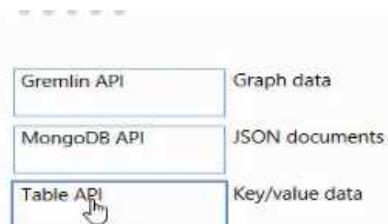
- Cassandra API
- Gremlin API
- MongoDB API
- Table API

Answer Area

- | | |
|-----|----------------|
| API | Graph data |
| API | JSON documents |
| API | Key/value data |

Answer:

Explanation:

Answer Area**Question: 153**

You need to store event log data that is semi-structured and received as the logs occur. What should you use?

- A. Azure Table storage
B. Azure Queue storage
C. Azure Files
- Azure Table storage is a NoSQL key-value store specifically designed for large amounts of semi-structured data. This means it can handle data that doesn't fit a rigid schema, which is common for event logs where fields might vary**

Answer: A

Explanation:

Question: 154

What should you use to automatically delete blobs from Azure Blob Storage?

- A. the change feed
B. a lifecycle management policy
C. soft delete
D. archive storage

Answer: D

Explanation:

Question: 155

What is a characteristic of a non-relational database?

- A. full support for Transact-SQL
B. a fixed schema
C. self describing entities

In many non-relational databases (especially document databases like MongoDB or JSON-based databases), the data itself often contains information about its structure and meaning. Each "entity" (e.g., a document) can define its own fields and their types, making it "self-describing" rather than relying on a separate, predefined schema.

Answer: C

Explanation:

Question: 156

HOTSPOT

Select the answer that correctly completes the sentence.

Answer Area

When provisioning an Azure Cosmos DB



, you need to specify which type of API you will use.

Account level- determines which API you'll use for interacting with all databases, containers, and items within that specific account

Answer:

Explanation:

Answer Area

When provisioning an Azure Cosmos DB



, you need to specify which type of API you will use.

Question: 157

You are building a retail kiosk system that will use a custom neural voice. You acquire audio samples and consent from the voice talent. You need to create a voice talent profile. What should you upload to the profile?

- A. a five-minute wav or mp3 file of the voice talent describing the kiosk system
- B. a five-minute .flac audio file and the associated transcript as a w file
- C. a .wav or mp3 file of the voice talent consenting to the creation of a synthetic version of their voice
- D. a .zip file that contains 10-second .wav files and the associated transcripts as .txt files

. a .wav or mp3 file of the voice talent

Answer: D

Explanation: **consenting to the creation of a synthetic version of their voice.**

Question: 158

You have an app named App1 that uses an Azure Cognitive Services model to identify anomalies in a time series data stream. You need to run App1 in a location that has limited connectivity. The solution must minimize costs. What should you use to host the model?

- A. Azure Kubernetes Services (AKS)
- B. a Kubernetes cluster hosted in an Azure Stack Hub integrated system
- C. Azure Container instances
- D. the Docker Engine **in a location with limited connectivity and minimizing costs**

Docker Engine is the most lightweight and cost-effective approach for local deployment

Answer: B

Explanation:

not incurring ongoing Azure infrastructure costs for the model's runtime (beyond the initial Cognitive Services resource cost for billing).

You plan to create an index for an Azure Cognitive Search service by using the Azure portal. The Cognitive Search service will connect to an Azure SQL database

The Azure SQL database contains a table named UserMessages. Each row in UserMessages has a field named MessageCopy that contains the text of social media messages sent by a user

Users will perform full text searches against the MessageCopy field, and the values of the field will be shown to the users-

You need to configure the properties of the index for the MessageCopy field to support the solution. Which attributes should you enable for the field?

- A. Searchable and Retrievable
- B. Sortable and Retrievable
- C. Searchable and Facetable
- D. Filterable and Retrievable

Searchable-This attribute allows the field's content to be included in full-text search queries

Retrievable-original content to be retrieved in search results.

Answer: A

Explanation:

Question: 160

HOTSPOT

You are developing a text processing solution.

You have the function shown below.

```
static void GetKeywords(TextAnalyticsClient textAnalyticsClient, string text)
{
    var response = textAnalyticsClient.RecognizeEntities (text);
    Console.WriteLine("Key words:");

    foreach (CategorizedEntity entity in response.Value)
    {
        Console.WriteLine($"{entity.Text}");
    }
}
```

Answer Area

Statements

The output will include the following words: our and included.

Yes

No



The output will include the following words: Paris, Eiffel, and Tower.



The function will output all the key phrases from the input string to the console.



Answer:

Explanation:

Answer Area

Statements

The output will include the following words: our and included.

Yes

No



The output will include the following words: Paris, Eiffel, and Tower.



The function will output all the key phrases from the input string to the console.



Question: 161

HOTSPOT

You have a library that contains thousands of images.

You need to tag the images as photographs, drawings, or clipart.

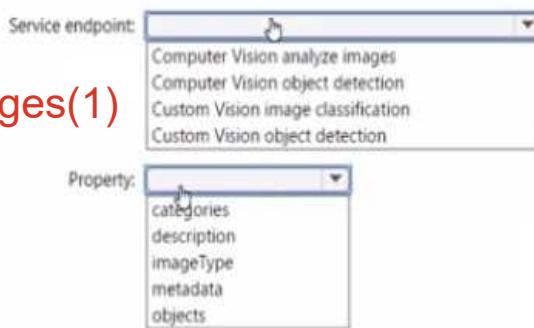
Which service endpoint and response property should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

computer vision analyze images(1)

imageType(3)

**Answer:**

Explanation:

Answer Area

Service endpoint: Computer Vision object detection
Property: description

Analyze Image -endpoint of the Azure AI Vision service.
This endpoint allows you to send an image and request various insights about its content.

Question: 162**HOTSPOT**

You are building a model to detect objects in images.

The performance of the model based on training data is shown in the following exhibit.

The screenshot shows the 'Performance' tab of the Azure Custom Vision service. At the top, there are tabs for 'Objects', 'Training Images', 'Performance' (which is selected), 'Predictions', 'Train', and 'Quick Test'. Below the tabs are buttons for 'Iterations', 'Publish', 'Delete', and 'Export'. A slider for 'Probability Threshold' is set to 50%, and a slider for 'Overlap Threshold' is set to 30%. The main area displays 'Iteration 1' results. The first section states: 'The percentage of false positives is [answer choice].'. A dropdown menu shows options: 0, 25, 30, 50, 100. The second section states: 'The value for the number of true positives divided by the total number of true positives and false negatives is [answer choice]%.'. Another dropdown menu shows options: 0, 25, 30, 50, 100.

Answer:

Explanation:

Answer Area

confirmed answer

The percentage of false positives is [answer choice].

0 25 30 50 100

The value for the number of true positives divided by the total number of true positives and false negatives is [answer choice]%.
50

Question: 163

You have an Azure IoT hub that receives series data from machinery. You need to build an app that

will perform the following actions:

- Perform anomaly detection across multiple correlated sensors
- Identify the root cause of process stops.
- Send incident alerts

The solution must minimize development time. Which Azure service should you use?

- A. Azure Metrics Advisor
B. Form Recognizer
C. Azure Machine Teaming
D. Anomaly Detector

with Azure AI Metrics Advisor, built on AI Anomaly Detector, and a part of Azure AI Services.
Quickly identify and fix problems through a powerful combination of monitoring in near-real time, adapting models to your scenario, and offering granular analysis with diagnostics and alerting.

Answer: D

Explanation:

Question: 164

You build a language model by using Conversational Language Understanding. The language model is used to search for information on a contact list by using an intent named FindContact. A conversational expert provides you with the following list of phrases to use for training

- Find contacts in London.
- Who do I know in Seattle?
- Search for contacts in Ukraine.

You need to implement the phrase list in Conversational Language Understanding.

Solution: You create a new utterance for each phrase in the FindContact intent.

- A. Yes
B. No

provide various example phrases that a user might say to express that intent. These example phrases are called utterances.
phrase list - a hint to understand vocabulary or relationship
create phraselist named:Account

Answer: B

Explanation:

Question: 165

for account or banking below vocabulary1)card

**2) "checking account" 3)"saving account" and
4)cheque**

You develop a Conversational Language Understanding model by using Language Studio

During testing, users receive incorrect responses to requests that do NOT relate to the capabilities of the model.

You need to ensure that the model identifies spurious requests.

What should you do?

- A. Enable active learning.
B. Add examples to the custom intents.
C. Add examples to the None intent
D. Add entities.

you teach it to recognize when a user's input is outside its capabilities.

Answer: A

Explanation:

Question: 166

You have an Azure Cognitive Search instance that indexes purchase orders by using Form Recognizer. You need to analyze the extracted information by using Microsoft Power BI. The solution must minimize development effort.

What should you add to the indexer?

- A. a table projection
- B. a projection group
- C. an object projection
- D. a file projection

is specifically designed to take repeating, structured data (like the line items from a purchase order extracted by Form Recognizer) and project them into a separate "table" within the search index.

Answer: C

Explanation:

This allows Power BI to easily connect to this projected table and treat each line item as a distinct row, making analysis straightforward

Question: 167

You develop a custom question answering project in Azure Cognitive Service for Language. The project will be used by a chatbot. You need to configure the project to engage in multi-turn conversations. What should you do?

- A. Add follow-up prompts.
- B. Enable active learning.
- C. Add alternate questions.
- D. Enable chit-chat.

Answer: A

Explanation:

Question: 168**HOTSPOT**

You are building an app that will enable users to upload images. The solution must meet the following requirements:

- Automatically suggest alt text for the images.
- Detect inappropriate images and block them.
- Minimize development effort.

You need to recommend a computer vision endpoint for each requirement.

What should you recommend? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Sending the image as raw binary data in the body of the HTTP request

Answer Area

Generate alt text:

```
https://westus.api.cognitive.microsoft.com/customvision/v3.1/prediction/projectId/classify/iterations/publishedName/image
https://westus.api.cognitive.microsoft.com/contentmoderator/moderate/v1.0/ProcessImage/Evaluate
https://westus.api.cognitive.microsoft.com/customvision/v3.1/prediction/projectId/classify/iterations/publishedName/image
https://westus.api.cognitive.microsoft.com/vision/v3.2/analyze/?visualFeatures=Adult,Description
```

Detect inappropriate content:

```
https://westus.api.cognitive.microsoft.com/vision/v3.2/analyze/?visualFeatures=Adult,Description
https://westus.api.cognitive.microsoft.com/contentmoderator/moderate/v1.0/ProcessImage/Evaluate
https://westus.api.cognitive.microsoft.com/customvision/v3.1/prediction/projectId/classify/iterations/publishedName/image
https://westus.api.cognitive.microsoft.com/vision/v3.2/analyze/?visualFeatures=Adult,Description
https://westus.api.cognitive.microsoft.com/vision/v3.2/describe?maxCandidates=1
```

<https://westus.api.cognitive.microsoft.com/vision/v3.2/analyze?visualFeatures=Adult,Description> Answer:

Explanation: <https://westus.api.cognitive.microsoft.com/vision/v3.2/analyze?visualFeatures=Adult,Description>

can fulfill both requirements simultaneously:

1) Generate alt text:

2) Detect inappropriate content:

Question: 169

HOTSPOT

You are building content for a video training solution.

You need to create narration to accompany the video content. The solution must use Custom Neural Voice.

What should you use to create a custom neural voice, and which service should you use to generate the narration? To answer, select the appropriate options in the answer area.

NOTE: Each correct answer is worth one point.

Answer Area

Custom neural voice: The Language Understanding portal
Microsoft Bot Framework Composer
The Azure portal
The Language Understanding portal
The Speech Studio portal

Narration: Language Understanding
Language Understanding
Speaker Recognition
Speech-to-text
Text-to-speech

Answer:

Explanation:

Answer Area

Custom neural voice: The Language Understanding portal
Speech studio portal

Narration: Language Understanding
text to speech

Speechstudio-This is the platform where you upload voice talent data, train, and deploy your unique custom voice model.

Question: 170 You need to measure the public perception of your brand on social media by using natural language processing. Which Azure service should you use?

- A. Content Moderator
- B. Form Recognizer
- C. Computer Vision
- D. Language service**

Language Service provide pre-built natural language processing (NLP) capabilities directly relevant to understanding public perception

Answer: D

Explanation:

Question: 171

You need to measure the public perception of your brand on social media by using natural language processing. Which Azure service should you use?

- A. Content Moderator
- B. Form Recognizer
- C. Computer Vision
- D. Language service

Answer: D

Explanation:

Question: 172

HOTSPOT

You have a bot that was built by using the Microsoft Bot Framework composer as shown in the following exhibit.



Use the drop-down menus 10 select the answer choice that completes each statement based on the information presented in the graphic.

NOTE: Each correct selection is worth one point.

Answer Area

If a user asks "what is the weather like in New York", the bot will [answer choice].

The GetWeather dialog uses a [answer choice] trigger.

Identify New York as a city entity
change to a different dialog
Identify New York as a city entity
identify New York as a state entity
respond with the weather in Seattle

Language Understanding Intent recogn...
Custom events
Dialog events
Language Understanding Intent recognized
QnA Intent recognized

Answer:

Explanation:



If a user asks "what is the weather like in New York", the bot will [answer choice].

The GetWeather dialog uses a [answer choice] trigger.


Question: 173

HOTSPOT

You are building a chatbot by using the Microsoft Bot Framework Composer. You have the dialog

Type here

When AI is used for high-stakes decisions that directly impact individuals (like their income), human oversight is critical to catch errors, ensure fairness, provide explanations, and allow for recourse.



For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
user.name is an entity.	<input type="radio"/>	<input checked="" type="radio"/>
The dialog asks for a user name and a user age and assigns appropriate values to the user.name and user.age properties.	<input type="radio"/>	<input checked="" type="radio"/>
The chatbot attempts to take the first non-null entity value for userName or personName and assigns the value to user.name.	<input checked="" type="radio"/>	<input type="radio"/>

Answer:

Explanation:

Answer Area

Statements	Yes	No
user.name is an entity.	<input type="radio"/>	<input checked="" type="radio"/>
The dialog asks for a user name and a user age and assigns appropriate values to the user.name and user.age properties.	<input checked="" type="radio"/>	<input type="radio"/>
The chatbot attempts to take the first non-null entity value for userName or personName and assigns the value to user.name.	<input checked="" type="radio"/>	<input type="radio"/>

Repeated

Question: 174

You are examining the Language service output of an application.

The text analyzed is: Our tour guide took us up the Space Needle during our trip to Seattle last week.

The response contains the data shown in the following table.

Text	Category	ConfidenceScore
Tour guide	PersonType	0.45
Space Needle	Location	0.38
Trip	Event	0.78
Seattle	Location	0.78
Last week	DateTime	0.80

Which Language service API is used to analyze the Text?

- A. Entity Linking
- B. Named Entity Recognition
- C. Key Phrase Extraction
- D. Sentiment Analysis

Repeated

Answer: B

Explanation:

Question: 175

DRAG DROP

You plan to build a chatbot to support task tracking.

You create a Conversational Language Understanding service named lu1.

You need to build a Conversational Language Understanding model to Integrate into the chatbot. The solution must minimize development time to build the model.

Which four actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions	Answer Area
Add the prebuilt domain ToDo.	
Add a new application.	
Add example utterances.	
Train the application.	
Publish the application.	

(Left) (Right)

Answer:

Explanation:

Actions	Answer Area
Add the prebuilt domain ToDo.	

(Left) (Right)

Repeated

The model needs a way to represent devices like "lights" and "thermostat"

An entity could be created to identify the type of device being referred to in the utterance

Question: 176

You are building a Conversational Language Understanding model.

You need to ensure that the model will support the following sample utterances:

- Set all the lights to on.
- Turn off the lights in the living room.
- What is the current thermostat temperature?
- Lower the temperature of the thermostat by five degrees.

Which three elements should you add to the model?

Each correct answer presents part of the solution. NOTE: Each correct selection is worth one point.

- A. a location Intent
- B. a change setting entity
- C. a device intent
- D. a change setting intent
- E. a query setting intent
- F. a device entity

changing a device setting and querying a device setting.
The utterances also involve devices (lights, thermostat) which require an entity to represent them.

1)Entity- lights,livingroom
off is state entity
2)utterance-
phrase - whole phrase
"turn on light in living room
3)intent - overall goal
change setting, turnon light

Answer: B, C, F

Explanation:

Question: 177

Type text here

DRAG DROP

You build a bot by using the Microsoft Bot Framework SDK.

You need to test the bot interactively on a local machine.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

NOTE: More than one order of answer choices is correct. You will receive credit for any of the correct orders you select.

Actions	Answer Area
Register the bot with the Azure Bot Service.	
Open the Bot Framework Composer.	
Build and run the bot.	
Open the Bot Framework Emulator.	
Connect to the bot endpoint.	

Drag and drop the actions from the left Actions list to the right Answer Area. Use the up and down arrows to rearrange the actions in the Answer Area.

Answer:

Explanation:

Actions	Answer Area
Register the bot with the Azure Bot Service.	1 Build and run the bot.
Open the Bot Framework Composer.	2 Open the Bot Framework Emulator.
	3 Connect to the bot endpoint.

A red checkmark is placed next to the "Actions" list, indicating it has been completed. The "Answer Area" list contains three numbered items: 1. Build and run the bot., 2. Open the Bot Framework Emulator., and 3. Connect to the bot endpoint. The up and down arrows are present on both sides of the lists.

confirmed**Question: 178**

You are building a social media extension that will convert text to speech. The solution must meet the following requirements:

- Support messages of up to 400 characters.
- Provide users with multiple voice options.
- Minimize costs.

You create an Azure Cognitive Services resource.

Which Speech API endpoint provides users with the available voice options?

A.

<https://uksouth.customvoice.apispeech.microsoft.eom/api/texttospeech/v3.0/longaudiosynthesis/voices> **custom voice management or long audio synthesis**

- B. <https://uksouth.tts.speech.microsoft.conT/cognitiveservices/voices/list> - deployment
C. <https://uksouth.voice.speech.microsoft.com/cognitiveservices/v1?deploymentId = {deploymentId}>
D. <https://uksouth.api.cognitive.microsoft.eom/speechtotext/v3.0/models/base> - listing models, not voices.

Answer: C

Explanation:

(tts.speech.microsoft.com) service, and the path /cognitiveservices/voices/list explicitly signifies an endpoint designed to list the available voices.

This is the standard endpoint for this purpose.

You have a Custom Vision service project that performs object detection. The project uses the General domain for classification and contains a trained model.

You need to export the model for use on a network that is disconnected from the internet.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

2

Actions
Change Domains to General (compact) .
Retrain the model.
Create a new classification model.
Change the classification type.
Export the model.

1

3

confirmed

Answer Area

1 Create a new classification model.
2 Change the classification type.
3 Export the model.

Up **Down**

Answer:

Explanation:

Actions
Change Domains to General (compact) .
Retrain the model.

To export a Custom Vision object detection model for offline use on a disconnected network, you need to ensure the model is trained

Question: 180 with an exportable domain and then proceed with the export

DRAG DROP

You have a question answering project in Azure Cognitive Service for Language.

You need to move the project to a Language service instance in a different Azure region.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

3

Actions
From the new Language service instance, import the project file.
From the new Language service instance, enable custom text classification.
From the new Language service instance, train and publish the project.
From the original Language service instance, export the existing project.
From the new Language service instance, regenerate the keys.
From the original Language service instance, train and publish the model.

2 **1** From the original Language service instance, export the existing project.

1



Answer:

Explanation:

Actions
From the new Language service instance, import the project file.
From the new Language service instance, enable custom text classification.
From the new Language service instance, train and publish the project.

Answer Area

1 From the original Language service instance, export the existing project.
2 From the new Language service instance, regenerate the keys.
3 From the original Language service instance, train and publish the model.



2. From the new Language service instance, import the project file. - export to desired azure region
3. From the new Language service instance, train and publish the project.

Question: 181

You are building a chatbot that will use question answering in Azure Cognitive Service for Language.

You have a PDF named Doc1.pdf that contains a product catalogue and a price list

You upload Doc1.pdf and train the model.

During testing, users report that the chatbot responds correctly to the following question What is the price of < product>?

The chatbot fails to respond to the following question How much does <product* cost?

You need to ensure that the chatbot responds correctly to both questions.

Solution: From Language Studio, you create an entity for cost, and then retrain and republish the model.

Does this meet the goal?

A. Yes

B. No

this kind of question variation is typically handled by adding alternate questions to an existing Question-Answer pair.

helps in extracting the word "cost" as a specific data point, but it does not teach the model that "How much does X cost?"

Answer: B

Question: 182

To solve this, you should add "How much does <product> cost?" as an alternate question to the Q& A pair that correctly answers "What is the price of <product>?". After adding the alternate question, retraining and republishing the model would then make the chatbot respond correctly to both.

DRAG DROP When AI is used for high-stakes decisions that directly impact individuals (like their income), human oversight is critical to catch errors, ensure fairness, provide explanations, and allow for recourse.

You are building an app that will scan confidential documents and use the Language service to analyze the contents.

You provision an Azure Cognitive Services resource.

You need to ensure that the app can make requests to the Language service endpoint. The solution must ensure that confidential documents remain on-premises.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions
Pull an image from Docker Hub.
Provision an on-premises Kubernetes cluster that has internet connectivity.
Provision an Azure Kubernetes Service (AKS) resource.
Run the container and specify an App ID and Client Secret.
Provision an on-premises Kubernetes cluster that is isolated from the internet.
Pull an image from the Microsoft Container Registry (MCR).
Run the container and specify an API key and the Endpoint URL of the Cognitive Services resource.

Answer Area
1. Pull an image from microsoftContainer registry
2. Provision an on-premise kubernetes cluster that has internet connectivity
3. Run the container and specify an API key and the Endpoint URL of the Cognitive Services resource.

Explanation:

Actions
Pull an image from Docker Hub.
Provision an on-premises Kubernetes cluster that has internet connectivity.
Provision an Azure Kubernetes Service (AKS) resource.
Run the container and specify an App ID and Client Secret.

Answer Area
1. Provision an on-premises Kubernetes cluster that is isolated from the internet.
2. Pull an image from the Microsoft Container Registry (MCR).
3. Run the container and specify an API key and the Endpoint URL of the Cognitive Services resource.

Question: 183

You have a factory that produces food products.

You need to build a monitoring solution for staff compliance with personal protective equipment (PPE) requirements. The solution must meet the following requirements:

- identify staff who have removed masks or safety glasses.
- Perform a compliance check every 15 minutes.
- Minimize development effort.
- Minimize costs.

Which service should you use?

- A. Face
- B. Computer Vision
- C. Azure Video Analyzer for Media (formerly Video indexer)

A. Face (Azure AI Face service): its primary focus is on facial detection and analysis. While it can detect the presence of a mask and general glasses, it lacks the specific object detection capabilities and customizability required for precise PPE compliance (especially for "safety glasses" and nuanced "removed" states) in an industrial setting.

Explanation:

Answer: A

Question: 184

HOTSPOT

You have an Azure subscription that has the following configurations:

- Subscription ID: 8d3591aa-96b8-4737-ad09-00f9b1ed35ad
- Tenant ID: 3edfe572-cb54-3ced-ae12-c5c177f39a12

You plan to create a resource that will perform sentiment analysis and optical character recognition

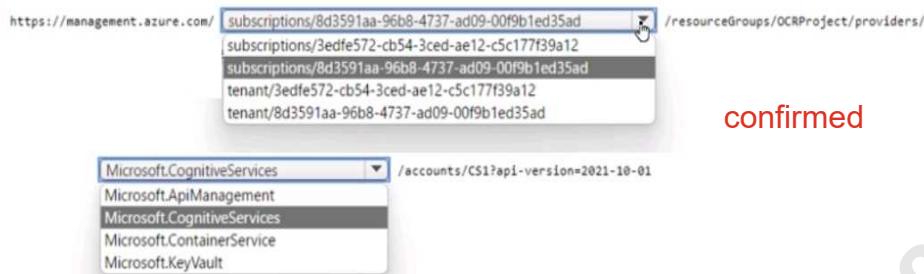
(OCR).

You need to use an HTTP request to create the resource in the subscription. The solution must use a single key and endpoint.

How should you complete the request? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area



confirmed

Answer:

Explanation:

Answer Area



Question: 185

You have a text-based chatbot.

You need to enable content moderation by using the Text Moderation API of Content Moderator.

Which two service responses should you use? Each correct answer presents part of the solution

NOTE: Each correct selection is worth one point.

- A. the adult classification score
- B. optical character recognition (OCR)
- C. personal data **The API can detect and identify personally identifiable information (PII)**
- D. text classification **within the text,**
- E. the racy classification score

text classification - broader feature within the Text Moderation API that categorizes content based on its likelihood of being inappropriate (e.g., indicating adult or racy content).

Answer: C, D

Explanation:

Question: 186

You are building an AI solution that will use Sentiment Analysis results from surveys to calculate bonuses for customer service staff. You need to ensure that the solution meets the Microsoft responsible AI principles. What should you do?

- A. Add a human review and approval step before making decisions that affect the staffs financial situation
- B. Include the Sentiment Analysis results when surveys return a low confidence score.
- C. Use all the surveys, including surveys by customers who requested that their account be deleted and their data be removed.
- D. Publish the raw survey data to a **here** location and provide the staff with access to the location.

When AI is used for high-stakes decisions that directly impact individuals (like their income), human oversight is critical to catch errors, ensure fairness, provide explanations, and allow **Answer: A**

Explanation: for recourse.

Question: 187

DRAG DROP

You have a Language Understanding solution that runs in a Docker container.

You download the Language Understanding container image from the Microsoft Container Registry (MCR).

You need to deploy the container image to a host computer.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions	Answer Area
From the Language Understanding portal, retrain the model.	 
From the host computer, run the container and specify the input directory.	
From the Language Understanding portal, export the solution as a package file.	
From the host computer, move the package file to the Docker input directory.	
From the host computer, build the container and specify the output directory.	

Answer:

Explanation:

Actions	Answer Area
From the Language Understanding portal, retrain the model.	 
From the host computer, run the container and specify the input directory.	
From the Language Understanding portal, export the solution as a package file.	
From the host computer, move the package file to the Docker input directory.	
From the host computer, build the container and specify the output directory.	

1. From the Language Understanding portal, export the solution as a package file.
2. From the host computer, move the package file to the Docker input directory.
3. From the host computer, run the container and specify the input directory.

When you run the container, you will specify the mounted input directory so that the Language Understanding service inside the container can load and use your custom model.

Question: 188

DRAG DROP

You are building a customer support chatbot.

You need to configure the bot to identify the following:

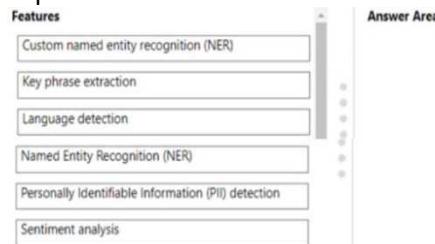
- Code names for internal product development
- Messages that include credit card numbers

The solution must minimize development effort.

Which Azure Cognitive Service for Language feature should you use for each requirement? To answer, drag the appropriate features to the correct requirements. Each feature may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content

NOTE: Each correct selection is worth one point.

Features	Answer Area
Custom named entity recognition (NER)	<div style="display: flex; justify-content: space-around;"> <div style="width: 45%;"> Identify code names for internal product development: <input type="text"/> </div> <div style="width: 45%;"> Identify messages that include credit card numbers: <input type="text"/> </div> </div>
Key phrase extraction	
Language detection	
Named Entity Recognition (NER)	
Personally identifiable information (PII) detection	
Sentiment analysis	
Text classification	

Explanation:**Question: 189**

You are building a chatbot.
You need to configure the bot to guide users through a product setup process.
Which type of dialog should you use?

- A. component
- B. waterfall
- C. adaptive
- D. action

This sequential nature makes it ideal for step-by-step procedures like a product setup.

Answer: B**Explanation:****Question: 190**

You have an Azure subscription that contains an Azure Cognitive Service for Language resource. You need to identify the URL of the REST interface for the Language service. Which blade should you use in the Azure portal?

- A. Identity
- B. Keys and Endpoint
- C. Properties
- D. Networking

Duplicate

Answer: B**Explanation:****Question: 191**

You are building a chatbot by using Microsoft Bot Framework Composer.
You need to configure the chatbot to present a list of available options. The solution must ensure that an image is provided for each option.
Which two features should you use? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point.

- A. an Azure function

- B. an adaptive card **flexible and platform-agnostic way to display rich UI**
 C. an entity
 D. a dialog **specially as part of a product setup process, occurs within the context of a dialog. The dialog orchestrates when and how these Adaptive Cards**
 E. an utterance

Answer: B, D

Explanation:

Question: 192

You are building a chatbot.
 You need to configure the chatbot to query a knowledge base.
 Which dialog class should you use?

- A. AdaptiveDialog
 B. QnAMakerDialog
 C. ComponentDialog
 D. SkillDialog

Answer: B

Explanation:

Question: 193

DRAG DROP

You need to analyze video content to identify any mentions of specific company names.
 Which three actions should you perform in sequence? To answer move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions
Add the specific company names to the exclude list.
Sign in to the Custom Vision website.
From Content model customization, select Language .
Sign in to the Azure Video Analyzer for Media website.
From Content model customization, select Brands .
Add the specific company names to the include list.

- Answer Area**
- 1.Sign in to the Azure Video Analyzer for Media website.
 - 2.From Content model customization, select **Brands**.
 - 3.Add the specific company names to the include list.

Explanation:

Actions
Add the specific company names to the exclude list.
Sign in to the Custom Vision website.
From Content model customization, select Language .

Answer Area
1 Sign in to the Azure Video Analyzer for Media website.
2 From Content model customization, select Brands .
3 Add the specific company names to the include list.

2.you can customize detection models.
To identify specific company names,
you need to work with the "Brands"



After selecting the "Brands" model for customization, you will add the specific company names you want to detect to an "include list" (or a custom brand list). This tells the service to actively search for and identify these defined brands in your video content.

Question: 194

You plan to build an app that will generate a list of tags for uploaded images. The app must meet the

following requirements:

- Generate tags in a user's preferred language.
- Support English, French, and Spanish.
- Minimize development effort

You need to build a function that will generate the tags for the app. Which Azure service endpoint should you use?

A. Custom Vision image classification **Custom Vision is designed for building custom image classification or object detection models.**

B. Content Moderator Image Moderation

C. Custom Translator

custom vision wont support multiple languages

D. Computer Vision Image Analysis

that can generate a list of relevant tags (keywords) based on the content of an image. Crucially, this service supports generating these tags in

Answer: A

Explanation: **multiple languages directly, including English, French, and Spanish, by specifying**

Question: 195

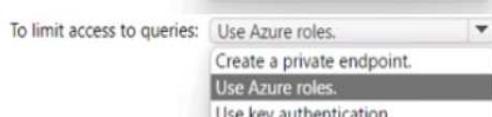
HOTSPOT

You have an Azure Cognitive Search resource named Search 1 that is used by multiple apps. You need to secure Search 1. The solution must meet the following requirements:

- Prevent access to Search1 from the internet.
- Limit the access of each app to specific queries

What should you do? To answer, select the appropriate options in the answer area. NOTE: Each correct answer is worth one point.

Answer Area



Answer:

Explanation:

Answer Area

To prevent access from the internet: **Create a private endpoint.**

To limit access to queries: **Use Azure roles.**

access control filters or row-level security on search index

Question: 196

You have a mobile app that manages printed forms.

You need the app to send images of the forms directly to Forms Recognizer to extract relevant information. For compliance reasons, the image files must not be stored in the cloud.

In which format should you send the images to the Form Recognizer API endpoint?

- A. raw image binary
- B. form URL encoded

Sending the image as raw binary data in the body of the HTTP request means the image data is processed directly by the API endpoint and is typically not persisted to cloud storage

D. JSON

Answer: A

Explanation:

Question: 197

DRAG DROP

You have a factory that produces cardboard packaging for food products. The factory has intermittent internet connectivity.

The packages are required to include four samples of each product.

You need to build a Custom Vision model that will identify defects in packaging and provide the location of the defects to an operator. The model must ensure that each package contains the four products.

Which project type and domain should you use? To answer, drag the appropriate options to the correct targets. Each option may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content

NOTE: Each correct selection is worth one point.

Options	Answer Area
Food	Project type: <input type="text"/>
General	Domain: <input type="text"/>
General (compact)	
Image classification	
Logo	
Object detection	

1.Object Detection
2.General(Compact)

Answer:

Explanation:

Options	Answer Area
Food	Project type: <input type="text"/> Object detection
General	Domain: <input type="text"/> Food
General (compact)	
Image classification	
Logo	
Object detection	

"General (Compact)" domain is optimized for common object detection scenarios and is specifically designed for deployment on edge devices, which aligns with the factory's "intermittent internet connectivity"

Question: 198

You have an app that analyzes images by using the Computer Vision API.

You need to configure the app to provide an output for users who are vision impaired. The solution must provide the output in complete sentences.

Which API call should you perform? Type text here

- A. readInStreamAsync
- B. describeImageInStreamAsync
- C. toggleFlagInStreamAsync

D. analyzeImageByDomainInStreamAsync

Answer: D

Explanation:

Question: 199

This API call from Azure AI Vision (Computer Vision) generates a human-readable sentence that describes the content of an image. This is specifically designed for scenarios like providing alt-text for accessibility, directly meeting the requirement for output in complete sentences for vision-impaired users.

You are designing a conversational interface for an app that will be used to make vacation requests.

The interface must gather the following data:

- The start date of a vacation
- The end date of a vacation
- The amount of required paid time off

The solution must minimize dialog complexity. Which type of dialog should you use?

- A. Skill
B. waterfall
C. adaptive
D. component

Waterfall dialogs are the most appropriate choice for this scenario. They are designed for linear, step-by-step conversations, which aligns perfectly with gathering the required information in a sequential manner. Each step in the waterfall can prompt for a specific piece of data: start date, end date, and then the amount of paid time off

Answer: D

Explanation:

Question: 200

You have a Language service resource that performs the following:

- Sentiment analysis
- Named Entity Recognition (NER)
- Personally Identifiable Information (PII) identification

You need to prevent the resource from persisting input data once the data is analyzed. Which query parameter in the Language service API should you configure?

- A. loggingOptOut
B. piiCategories
C. showStats -set to true request indicate request statistics in the response, but not affect data saved.
D. Model-version model version to do analysis

This query parameter, when set to true, prevents the Language service from logging or persisting the input text after analysis.

This directly addresses the requirement of preventing data persistence.

piiCategories: This parameter is used to specify which categories of Personally Identifiable Information (PII) to detect

Answer: A

Explanation:

Question: 201

You have an Azure subscription that contains a Language service resource named ta1 and a virtual network named vnet1. You need to ensure that only resources in vnet1 can access ta1. What should you configure?

- A. a network security group (NSG) for vnet1 used to filter network traffic within a virtual network

- B. Azure Firewall for vnet1 used to inspect and filter network traffic. It can also be used to restrict traffic
C. the virtual network settings for ta1 -deploy and run Language service APIs locally,
D. a Language service container for ta1either on-premises or in other cloud

Vnet - allows to create own private network in azure like isolated network that you control

subnet - logical division in vnet e.g 10.0.1.0/24 with in vnet

Answer: B

Explanation: you are instructing Azure that only traffic originating from vnet1's subnets can reach the ta1 Language service resource.

Question: 202 This is done by adding the allowed network's subnets to the service's networking configuration.

This eliminates the need to configure an NSG rule to filter traffic to the language service based on IP addresses

HOTSPOT

You have a collection of press releases stored as PDF files.

You need to extract text from the files and perform sentiment analysis.

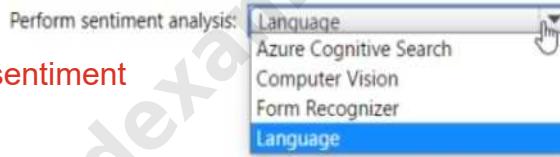
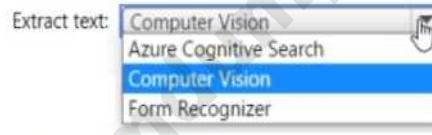
Which service should you use for each task? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Form Recognizer (Azure document Intelligence)
- text,tables, and other structured data from pdf.

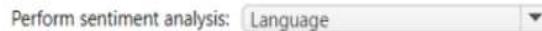
Azure AI Language or Language to pull sentiment



Answer:

Explanation:

Answer Area



Question: 203

HOTSPOT

You are building a chatbot.

You need to use the Content Moderator service to identify messages that contain sexually explicit language.

Which section in the response from the service will contain the category score, and which category will be assigned to the message? To answer, select the appropriate options in the answer area,

NOTE: Each correct selection is worth one point.

Answer Area

classification - section contain category score

1

Section: **Classification**

Classification
pii
Terms

Category: **1**

3
1
2
3

Answer:

Explanation:

Answer Area

Section: **Classification**

Category: **3**

Question: 204

HOTSPOT

You are building a solution that students will use to find references for essays. You use the following code to start building the solution.

```
using Azure;
using System;
using Azure.AI.TextAnalytics;

private static readonly AzureKeyCredential credentials = new AzureKeyCredential("<key>");
private static readonly Uri endpoint = new Uri("<endpoint>");

static void EntityLinker(TextAnalyticsClient client)
{
    var response = client.RecognizeLinkedEntities(
        "Our tour guide took us up the Space Needle during our trip to Seattle last week.");
}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No. NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
The code will detect the language of documents.	<input type="radio"/>	<input checked="" type="radio"/>
The url attribute returned for each linked entity will be a Bing search link.	<input type="radio"/>	<input checked="" type="radio"/>
The matches attribute returned for each linked entity will provide the location in a document where the entity is referenced.	<input checked="" type="radio"/>	<input type="radio"/>

Answer:

Explanation:

This method takes a piece of text and tries to find "things" (entities) mentioned in it, like people, places, organizations, etc.

Then, it goes a step further by attempting to link these entities to their real-world counterparts, such as Wikipedia entries, to provide more context and information

Answer Area

Statements	Yes	No
The code will detect the language of documents.	<input type="radio"/>	<input checked="" type="radio"/>
The <code>url</code> attribute returned for each linked entity will be a Bing search link.	<input checked="" type="radio"/>	<input type="radio"/>
The <code>matches</code> attribute returned for each linked entity will provide the location in a document where the entity is referenced.	<input checked="" type="radio"/>	<input type="radio"/>

Question: 205

You are building a bot by using Microsoft Bot Framework.

You need to configure the bot to respond to spoken requests. The solution must minimize development effort.

What should you do?

- A. Deploy the bot to Azure and register the bot with a Direct Line Speech channel
- B. Integrate the bot with Cortana by using the Bot Framework SDK.
- C. Create an Azure function that will call the Speech service and connect the bot to the function.
- D. Deploy the bot to Azure and register the bot with a Microsoft Teams channel.

Direct Line Speech is a channel specifically designed for voice-enabled bot interactions.

Answer: B

Explanation:

Question: 206

You have a chatbot that was built by using Microsoft Bot Framework and deployed to Azure.

You need to configure the bot to support voice interactions. The solution must support multiple client apps.

Which type of channel should you use?

- Direct Line Speech is designed specifically for voice-enabled bot interactions.
It handles the complexities of speech-to-text and text-to-speech conversion, making it easier to build voice-first experiences**
- A. Cortana
 - B. Microsoft Teams
 - C. Direct Line Speech

It enables low-latency, high-reliability connections with client applications using the Speech SDK, optimized for conversational experiences.

Answer: C

Explanation:

Question: 207

You are developing a monitoring system that will analyze engine sensor data, such as rotation speed, angle, temperature, and pressure. The system must generate an alert in response to atypical values.

What should you include in the solution?

- A. Application Insights in Azure Monitor
- B. metric alerts in Azure Monitor
- C. Multivariate Anomaly Detection
- D. Univariate Anomaly Detection

For example, if pressure and temperature are normally related, a Multivariate Anomaly Detection system will likely detect it as an anomaly if both metrics are unusually high at the same time,

Answer: C

Explanation:

Question: 208**DRAG DROP**

You develop an app in O named App1 that performs speech-to-speech translation.

You need to configure App1 to translate English to German.

How should you complete the speechTranslationConfig object? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Values
addTargetLanguage
speechSynthesisLanguage
speechRecognitionLanguage
voiceName

Answer Area
var translationConfig = SpeechTranslationConfig.FromSubscription(SPEECH_SUBSCRIPTION_KEY, SPEECH_SERVICE_REGION); translationConfig. <input type="text"/> = "en-US"; translationConfig. <input type="text"/> ("de");

Explanation:

Values
addTargetLanguage
speechSynthesisLanguage
speechRecognitionLanguage
voiceName

Answer Area
var translationConfig = SpeechTranslationConfig.FromSubscription(SPEECH_SUBSCRIPTION_KEY, SPEECH_SERVICE_REGION); translationConfig. <input type="text"/> = "en-US"; translationConfig. <input type="text"/> ("de");

Question: 209

You train a Conversational Language Understanding model to understand the natural language input of users.

You need to evaluate the accuracy of the model before deploying it.

What are two methods you can use? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. From the language authoring REST endpoint, retrieve the model evaluation summary.
- B. From Language Studio, enable Active Learning, and then validate the utterances logged for review.
- C. From Language Studio, select Model performance.
- D. From the Azure portal, enable log collection in Log Analytics, and then analyze the logs.

Answer: A, C

Explanation:

Question: 210

You are building a Language Understanding solution.

You discover that many intents have similar utterances containing airport names or airport codes.

You need to minimize the number of utterances used to train the model.

Which type of custom entity should you use?

- A. Patterna.any
- B. machine-learning
- C. list
- D. regular expression

"Book a flight to {destination}"
where destination is a Pattern.any entity

list wont be best fit because need to add every airport start & destination
which is too many.

Answer: C

Explanation:

Question: 211

You use the Microsoft Bot Framework Composer to build a chatbot that enables users to purchase items.

You need to ensure that the users can cancel in-progress transactions. The solution must minimize development effort.

What should you add to the bot?

- A. a language generator
- B. a custom event
- C. a dialog trigger
- D. a conversation activity

just for sending
and receiving messages and for
managing conversation flow
interruptions

Dialog triggers allow you to interrupt the normal flow of a dialog and execute a specific set of actions when a certain condition is met. You can define an "Intent recognized" trigger to capture user utterances expressing a desire to cancel (e.g., "cancel," "abort," "stop")

Answer: D

Explanation:

Question: 212

You have an Azure subscription that contains an Anomaly Detector resource. You deploy a Docker host server named Server 1 to the on-premises network. You need to host an instance of the Anomaly Detector service on Server 1. Which parameter should you include in the docker run command?

- A. Fluentd
- B. Billing
- C. Http Proxy
- D. Mounts

It specifies the endpoint URI of your Anomaly Detector resource in Azure, enabling the container to connect to Azure for billing and usage tracking.

Answer: B

Explanation:

Question: 213

HOTSPOT

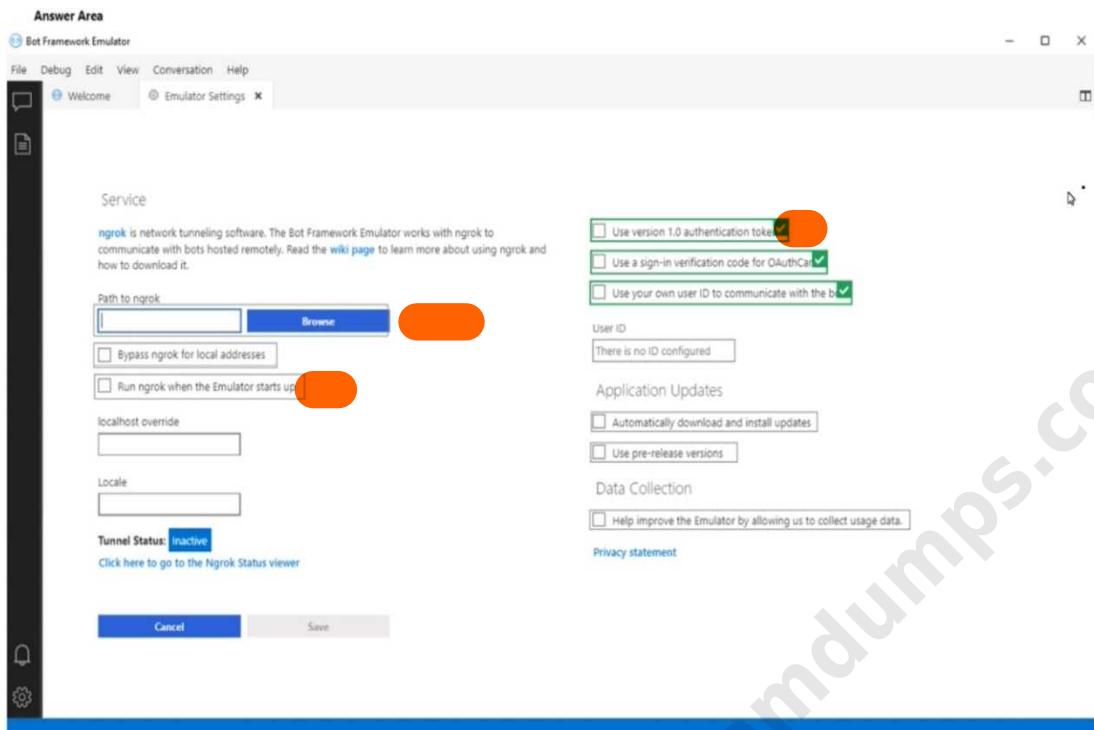
You have a chatbot.

You need to test the bot by using the Bot Framework Emulator. The solution must ensure that you are prompted for credentials when you sign in to the bot.

Which three settings should you configure? To answer, select the appropriate settings in the answer

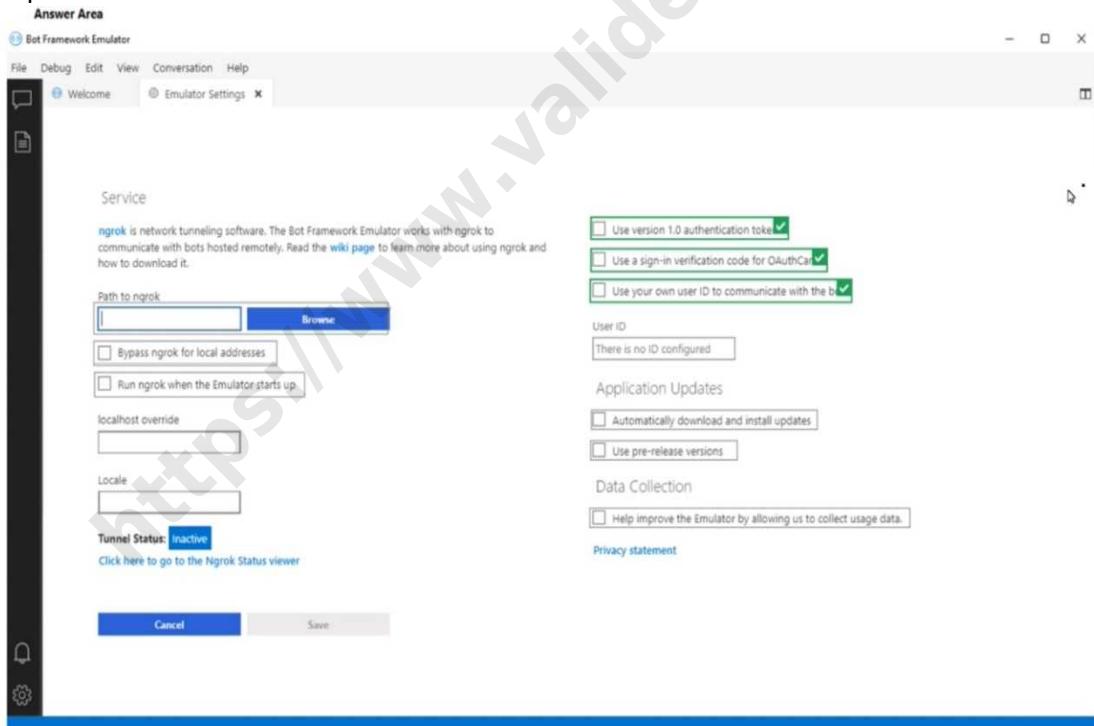
area.

NOTE Each correct selection is worth one point.



Answer:

Explanation:



Question: 214

You have an Azure Cognitive Search solution and an enrichment pipeline that performs Sentiment Analysis on social media posts.

You need to define a knowledge store that will include the social media posts and the Sentiment Analysis results.

Which two fields should you include in the definition? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- only for storageCOnnection String connection**
- A. storageContainer
 - B. tables
 - C. storaeetConnectionString
 - D. files
 - E. objects
- files-This field is used when you are storing files in the knowledge store.
This is not needed for storing structured results of sentiment analysis or the actual data of social media posts themselves.

social media posts stored as json contains various fields, object field to store raw social media posts.

Explanation: **ConnectionString -same as Storage Container** **Answer: D, E**
used to connect to storage account but not define fields of knowledge store.

Question: 215

HOTSPOT

You are building a call handling system that will receive calls from French-speaking and German-speaking callers. The system must perform the following tasks;

- Capture inbound voice messages as text.
- Replay messages in English on demand.

Which Azure Cognitive Services should you use? To answer, select the appropriate options in the answer area

a. NOTE: Each correct selection is worth one point.

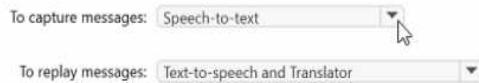
Answer Area



Answer:

Explanation:

Answer Area



Question: 216

HOTSPOT

You are building an app that will process incoming email and direct email messages to either French

or English language support teams.

Which Azure Cognitive Services API should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area



eastus.api.cognitive.microsoft.com

/text/analytics3.1/languages

Explanation: You would send the body of the incoming email to the /text/analytics3.1/languages endpoint

Answer Area



confirmed.

Answer:

Question: 217

You are building a Chatbot by using the Microsoft Bot Framework SDK. The bot will be used to accept food orders from customers and allow the customers to customize each food item. You need to configure the bot to ask the user for additional input based on the type of item ordered. The solution must minimize development effort. Which two types of dialogs should you use? Each correct answer presents part of the solution. NOTE: Each correct selection is worth one point.

- A. adaptive You create an adaptive dialog for the overall food ordering process
- B. action itself is not an "action" dialog. Actions are the tasks performed within a dialog step.
- C. waterfall execute sequentially, They don't have built in mechanisms for
- D. prompt interruptions and dynamic dialog selection.
- E. input

prompt -to ask the user specific questions about the **Answer: B, C**

Explanation: customization options for each food item
TextPrompt, ChoicePrompt, ConfirmPrompt, etc.

Question: 218

You are building an app that will include one million scanned magazine articles. Each article will be stored as an image file. You need to configure the app to extract text from the images. The solution must minimize development effort. What should you include in the solution?

- A. Computer Vision Image Analysis
- B. the Read API in Computer Vision
- C. Form Recognizer
- D. Azure Cognitive Service for Language

Form Recognizer is excellent for extracting text and structured data from forms and documents with predefined layouts. However, since magazine articles can have various layouts, Form Recognizer might not be the most suitable or efficient option for this use case.

The Read API (CV)- specifically designed for Optical Character Recognition (OCR) and is optimized for extracting text from images, including scanned documents.
It can handle a wide range of image formats and text orientations.

Answer: B

Question: 219

You have an Azure subscription that contains an AI enrichment pipeline in Azure Cognitive Search and an Azure Storage account that has 10 G8 of scanned documents and images. You need to index the documents and images in the storage account. The solution must minimize how long it takes to build the index. What should you do?

- A. From the Azure portal, configure parallel indexing.
- B. Create a text-based indexer by using the REST API.
- C. From the Azure portal, configure scheduled indexing.
- D. Configure field mappings by using the REST API.

Answer: A

Explanation:

The most efficient way to index a large amount of data like 10 GB of scanned documents and images in Azure Cognitive Search, while minimizing the build time, is by configuring parallel indexing

Question: 220

You are building a flight booking bot by using the Microsoft Bot Framework SDK.

The bot will ask users for their departure date. The bot must repeat the question until a valid date is given, or the users cancel the transaction.

Which type of dialog should you use?

- A. prompt
- B. action
- C. waterfall
- D. adaptive

needs to gather information step-by-step. They consist of a series of asynchronous steps that execute in order. This makes them perfect for guiding users through a multi-turn interaction, such as booking a flight.

Answer: A

Explanation:

Question: 221

You create five bots by using Microsoft Bot Framework Composer.

You need to make a single bot available to users that combines the bots. The solution must support dynamic routing to the bots based on user input.

Which three actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Create an Orchestrator model.
- B. Change the Recognizer/Dispatch type. to configure main bot to orchestra recognizer , by changing recognizer type you enable bot to leverage orchestrator routing capabilities.
- C. Create a composer extension.
- D. Enable WebSockets.
- E. Create a custom recognizer JSON file.
- F. Install the Orchestrator package.

orchestrate not enabled default so enable and install necessary components &libraries

Orchestrator is a language understanding model that helps route user input to the appropriate bot (or "skill") based on the detected intent. It acts as a central dispatcher

Answer: A, B, F

Explanation:

Question: 222**HOTSPOT**

You are building an app by using the Speech SDK. The app will translate speech from French to German by using natural language processing.

You need to define the source language and the output language.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area**Answer:**

Explanation:

Answer Area

```
var speechTranslationConfig =
    SpeechTranslationConfig.FromSubscription(speechKey, speechRegion);
speechTranslationConfig. SpeechRecognitionLanguage = "fr"
speech_translation_config. SpeechSynthesisLanguage = ("de")
```

SpeechRecognitinoLanguage

addTargetLanguage

Question: 223**DRAG DROP**

You have a collection of Microsoft Word documents and PowerPoint presentations in German.

You need to create a solution to translate the file to French. The solution must meet the following requirements:

- * Preserve the original formatting of the files.
- * Support the use of a custom glossary.

You create a blob container for German files and a blob container for French files. You upload the original files to the container for German files.

Which three actions should you perform in sequence actions from the list of actions to the answer area and arrange them in the correct order.

Actions
Perform an asynchronous translation by using the list of files to be translated.
Upload a glossary file to the container for German files.
Upload a glossary file to the container for French files.
Generate a list of files to be translated.
Define a document translation specification that has a French target.
Perform an asynchronous translation by using the document translation specification.

Answer Area



- 1.upload a glossary file to the container for german files
- 2.define a document translation specification that has a french target.

Answer:

Explanation:

Actions
Perform an asynchronous translation by using the list of files to be translated.
Upload a glossary file to the container for German files.
Upload a glossary file to the container for French files.

Answer Area



- 1 Generate a list of files to be translated.
- 2 Define a document translation specification that has a French target.
- 3 Perform an asynchronous translation by using the document translation specification.

Question: 224

DRAG DROP

You have an app that manages feedback.

You need to ensure that the app can detect negative comments by using the Sentiment Analysis API in Azure Cognitive Service for Language. The solution must ensure that the managed feedback remains on your company's internal network.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

NOTE: More than one order of answer choices is correct You will receive credit for any of the correct orders you select.

Actions
Provision the Language service resource in Azure.
Deploy a Docker container to an Azure container instance.
Deploy a Docker container to an on-premises server.
Identify the Language service endpoint URL and query the prediction endpoint.
Run the container and query the prediction endpoint.

Answer Area



Explanation:

Provision the Language service resource in Azure.

Deploy a Docker container to an on-premises server.

Run the container and query the prediction endpoint.

- 1.Provision Language service resource in azure.
- 2.Deploy the docker container to an on-premises server
- 3.run the container and query prediction endpoint

According to the [Microsoft documentation](#), the Language service is a cloud-based service that provides various natural language processing features, such as sentiment analysis, key phrase extraction, named entity recognition, etc. You can provision the Language service resource in Azure by following the steps in [Create a Language resource](#). You will need to provide a name, a subscription, a resource group, a region, and a pricing tier for your resource. You will also get a key and an endpoint for your resource, which you will use to authenticate your requests to the Language service API.

According to the [Microsoft documentation](#), you can also use the Language service as a container on your own premises or in another cloud. This option gives you more control over your data and network, and allows you to use the Language service without an internet connection. You can deploy a Docker container to an on-premises server by following the steps in [Deploy Language containers](#). You will need to have Docker installed on your server, pull the container image from the Microsoft Container Registry, and run the container with the appropriate parameters. You will also need to activate your container with your key and endpoint from your Azure resource.

According to the [Microsoft documentation](#), once you have deployed and activated your container, you can run it and query the prediction endpoint to get sentiment analysis results. The prediction endpoint is a local URL that follows this format: `http://<container IP address>:<port>/text/analytics/v3.1-preview.4/sentiment`. You can send HTTP POST requests to this endpoint with your text input in JSON format, and receive JSON responses with sentiment labels and scores for each document and sentence in your input.

Question: 225

You are building an app that will share user images.

You need to configure the app to perform the following actions when a user uploads an image:

- Categorize the image as either a photograph or a drawing.
- Generate a caption for the image.

The solution must minimize development effort.

Which two services should you include in the solution? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. object detection in Computer Vision
- B. content tags in Computer Vision
- C. image descriptions in Computer Vision
- D. image type detection in Computer Vision
- E. image classification in Custom Vision

image desc-service automatically generates a human-readable caption that describes the content of an image

image type -This service categorizes an image by its type, such as a photograph or drawing.

This fulfills the requirement to categorize the image as either a photograph or a drawing

Answer: CD

Explanation: content tag- helpful for image organization, these tags don't directly address the need for a descriptive caption

According to the [Microsoft documentation](#), Computer Vision is a cloud-based service that provides developers with access to advanced algorithms for processing images and returning information. By uploading an image or specifying an image URL, Computer Vision algorithms can analyze visual content in different ways based on inputs and user choices.

According to the [Microsoft documentation](#), image type detection is one of the features of Computer Vision that can categorize an image as either a photograph or a drawing. You can use the image type detection feature by calling the Analyze Image API with the visualFeatures parameter set to ImageType. The API will return a JSON response with an imageType field that indicates whether the image is a photo or a clipart.

According to the [Microsoft documentation](#), image descriptions is another feature of Computer Vision that can generate a caption for an image. You can use the image descriptions feature by calling the Analyze Image API with the visualFeatures parameter set to Description. The API will return a JSON response with a description field that contains a list of captions for the image, each with a confidence score.

Therefore, by using these two features of Computer Vision, you can achieve your app requirements

with minimal development effort. You don't need to use any other services, such as object detection, content tags, or Custom Vision, which are designed for different purposes.

Question: 226

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have a chatbot that uses question answering in Azure Cognitive Service for Language

Users report that the responses of the chatbot lack formality when answering spurious questions

You need to ensure that the chatbot provides formal responses to spurious questions.

Solution: From Language Studio, you change the chitchat source to **qna_chitchat_professional.tsv**, and then retrain and republish the model.

Does this meet the goal?

- A. Yes
- B. No

formal response, the correct dataset would be qna_chitchat_professional.tsv

Answer: B

Explanation:

Question: 227

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have a chatbot that uses question answering in Azure Cognitive Service for Language

Users report that the responses of the chatbot lack formality when answering spurious questions

You need to ensure that the chatbot provides formal responses to spurious questions.

Solution: From Language Studio, you change the chitchat source to **qna_chitchit_friendly.tsv**, and then retrain and republish the model.

Does this meet the goal?

only different so NO

- A. Yes
- B. No

The qna_chitchat_friendly.tsv file is designed to provide friendly and informal responses for chit-chat

Answer: B

Explanation:

Question: 228

Note: This question is part of a series of questions that present the same scenario. Each question in

the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have a chatbot that uses question answering in Azure Cognitive Service for Language

Users report that the responses of the chatbot lack formality when answering spurious questions

You need to ensure that the chatbot provides formal responses to spurious questions.

Solution: From Language Studio, you remove all the chit-chat question and answer pairs, and then retrain and republish the model

Does this meet the goal?

A. Yes

B. No

Removing all chit-chat question and answer pairs
will prevent the chatbot from engaging in any chit-chat,
including responding formally to spurious questions

Answer: B

Explanation:

Removing all the chit-chat question and answer pairs from the project will not ensure that the chatbot provides formal responses to spurious questions. It will only make the chatbot unable to handle any chit-chat scenarios, which may result in a poor user experience and a loss of engagement. Instead, you should choose a chit-chat personality that matches the tone and style of your chatbot, such as Professional or Caring. [You can also edit the chit-chat questions and answers to suit your specific needs, or add new ones that are not in the predefined data set12](#). This way, you can ensure that the chatbot responds appropriately to spurious questions, while still maintaining a conversational and engaging interaction with the user.

Question: 229

HOTSPOT

You are building an app that will share user images.

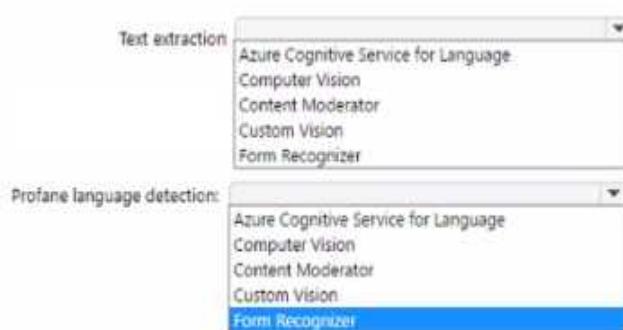
You need to configure the app to meet the following requirement's

- uploaded images must be scanned and any text must be extracted from the images.
- Extracted text must be analyzed for the presence of profane language.
- The solution must minimize development effort.

what should you use for each requirement? To answer, select the apocopate opinions in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area



Answer:

Explanation:

Text Extraction: b) Computer Vision

Computer Vision is a service that can analyze images and extract text from them using optical character recognition (OCR) or read API1. OCR can detect and extract printed or handwritten text from images in various languages and formats, such as PDF, TIFF, or JPEG2. Read API can perform asynchronous batch processing of text from multiple images, and can also handle text in tables, mixed languages, and rotated text3.

Profane Language Detection: c) Content Moderator

"profane" generally means offensive.

Content Moderator is a service that can detect and filter out potentially offensive or unwanted content in text, images, and videos. For text, Content Moderator can identify profane or abusive language, personal data, and custom terms that you define. You can also use Content Moderator to review and moderate the content manually or automatically using workflows.

Question: 230

You are building a solution that will detect anomalies in sensor data from the previous 24 hours.

You need to ensure that the solution scans the entire dataset, at the same time, for anomalies.

Which type of detection should you use?

- A. batch
- B. streaming
- C. change point

Answer: A

Explanation:

Batch anomaly detection is a type of anomaly detection that scans the entire dataset at once for outliers and unusual patterns. Batch anomaly detection is suitable for offline analysis of historical data, such as sensor data from the previous 24 hours. [Batch anomaly detection can use various techniques, such as statistical methods, machine learning methods, or hybrid methods, to identify anomalies in the data123.](#)

Question: 231

You have a 20-GB file named file1. avi that is stored on a local drive.

You need to index file1. avi by using the Azure Video indexer website.

What should you do first?

- A. Upload File1. avi to an Azure storage queue. [only for message queueing](#)
- B. upload File1. avi to the [www.youtube.com](#) seepage [only for video sharing.](#)
- C. Upload file1. avi to the Azure video indexer website.
- D. Upload file1. avi to Microsoft OneDrive. [file storage and synchronization service does not provide indexing](#)

Answer: D

Explanation:

[This is because the Azure Video Indexer website allows you to upload videos from a URL or from your file system, but there are some limitations and considerations for each option1.](#)

Type text here

If you upload from your file system, the size of the file is limited to 2 GB, which is less than the size of file1.avi (20 GB). Therefore, this option is not feasible.

If you upload from a URL, the size of the file is limited to 30 GB, which is enough for file1.avi. However, the URL must be publicly accessible and valid, and the file must be accessible. [You cannot use URLs from streaming services such as YouTube1](#). Therefore, options A and B are not valid.

The best option is to upload file1.avi to a cloud storage service such as Microsoft OneDrive, and then generate a shareable link for the file. You can then paste the link in the Azure Video Indexer website and upload the video from the URL. [This way, you can avoid the file size limitation and ensure that the file is accessible and valid2](#).

Question: 232

DRAG DROP

You Build a bot in JavaScript.

From the Azure Command-Line interface (CLI), you run the following command.

az bot prepare-deploy

You need to deploy the bot to Azure.

Which three Azure CLI commands should you run in sequence? To answer, move the appropriate commands from the list of commands to the answer area and arrange them in the client order.

Commands	Answer Area
az ad app credential	
az ad app create	
az deployment group create	
az webapp deployment source config-zip	
az ad app update	



Answer:

Explanation:

To deploy the bot to Azure, you should run the following three Azure CLI commands in sequence:

az deployment group create This command will create the Azure resources for your bot using an ARM template and a parameters file. You need to specify the resource group name, the template file path, and the parameters file path. For example:

az deployment group create --resource-group myResourceGroup --template-file "deploymentTemplates\template-with-preexisting-rg.json" --parameters "deploymentTemplates\parameters-for-template-BotApp-with-rg.json"

[This command will also output the app ID and password of your bot, which you will need for the next command1.](#)

az webapp deployment source config-zip This command will deploy your bot code to the app service that you created in the previous step. You need to specify the resource group name, the app service name, and the zip file path of your bot code. For example:

az webapp deployment source config-zip --resource-group myResourceGroup --name myBotAppService --src "code.zip"

[This command will also output the URL of your bot endpoint, which you will need for the next command2.](#)

az ad app update This command will update your bot registration with the endpoint URL of your bot. You need to specify the app ID of your bot and the endpoint URL. For example:

az ad app update --id myBotAppId --set replyUrls="https://myBotAppService.azurewebsites.net/api/messages"

[This command will complete the deployment process and make your bot ready to be tested3.](#)

Question: 233

You are building a chatbot.

You need to ensure that the bot will recognize the names of your company's products and codenames. The solution must minimize development effort.

Which Azure Cognitive Service for Language service should you include in the solution?

- A. custom text classification
- B. entity linking
- C. custom Named Entity Recognition (NER)
- D. key phrase extraction

duplicate

Explanation:

Answer: C

Question: 234

You are developing an app that will use the Decision and Language APIs.

You need to provision resources for the app. The solution must ensure that each service is accessed by using a single endpoint and credential.

Which type of resource should you create?

- A. Language
- B. Speech
- C. Azure Cognitive Services
- D. Content Moderator

unified resource that allows you to access multiple cognitive services, including Language and Decision APIs, using a single endpoint and API key.

Answer: C

Explanation:

Question: 235

You have an Azure subscription that contains an Azure App Service app named App1.

You provision a multi-service Azure Cognitive Services resource named CSAccount1.

You need to configure App1 to access CSAccount1. The solution must minimize administrative effort.

What should you use to configure App1?

- A. a system-assigned managed identity and an X.509 certificate
- B. the endpoint URI and an OAuth token complex auth flow and also adds administrative complexity
- C. the endpoint URI and a shared access signature (SAS) token commonly used for storage accounts and not recommended for Cognitive Services
- D. the endpoint URI and subscription key

Answer: C

Explanation:

Question: 236

You have an Azure Cognitive Services model named Model that identifies the intent of text input.

You develop an app in C# named App1.

You need to configure App1 to use Model1.

Which package should you add to App1?

A. Azure.AI.Language.Conversations

B. SpeechServicesToolkit **used for speech-related tasks, such as speech recognition and synthesis**

C. Universal.Microsoft.CognitiveServices.Speech **This package is also related to speech services**

D. Xamarin.Cognitive.Speech **and is not relevant for intent recognition from text**

package is specific to Xamarin development for mobile applications and is not the correct package

Answer: C

Explanation:

Question: 237

You create a bot by using the Microsoft Bot Framework SDK.

You need to configure the bot to respond to events by using custom text responses.

What should you use?

A. an adaptive card

B. an activity handler

C. a dialog

D. a skill

Answer: B

Explanation:

Question: 238

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have a chatbot that uses question answering in Azure Cognitive Service for Language.

Users report that the responses of the chatbot lack formality when answering random questions that are outside the scope of the knowledge base.

You need to ensure that the chatbot provides formal responses to these spurious questions.

Solution: From Language Studio, you modify the question and answer pairs for the custom intents, and then retrain and republish the model.

Does this meet the goal?

A. Yes

Use a Formal Chit-Chat Model designed for formalqna_chitchat_professional.tsv

B. No

Define a Formal Fallback Response

Custom Knowledge Base

Answer: B

Explanation:

Question: 239

DRAG DROP

You have a Docker host named Host1 that contains a container base image.

You have an Azure subscription that contains a custom speech-to-text model named model1.

You need to run model1 on Host1.

model already exist so retrain the model not required.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions
Configure disk logging.
Export model1 to Host1.
Request approval to run the container.
Retrain the model.
Run the container.

Answer Area



1. Request approval to run the container
2. Export model1 to Host1
3. Run the container

environments where there's a formal approval process for running containers (due to security, compliance, or resource management concerns)

Answer:

verified

Explanation:

According to the course AI-102T00: Designing and Implementing a Microsoft Azure AI Solution1, the correct sequence of actions to run model1 on Host1 is:

Export model1 to Host1

Before you can run the model, it needs to be available on the host machine

Run the container

you can run a Docker container using the appropriate base image and configure it to use the exported model persist logs from the container

Configure disk logging

The explanation and references are as follows:

Export model1 to Host1: This step is required to deploy the custom speech-to-text model to the Docker host. [You can use the Azure portal or the Azure CLI to export the model as a container image2.](#)

Run the container: This step is required to start the container and run the model on the Docker

host. [You can use the Docker CLI to run the container image3.](#)

Configure disk logging: This step is optional but recommended to monitor the performance and health of the container. [You can use the Docker CLI to configure disk logging for the container4.](#)

Question: 240

DRAG DROP

You are building a transcription service for technical podcasts.

Testing reveals that the service fails to transcribe technical terms accurately.

You need to improve the accuracy of the service.

Which five actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

NOTE: More than one order of answer choices is correct. You will receive credit for any of the correct orders you select.

Actions
Create a Speaker Recognition model.
Create a Conversational Language Understanding model.
Create a Custom Speech project.
Create a speech-to-text model.
Upload training datasets.
Train the model.
Deploy the model.

Answer Area



1. Create Custom Speech project
2. Create speech-to-text model
3. Upload training datasets
4. Train model
5. Deploy model

Answer:

Explanation:

Actions
Create a Speaker Recognition model.
Create a Conversational Language Understanding model.

Answer Area



1. Create a Custom Speech project.
2. Create a speech-to-text model.
3. Upload training datasets.
4. Train the model.
5. Deploy the model.

<https://learn.microsoft.com/en-us/azure/cognitive-services/speech-service/custom-speech-overview#how-does-it-work>

With Custom Speech, you can upload your own data, test and train a custom model, compare accuracy between models, and deploy a model to a custom endpoint.

- Create a project and choose a model. Use a Speech resource that you create in the Azure portal. If you will train a custom model with audio data, choose a Speech resource region with dedicated hardware for training audio data.

- Upload test dat

a. Upload test data to evaluate the speech to text offering for your applications, tools, and products.

- Train a model. Provide written transcripts and related text, along with the corresponding audio data. Testing a model before and after training is optional but recommended.

- Deploy a model. Once you're satisfied with the test results, deploy the model to a custom endpoint. With the exception of batch transcription, you must deploy a custom endpoint to use a Custom Speech model.

Question: 241

You have a file share that contains 5,000 images of scanned invoices.

You need to analyze the images. The solution must extract the following data:

- Invoice items
- Sales amounts
- Customer details

What should you use?

A. Custom Vision

B. Computer Vision computer vision-image classification and object detection. It is not intended for extracting structured data from forms or invoices

C. Immersive Reader -designed to enhance the reading experience by providing features like text-to-speech, line focus, and translation. It is not relevant for analyzing images or extracting data from documents.

D. Form Recognizer Form Recognizer is the best service to use for extracting structured data from scanned invoices, such as invoice items, sales amounts,

and customer details. It's purpose-built for document analysis, making it the most efficient and accurate solution.

Answer: C

Explanation:

Question: 242

You have an Azure Cognitive Search solution and a collection of blog posts that include a category

field. You need to index the posts. The solution must meet the following requirements:

- Include the category field in the search results.
- Ensure that users can search for words in the category field.
- Ensure that users can perform drill down filtering based on category.

Which index attributes should you configure for the category field?

facetable- drill-down filtering based on the categories

- A. searchable, facetable, and retrievable
- B. retrievable, filterable, and sortable sortable is not good for drill down requirement
- C. retrievable, facetable, and key key - only for unique identifier
- D. searchable, sortable, and retrievable

Answer: B

Explanation:

Question: 243

You have an Azure subscription that contains a multi-service Azure Cognitive Services Translator resource named Translator1.

You are building an app that will translate text and documents by using Translator1.

You need to create the REST API request for the app.

Which headers should you include in the request?

region to direct to correct azure region when using regional oppose to global one)

- A. the subscription key and the client trace ID
- B. the subscription key, the subscription region, and the content type
 subscription key essential for authenticating with cognitive services translation service
- C. the resource ID and the content language for translation accepted one are applicaton/json,text/plain depends on how text is sent.
- D. the access control request, the content type, and the content length

Answer: B

Explanation:

Question: 244

You are building an app that will use the Azure Video Indexer service.

You plan to train a language model to recognize industry-specific terms.

You need to upload a file that contains the industry-specific terms.

Which file format should you use?

- A. PDF **PDFs are document files designed for displaying content, not for training language models.**
- B. XML **XML is used for structuring data and would be unsuitable for providing the raw text data needed for language model training in Video Indexer.**
- C. TXT **custom language model training feature primarily utilizes plain text files (.txt) for adaptation text.**
- D. XLS **XLS (or XLSX) files are spreadsheet files. They are not the appropriate format for uploading industry-specific terms for training a language model in Video Indexer.**

Answer: C

Explanation:

Question: 245

HOTSPOT

You are building an Azure web app named App1 that will translate text from English to Spanish.

You need to use the Text Translation REST API to perform the translation. The solution must ensure that you have data sovereignty in the United States.

How should you complete the URI? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Explanation:

Answer:

- api.cognitive.microsofttranslator.com
global endpoint nam is for north america
1. api-nam.cognitive.microsofttranslator.com
 2. translate

<https://learn.microsoft.com/en-us/azure/cognitive-services/Translator/reference/v3-0-reference#base-urls>

Requests to Translator are, in most cases, handled by the datacenter that is closest to where the request originated. If there's a datacenter failure when using the global endpoint, the request may be routed outside of the geography.

To force the request to be handled within a specific geography, use the desired geographical endpoint. All requests are processed among the datacenters within the geography.

- United States

api-nam.cognitive.microsofttranslator.com

<https://learn.microsoft.com/en-us/azure/cognitive-services/translator/reference/rest-api-guide>

- translate

Translate specified source language text into the target language text.

Question: 246

HOTSPOT

You have an app named App1 that uses Azure AI Document Intelligence to analyze medical You have an app named App1 that uses Azure AI Document Intelligence for patients. You send a request to App1 and receive the following response.

```
{  
    "status": "succeeded",  
    "createdDateTime": "2023-09-14T21:01:02Z",  
    "lastUpdatedDateTime": "2023-09-14T21:01:03Z",  
    "analyzeResult": {  
        "apiVersion": "2023-07-31",  
        "modelId": "prebuilt-healthInsuranceCard.us",  
        "stringIndexType": "utf16CodeUnit",  
        "content": "Blood Pressure 118/72",  
        "pages": [  
            {  
                ...  
                "words": [  
                    {  
                        "content": "Blood",  
                        "polygon": [ ... ],  
                        "confidence": 0.766,  
                        "span": { ... }  
                    },  
                    {  
                        "content": "Pressure",  
                        "polygon": [ ... ],  
                        "confidence": 0.716,  
                        "span": { ... }  
                    },  
                    {  
                        "content": "118/72",  
                        "polygon": [ ... ],  
                        "confidence": 0.761,  
                        "span": { ... }  
                    }  
                ]  
            ],  
            ...  
            "documents": [  
                {  
                    "docType": "healthInsuranceCard.us",  
                    "boundingRegions": [ ... ]  
                }  
            ],  
            "fields": {},  
            "confidence": 1,  
            "spans": [ ... ]  
        }  
    }  
}
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth point.

Answer Area

Statements	Yes	No
The chosen model is suitable for the intended use case.	<input type="radio"/>	<input checked="" type="radio"/>
The text content was recognized with greater than 70 percent confidence.	<input checked="" type="radio"/>	<input type="radio"/>
The form elements were recognized with greater than 70 percent confidence.	<input type="radio"/>	<input checked="" type="radio"/>

No
Yes
No

Answer:

Explanation:

Answer Area

Statements	Yes	No
The chosen model is suitable for the intended use case.	<input type="radio"/>	<input checked="" type="radio"/>
The text content was recognized with greater than 70 percent confidence.	<input checked="" type="radio"/>	<input type="radio"/>
The form elements were recognized with greater than 70 percent confidence.	<input type="radio"/>	<input checked="" type="radio"/>

Question: 247

HOTSPOT

You plan to deploy an Azure OpenAI resource by using an Azure Resource Manager (ARM) template.

You need to ensure that the resource can respond to 600 requests per minute.

How should you complete the template? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Azure OpenAI allows you to manage how frequently your application can make inferencing requests. Your rate limits are based on Tokens-per-Minute (TPM). For example, if you have a capacity of 1, this equals 1,000 TPM, and the rate limit of requests you can make per minute (RPM) is calculated using a ratio. For every 1,000 TPM, you can make 6 RPM.

Answer Area

```
{
  "type": "Microsoft.CognitiveServices/accounts/deployments",
  "apiVersion": "2023-05-01",
  "name": "arm-aoai-sample-resource/arm-je-std-deployment",
  "dependsOn": [
    "[resourceId('Microsoft.CognitiveServices/accounts', 'arm-aoai-sample-resource')]"
  ],
  "sku": {
    "name": "Standard",
    "capacity": 600
  },
  "properties": {
    "model": {
      "format": "OpenAI",
      ...
    }
  }
}
```

capacity 100

So, if you need to process 600 requests every minute, you'll require a TPM that supports that many RPM. Using the ratio, for 600 RPM, you need 100,000 TPM (because 600 divided by 6 equals 100, and 100 multiplied by 1,000 equals 100,000). In this scenario, you would set the capacity to 100, since each capacity unit equals 1,000 TPM.

Explanation:

Answer:

Question: 248

You have an Azure subscription. The subscription contains an Azure OpenAI resource that hosts a GPT-4 model named Model1 and an app named App1. App1 uses Model1.

You need to ensure that App1 will NOT return answers that include hate speech.

What should you configure for Model1?

Repeated Attempts to jailbreak the system like leetcode violations or repeatedly violating rules

- A. the Frequency penalty parameter
- B. abuse monitoring
- C. a content filter
- D. the Temperature parameter

discourages the model from repeating the same tokens or phrases

Abuse Monitoring -instances of recurring content or behaviors that might violate the service's Code of Conduct. It's a crucial aspect of responsible AI, but it doesn't prevent hate speech in real-time. Instead, it focuses on identifying patterns of abuse and taking actions like notifying the customer or suspending the service.

It's a reactive measure, not a preventive one

randomness and creativity of the model's output

Answer: B

Explanation:

Question: 249

You have an Azure subscription. The subscription contains an Azure OpenAI resource that hosts a GPT-3.5 Turbo model named Model1.

You configure Model1 to use the following system message: "You are an AI assistant that helps people solve mathematical puzzles. Explain your answers as if the request is by a 4-year-old."

Which type of prompt engineering technique is this an example of?

A. few-shot learning

the model with information to shape its understanding of the task and its style of response.

B. affordance

Instead of simply asking, "Write a story about a dragon," you could prime the model with, "Write a fantasy story about a dragon who lives in a crystal cave and befriends a young wizard". This primes the model to generate a story within a specific genre and setting

C. chain of thought

D. priming

Answer: D

Explanation:

Question: 250

HOTSPOT

You build a chatbot by using Azure OpenAI Studio.

You need to ensure that the responses are more deterministic and less creative.

Which two parameters should you configure? To answer, select the appropriate parameters in the answer area.

NOTE: Each correct answer is worth one point.

Imagine you have a bag of marbles, each representing a word, and the size of each marble represents its probability. Top P is like picking marbles from the bag, starting with the largest (most probable) and continuing to pick smaller marbles until the total size of the marbles you've picked reaches a certain threshold (P)

Answer Area
Answer:

Explanation:

Answer Area

Temperature

Top P

high Top P(close to 1)
less probable ones

Low Top P(close to 0)
most probable ones

1. Open the Custom Vision Project from portal
Top p - Unlike Top K (which selects from a fixed number of words),
Top P allows the vocabulary size to change based on the probability distribution.

Question: 251

You are building a chatbot for a travel agent. The chatbot will use the Azure OpenAI GPT 3.5 model and will be used to make travel reservations.

You need to maximize the accuracy of the responses from the chatbot.

What should you do?

The chatbot needs real-time, accurate information about flights, hotels, availability, and pricing from the travel agent's database to be useful. Without this connection, it's just guessing.

- A. Configure the model to include data from the travel agent's database.
important step than others
- B. Set the Top P parameter for the model to 0.
always selecting the most probable token.
However, "most probable" doesn't equal "most accurate"
- C. Set the Temperature parameter for the model to 0.
- D. Modify the system message used by the model to specify that the answers must be accurate.

Answer: A

Explanation:

Question: 252

HOTSPOT

You are developing a text processing solution.

You have the function shown below.

```
static void GetKeyWords(TextAnalyticsClient textAnalyticsClient, string text)
{
    var response = textAnalyticsClient.RecognizeEntities(text);
    Console.WriteLine("Key words:");

    foreach (CategorizedEntity entity in response.Value)
    {
        Console.WriteLine($"\\t{entity.Text}");
    }
}
```

For the second argument, you call the function and specify the following string.

Our tour of Paris included a visit to the Eiffel Tower.

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area**Statements**

The output will include the following words: our and included.

Yes	No
<input type="radio"/>	<input checked="" type="radio"/>

The output will include the following words: Paris, Eiffel, and Tower.

<input checked="" type="radio"/>	<input type="radio"/>
----------------------------------	-----------------------

The function will output all the key phrases from the input string to the console.

<input type="radio"/>	<input checked="" type="radio"/>
-----------------------	----------------------------------

Duplicate

Answer:

Explanation:

Answer Area**Statements**

The output will include the following words: our and included.

Yes	No
<input checked="" type="radio"/>	<input type="radio"/>

The output will include the following words: Paris, Eiffel, and Tower.

<input type="radio"/>	<input checked="" type="radio"/>
-----------------------	----------------------------------

The function will output all the key phrases from the input string to the console.

<input type="radio"/>	<input checked="" type="radio"/>
-----------------------	----------------------------------

Question: 253

HOTSPOT

You have an Azure subscription that contains an Azure AI Document Intelligence resource named DM.

You build an app named App1 that analyzes **PDF files for handwritten content** by using DM.

You need to ensure that App1 will recognize the handwritten **content**.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

simple document with just text, you'd use **prebuilt-read**.

document containing tables and text, like an invoice, you'd use **prebuilt-layout**.

Answer Area

```
Uri fileUri = new Uri("<fileUri>");  
AnalyzeDocumentOperation operation = await client.AnalyzeDocumentFromUriAsync(WaitUntil.Completed,  
    AnalyzeResult result = operation.Value;  
    foreach (DocumentStyle style in result.Styles)  
    {  
        bool isHandwritten = style.IsHandwritten.HasValue && style.IsHandwritten == true;  
        if (isHandwritten && style.Confidence > 0.75)  
        {  
            Console.WriteLine($"Handwritten content found:");  
            foreach (DocumentSpan span in style.Spans)  
            {  
                Console.WriteLine($" Content: {result.Content.Substring(span.Index, span.Length)}");  
            }  
        }  
    }  
}
```

0.75 prebuilt-read for text use prebuilt read

Explanation:**Answer Area**

```
Uri fileUri = new Uri("<fileUri>");  
AnalyzeDocumentOperation operation = await client.AnalyzeDocumentFromUriAsync(WaitUntil.Completed, "prebuilt-document", fileUri);  
AnalyzeResult result = operation.Value;  
foreach (DocumentStyle style in result.Styles)  
{  
    bool isHandwritten = style.IsHandwritten.HasValue && style.IsHandwritten == true;  
    if (isHandwritten && style.Confidence > 0.75)  
    {  
        Console.WriteLine($"Handwritten content found:");  
        foreach (DocumentSpan span in style.Spans)  
        {  
            Console.WriteLine($" Content: {result.Content.Substring(span.Index, span.Length)}");  
        }  
    }  
}
```

Question: 254

You build a chatbot that uses the Azure OpenAI GPT 3.5 model.

You need to improve the quality of the responses from the chatbot. The solution must minimize development effort.

What are two ways to achieve the goal? Each correct answer presents a complete solution.

NOTE: Each correct answer is worth one point.

- A. Fine-tune the model.

- B. Provide grounding content.
- C. Add sample request/response pairs.
- D. Retrain the language model by using your own data.
- E. Train a custom large language model (LLM).

Provide grounding content: This involves feeding the chatbot with relevant domain-specific information and data. This can include documents, articles, FAQs, or any other content related to the chatbot's purpose. By providing this context, the chatbot can better understand the user's intent and respond in a more relevant and informative way.

Answer: B, C

Explanation:

Question: 255

HOTSPOT

Add sample request/response pairs: This technique involves providing the chatbot with a set of pre-defined questions and their corresponding answers. This helps the chatbot learn the conversation patterns and phrasing related to its specific domain. By analyzing these examples, the chatbot can improve its ability to generate natural and consistent responses to user queries.

You have an Azure subscription that contains an Azure OpenAI resource named AH.

You build a chatbot that will use AI1 to provide generative answers to specific questions.

You need to ensure that the responses are more creative and less deterministic.

How should you complete the code? To answer, select the appropriate options in the answer area.


NOTE: Each correct selection is worth one point.

Answer Area

```
new ChatCompletionsOptions()
{
    Messages =
    {
        new ChatMessage(
            ChatRole.User, "", ""),
        ChatRole.Assistant,
        ChatRole.Function,
        ChatRole.System,
        ChatRole.User
    },
    Temperature = (float)1.0,
    PresencePenalty,
    Temperature,
    TokenSelectionBiases
};

MaxTokens = 800,
});
```

User
Temperature



Answer:

Explanation:

Answer Area

```
new ChatCompletionsOptions()
{
    Messages =
    {
        new ChatMessage( ChatRole.User , @"" ),
    },
    Temperature = (float)1.0,
    MaxTokens = 800,
});
```

Question: 256

DRAG DROP

You have an Azure subscription that contains an Azure OpenA1 resource named AH.

You plan to build an app named App1 that will write press releases by using AM.

You need to deploy an Azure OpenA1 model for App1. The solution must minimize development effort.

Which three actions should you perform in sequence in Azure OpenAI Studio? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions	Answer Area
Create a deployment that uses the text-embedding-ada-002 model.	
Apply the Default system message template.	
Create a deployment that uses the GPT-35 Turbo model.	
Apply the Marketing Writing Assistant system message template.	
Deploy the solution to a new web app.	

Drag and drop arrows: > (right), < (left), ↑ (up), ↓ (down).

Answer:

Explanation:

Actions	Answer Area
Create a deployment that uses the text-embedding-ada-002 model.	1 Create a deployment that uses the GPT-35 Turbo model.
Apply the Default system message template.	2 Apply the Marketing Writing Assistant system message template.
	3 Deploy the solution to a new web app.

Drag and drop arrows: > (right), < (left), ↑ (up), ↓ (down).

Question: 257

HOTSPOT

A1 You have an Azure OpenAI resource named AH that hosts three deployments of the GPT 3.5 model. Each deployment is optimized for a unique workload.

You plan to deploy three apps. Each app will access AM by using the REST API and will use the deployment that was optimized for the apps intended workload.

You need to provide each app with access to AH and the appropriate deployment. The solution must ensure that only the apps can access AM.

What should you use to provide access to AM. and what should each app use to connect to its appropriate deployment? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Provide access to AI1 by using:

An API key
 A bearer token
 A shared access signature (SAS) token

Connect to the deployment by using:

A deployment endpoint
 An API key
 A deployment name
 A deployment type

Explanation:

Answer Area

Provide access to AI1 by using:

An API key

Connect to the deployment by using:

A deployment endpoint

Answer:

Question: 258

HOTSPOT

You are building a chatbot.

You need to use the Content Moderator API to identify aggressive and sexually explicit language.

Which three settings should you configure? To answer, select the appropriate settings in the answer area.

NOTE: Each correct selection is worth one point.

Microsoft Cognitive Services

▶ Image

▶ Text

classify

Content Moderator - Moderate

Text - Screen

The operation detects profanity in more than 100 languages and match against custom and shared blacklists.

Host

Type text here

Name: uksouth.api.cognitive.micros

Query parameters

autocorrect	Value	X Remove parameter
PII	Value	X Remove parameter
listId	Value	X Remove parameter
classify	false	X Remove parameter
language	Value	X Remove parameter

+ Add parameter

Headers

Content-Type	text/plain	X Remove header
Ocp-Apim-Subscription-Key	Value	?

+ Add header

Explanation:

Answer:

Answer Area**Content Moderator - Moderate****Text - Screen**

The operation detects profanity in more than 100 languages and match against custom and shared blacklists.

Host**Name**

Resource Name

Query parameters**autocorrect**

✖ Remove parameter
PII

✖ Remove parameter
listId

✖ Remove parameter
classify

✖ Remove parameter
language

✖ Remove parameter
+ Add parameter

autocorrect- Improves text quality and can contribute to better analysis.

1.classify

2.laguage

3.ocp-apim-subscription-key

listid-Enables customization and fine-tuning of moderation based on specific needs.

classify-The primary driver for detecting potentially offensive content.

language- Essential for language-specific accuracy, especially for profanity detection.

ocp-apim-subscription-key - authentication key for accessing contentmoderation api

Question: 259**DRAG DROP**

You have a monitoring solution that uses the Azure AI Anomaly Detector service.

You provision a server named Server1 that has intermittent internet access.

You need to deploy the Azure AI Anomaly Detector to Server 1.

Which four actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions
From Server1, run the docker push command.
Query the prediction endpoint of the Azure AI Anomaly Detector in Azure.
Install the Docker Engine on Server1.
From Server1, run the docker pull command.
From Server1, run the docker run command.
Query the prediction endpoint on Server1.

Answer Area

Explanation:

Actions
From Server1, run the docker push command.
Query the prediction endpoint of the Azure AI Anomaly Detector in Azure.

Answer Area
1 Install the Docker Engine on Server1.
2 From Server1, run the docker pull command.
3 From Server1, run the docker run command.
4 Query the prediction endpoint on Server1.

Answer:

Azure AI Content Safety: This service focuses on detecting harmful content.

Question: 260

You have an Azure Cognitive Search solution and a collection of handwritten letters stored as JPEG files.

You plan to index the collection. The solution must ensure that queries can be performed on the contents of the letters.

You need to create an indexer that has a skillset.

Which skill should you include?

The OCR skill processes the image, identifies text, and converts it into a machine-readable format. This extracted text can then be indexed by Azure Cognitive Search, making it searchable.

- A. key phrase extraction
- B. optical character recognition (OCR)**
- C. document extraction
- D. image analysis

The OCR skill uses machine learning models from Azure AI Vision to recognize printed and handwritten text within image files, including JPEGs.

Answer: B

Explanation:

Question: 261

HOTSPOT

You plan to provision Azure Cognitive Services resources by using the following method.

```
{  
    CognitiveServicesAccount parameters =  
        new CognitiveServicesAccount(null, null, kind, location, name,  
            new CognitiveServicesAccountProperties(), new Sku(tier));  
    result = client.Accounts.Create(resource_group_name, tier, parameters);  
}
```

You need to create a Standard tier resource that will convert scanned receipts into text.

How should you call the method? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

formRecognizer
"s0","eastus"



To provision an Azure resource, you need to use the full, official Azure region name or code, not shortened versions or combinations with service tiers.
Therefore, the correct option is eastus.

Answer:

Explanation:

Answer Area

provision_resource("res1", "FormRecognizer", "eastus", "S1")

S0 (Standard Tier): A paid tier that provides higher usage limits and more features.

Question: 262

You are developing a system that will monitor temperature data from a data stream. The system must generate an alert in response to atypical values. The solution must minimize development effort.

What should you include in the solution?

- A. Univariate Anomaly Detection
- B. Azure Stream Analytics
- C. metric alerts in Azure Monitor monitoring temperature data from a stream and generating alerts
- D. Multivariate Anomaly Detection

Duplicate

Answer: D

Question: 263

HOTSPOT

You have an Azure subscription that contains an Azure OpenAI resource. You configure a model that has the following settings:

- Temperature: 1
- Top probabilities: 0.5
- Max response tokens: 100

You ask the model a question and receive the following response.

```
{  
    "choices": [  
        {  
            "finish_reason": "stop",  
            "index": 0,  
            "message": {  
                "content": "The founders of Microsoft are Bill Gates and Paul Allen. They co-founded the company in 1975.",  
                "role": "assistant"  
            }  
        }  
    ],  
    "created": 1679014554,  
    "id": "chatcmpl-6usfn2yyjkbmESe3G4jaQR6bDSc01",  
    "model": "gpt-3.5-turbo-0301",  
    "object": "chat.completion",  
    "usage": {  
        "completion_tokens": 86,  
        "prompt_tokens": 37,  
        "total_tokens": 123  
    }  
}
```

Token Limit: If the language model's response would naturally exceed the max_tokens value, the model will stop generating tokens once it reaches that limit.

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth point.

Answer Area

completion_tokens - charged based on this token count

Statements

The subscription will be charged 86 tokens for the execution of the session.

Yes**No**

The text completion was truncated because the Max response tokens value was exceeded.



The prompt_tokens value will be included in the calculation of the Max response tokens value.



Answer:

Answer Area

Statements

The subscription will be charged 86 tokens for the execution of the session.

Yes**No**

The text completion was truncated because the Max response tokens value was exceeded.



The prompt_tokens value will be included in the calculation of the Max response tokens value.



Question: 264

HOTSPOT

You have an Azure subscription that contains an Azure OpenA1 resource named All.

You plan to develop a console app that will answer user questions.

You need to call All and output the results to the console.

How should you complete the code? To answer, select the appropriate options in the answer area

NOTE: Each correct selection is worth one point.

Answer Area

```
OpenAIclient client =  
    new OpenAIclient(new Uri(endpoint), new AzureKeyCredential(key));  
  
Response<Completions> response =  
    client.GetCompletions (deploymentName, "What is Microsoft Azure?");  
  
Console.WriteLine  
    (response.Value.Choices[0].Text);  
(response.Value.Choices[0].Text);  
(response.Value.Id);  
(response.Value.PromptFilterResults);
```



Answer:

Answer Area

```
OpenAIclient client =  
    new OpenAIclient(new Uri(endpoint), new AzureKeyCredential(key));  
  
Response<Completions> response =  
    client.GetCompletions (deploymentName, "What is Microsoft Azure?");  
Console.WriteLine  
    (response.Value.Choices[0].Text);
```

Question: 265

HOTSPOT

You have an Azure subscription.

You need to deploy an Azure AI Document Intelligence resource.

How should you complete the Azure Resource Manager (ARM) template? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

2.Upload Sample Images

Answer Area

```
{  
    "$schema": "https://schema.management.azure.com/schemas/2019-04-01/deploymentTemplate.json#",  
    "contentVersion": "1.0.0.0",  
    "parameters": {},  
    "variables": {},  
    "resources": [  
        {  
            "type": "Microsoft.CognitiveServices/accounts",  
            "Microsoft.CognitiveSearch",  
            "Microsoft.CognitiveServices",  
            "Microsoft.MachineLearning",  
            "Microsoft.MachineLearningServices",  
            "apiVersion": "2023-05-01",  
            "name": "DocumentIntelligenceDemo",  
            "location": "westeurope",  
            "sku": {  
                "name": "F0"  
            },  
            "kind": "FormRecognizer",  
            "AiBuilder",  
            "CognitiveSearch",  
            "FormRecognizer",  
            "OpenAI",  
        }  
    ]  
}
```

Answer:

Answer Area

```
{  
    "$schema": "https://schema.management.azure.com/schemas/2019-04-01/deploymentTemplate.json#",  
    "contentVersion": "1.0.0.0",  
    "parameters": {},  
    "variables": {},  
    "resources": [  
        {  
            "type": "Microsoft.CognitiveServices/accounts",  
            "apiVersion": "2023-05-01",  
            "name": "DocumentIntelligenceDemo",  
            "location": "westeurope",  
            "sku": {  
                "name": "F0"  
            },  
            "kind": "FormRecognizer",  
        }  
    ]  
}
```

Question: 266

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You are building a chatbot that will use question answering in Azure Cognitive Service for Language.

You have a PDF named Doc1.pdf that contains a product catalogue and a price list

You upload Doc1.pdf and train the model.

During testing, users report that the chatbot responds correctly to the following question: What is the price of <product>?

The chatbot fails to respond to the following question: How much does <product> cost?

You need to ensure that the chatbot responds correctly to both questions.

Solution: From Language Studio, you add alternative phrasing to the question and answer pair, and then retrain and republish the model.

Does this meet the goal?

- A. Yes
- B. No

Answer: A

Question: 267

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it as a result, these questions will not appear in the review screen.

You are building a chatbot that will use question answering in Azure Cognitive Service for Language.

You upload Doc1.pdf and train that contains a product catalogue and a price list.

During testing, users report that the chatbot responds correctly to the following question: What is the price of <product>?

The chatbot fails to respond to the following question: How much does <product> cost?

You need to ensure that the chatbot responds correctly to both questions.

Solution: from Language Studio, you create an entity for price, and then retrain and republish the model.

Does this meet the goal?

Add synonym, Add alternate Questions and pattern/matching

Duplicate

- A. Yes
- B. No

While the chatbot will now recognize "price" as an entity, it still might not connect it to the phrase "how much does cost" which uses a different structure and vocabulary.

To handle such variations, you typically need to introduce **Answer: B** synonyms or alternate phrasings for "price" like "cost," "amount," etc. within the question-answering project or knowledge

Question: 268

The chatbot may understand "What is the price of ?" due to the explicit use of the word "price".

However, it might not recognize that "How much does cost?"

You are building an internet-based training solution. The solution requires that a user's camera and microphone remain enabled.

You need to monitor a video stream of the user and detect when the user asks an instructor a question. The solution must minimize development effort.

What should you include in the solution?

Object detection identifies objects in images or video frames.
It's not directly relevant to detecting questions in speech. Y

- A. object detection in Azure AI Custom Vision
- B. the Face service in Azure AI Vision While it could detect facial expressions or head movements, these are not reliable indicators of a question being asked.
- C. language detection in Azure AI Language Service
- D. speech-to-text in the Azure AI Speech service

Answer: D

Explanation:

Question: 269

Speech-to-text for Question Detection: The most direct way to detect if a user is asking a question is to analyze their speech. The Azure AI Speech service provides speech-to-text capabilities that can transcribe the user's spoken words into text in This allows the system to then analyze the text for question indicators (e.g., question words like "what," "how," "when," or question marks).

HOTSPOT

You are building an app that will answer customer calls about the status of an order. The app will query a database for the order details and provide the customers with a spoken response.

You need to identify which Azure AI service APIs to use. The solution must minimize development effort.

Which object should you use for each requirement? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

SpeechRecognizer

SpeechSynthesizer

Speech Synthesizer: This is the right choice for providing a spoken response. The Speech Synthesizer converts text (the order details retrieved from your database) into synthesized speech, which can be played back to the customer over the phone.

Speech recognizer :audio input

(the customer's speech from the phone call) and converts it into text

Answer:

Explanation:

Answer Area

Convert customer calls into text queries: **TranslationRecognizer**

Provide customers with the order details: **SpeechSynthesizer**

Question: 270**HOTSPOT**

You are building an app that will automatically translate speech from English to French, German, and Spanish by using Azure AI service.

You need to define the output languages and configure the Azure AI Speech service.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

SpeechRecognizer: This class is for speech-to-text only and lacks built-in translation.

IntentRecognizer: This class is for identifying the intent behind spoken language, not translation.

Answer Area

```
static async Task TranslateSpeechAsync()
{
    var config = SpeechTranslationConfig.FromSubscription(SPEECH_SUBSCRIPTION_KEY, SPEECH_SERVICE_REGION);
    var languages = new List<string>
    {
        ["fr", "de", "es"]
        ["en-GB"]
        ["en", "fr", "de", "es"]
        ["fr", "de", "es"]
        ["French", "German", "Spanish"]
    };
    languages.ForEach(config.AddTargetLanguage);
    using var recognizer = new
    {
        ...
    }
}
```

TranslationRecognizer -This is designed specifically for translating spoken language into one or more target languages using the Azure Speech service.

Answer:

Explanation:

SpeakerRecognizer: This class is for speaker identification and verification, not translation

Answer Area

```
static async Task TranslateSpeechAsync()
{
    var config = SpeechTranslationConfig.FromSubscription(SPEECH__SUBSCRIPTION__KEY, SPEECH__SERVICE__REGION);
    var languages = new List<string> {"fr", "de", "es"};
    languages.ForEach(config.AddTargetLanguage());
    using var recognizer = new TranslationRecognizer();
    ...
}
```

Question: 271

You have a Microsoft OneDrive folder that contains a 20-GB video file named FileVavi. You need to index File1.avi by using the Azure Video Indexer website. What should you do?

- A. Upload File1.avi to the www.youtube.com webpage. and then copy the URL of the video to the Azure AI Video Indexer website
- B. From OneDrive. create a download link, and then copy the link to the Azure AI Video Indexer website.
- C. From OneDrive, create a sharing link for File1.avi and then copy the link to the Azure AI Video Indexer website.
- D. Download File1.avi to a local computer, and then upload the file to the Azure AI Video Indexer website.

Duplicate

Answer: D

Explanation:

Question: 272

You are developing an app that will use the Speech and language APIs.

You need to provision resources for the app. The solution must ensure that each service is accessed by using a single endpoint and credential

Which type of resource should you create?

A. Azure AI Content Safety

Azure AI Content Safety: This service focuses on detecting harmful content and is not related to Speech or Language APIs.

B. Azure AI service

Azure AI Service: This service only provides access to speech-related APIs, not the Language API.

C. Azure AI Speech

Azure AI Language: This service only provides access to language-related APIs, not the Speech API.

D. Azure AI Language

Azure AI service resource is designed to provide a single endpoint and set of credentials for accessing multiple Azure AI services, including Speech and Language. This aligns with your requirement to have a unified access point for both APIs.

Answer: B

Explanation:

Custom Vision Portal: Great for quick prototyping, experimentation, and simple image classification projects.

Question: 273

DRAG DROP

Visual Studio Project (and API/SDK): More suitable for integrating Custom Vision into complex applications, customizing models, automating processes, and using the model offline

You have an app that uses Azure AI and a custom trained classifier to identify products in images. You need to add new products to the classifier. The solution must meet the following requirements:

- Minimize how long it takes to add the products
- Minimize development effort.

Which five actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions
From the Custom Vision portal, open the project.
From the Azure Machine Learning studio, open the workspace.
From Vision Studio, open the project.
Upload sample images of the new products.
Label the sample images.
Retrain the model.
Publish the model.

Answer Area



1. Open the Custom Vision Project Portal: _____

2. Upload Sample Images

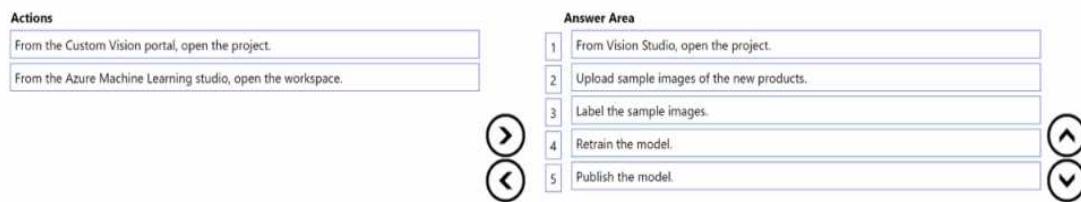
Answer:

3. Label the Sample Images:

4. Retrain the Model

5. Retrain the Model

Type text here



The Azure Custom Vision portal is generally preferred for tasks like training and managing custom image classifiers because it simplifies the process for those who are not familiar with the command-line interface.

Question: 274

You have an Azure subscription that contains an Azure AI service resource named CSAccount1 and a virtual network named VNet1. CSAccount1 is connected to VNet1.

You need to ensure that only specific resources can access CSAccount1. The solution must meet the following requirements:

- Prevent external access to CSAccount1
- Minimize administrative effort

Which two actions should you perform? Each correct answer presents part of the solution. NOTE: Each correct answer is worth one point.

While you might adjust virtual network settings, it won't directly restrict access to CSAccount1 like a service endpoint.

A. In VNet1, modify the virtual network settings.

B. In VNet1, enable a service endpoint for CSAccount1

Service endpoints allow secure connections from your virtual network to Azure AI services

C. In CSAccount1, configure the Access control (IAM) settings.

While IAM settings are crucial for managing permissions, they don't directly block network access

D. In VNet1, create a virtual subnet.

E. In CSAccount1, modify the virtual network settings.

After setting up the service endpoint in VNet1 you need to allow the virtual network access in CSAccount1 by modifying its virtual network settings and allowing access from VNet1 this will also prevent any external or other network access

Answer: B, D

Explanation:

Question: 275

HOTSPOT

You have an Azure subscription

You need to create a new resource that will generate fictional stories in response to user prompts.

The solution must ensure that the resource uses a customer-managed key to protect data.

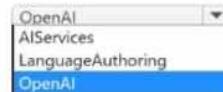
How should you complete the script? To answer, select the appropriate options in the answer area.

**1. Regenerate the secondary admin key
2. update the app to use secondary key
3. OpenAI**

NOTE: Each correct selection is worth one point.

Answer Area

```
az cognitiveservices account create -n myresource -g myResourceGroup --kind OpenAI
--encryption {
    --api-properties {
        --assign-identity keyVault",
        encryption
    }
    "keyName": "KeyName",
    "keyVersion": "secretVersion",
    "keyVaultUri": "https://issue23056kv.vault.azure.net/"
}
```



--sku S -l WestEurope

--kind openAI: This specifies that you want to create an OpenAI resource which can be used to generate fictional stories.

Answer:

Explanation:

Answer Area **encryption-key-name <key_name>:** Specifies the name of the key to be used for encryption from Azure Key Vault.

```
az cognitiveservices account create -n myresource -g myResourceGroup --kind OpenAI
--encryption {
    "keySource": "Microsoft.KeyVault",
    "keyVaultProperties": {
        "keyName": "KeyName",
        "keyVersion": "secretVersion",
        "keyVaultUri": "https://issue23056kv.vault.azure.net/"
}
```

Encryption

OpenAI

Question: 276

HOTSPOT

You have an Azure subscription that contains an Azure AI Content Safety resource named CS1. You need to use the SDK to call CS1 to identify requests that contain harmful content. How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point

Answer Area

```
var client = new TextCategoriesAnalysis(new Uri(endpoint), new AzureKeyCredential(key));
var request = new AddOrUpdateTextBlocklistItemsOptions("what is the weather forecast for Seattle");
Response<AnalyzeTextOptions> response = client.AnalyzeTextBlocklistMatch(TextCategoriesAnalysis)
```

ContentSafetyclient

AnalyzeTextOptions

Explanation:**Answer:**

Regenerate the secondary key

Answer Area

```
var client = new TextCategoriesAnalysis(new Uri(endpoint), new AzureKeyCredential(key));  
var request = new AddOrUpdateTextBlocklistItemsOptions("what is the weather forecast for Seattle");  
Response<AnalyzeTextResult> response;  
response = client.AnalyzeText(request);
```

Confirmed

Question: 277

HOTSPOT

You have an Azure subscription that contains an Azure AI Content Safety resource named CS1.

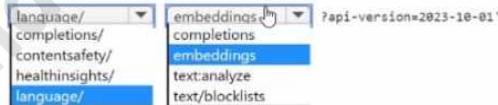
You need to call CS1 to identify whether a user request contains hateful language.

How should you complete the command? To answer, select the appropriate options in the answer area

NOTE: Each correct selection is worth one point

Answer Area

```
curl --location --request POST 'https://cs1.cognitiveservices.azure.com/  
--header 'Ocp-Apim-Subscription-Key: <your_subscription_key>' \  
--header 'Content-Type: application/json' \  
--data-raw '{  
    "text": "What is the weather forecast for Seattle",  
    "categories": ["Hate"]  
    "blocklistNames": [  
        "string"  
    ],
```



contentsafety

textanalyze

/contentsafety indicates that you are using the Content Safety service.

text:analyze specifies that you are analyzing text content.

Explanation:**Answer:**

analyze is the specific operation to perform.

language or healthinsights not related to content safety

embeddings,completions,/text/blocks are not correctpath or valid for this.

Answer Area

```
curl --location --request POST 'https://cs1.cognitiveservices.azure.com/  
--header 'Ocp-Apim-Subscription-Key: <your_subscription_key>' \  
--header 'Content-Type: application/json' \  
--data-raw '{  
    "text": "What is the weather forecast for Seattle",  
    "categories": ["Hate"]  
    "blocklistNames": [  
        "string"  
    ],  
    "haltOnBlocklistHit": true,  
    "outputType": "FourSeverityLevels"  
}'
```

Question: 278

You have an Azure subscription that contains an Azure AI Document Intelligence resource named DM. DI1 uses the Standard S0 pricing tier

You have the files shown in the following table.

Name	Size	Description
File1.pdf	800 MB	Contains scanned images
File2.jpg	1 KB	An image that has 25 x 25 pixels
File3.tiff	5 MB	An image that has 5000 x 5000 pixels

Which files can you analyze by using DI1?

A. File1.pdf only

File1.pdf: PDF is a supported format.

File2.jpg: JPG is a supported format.

File3.tiff: TIFF is a supported format.

B. File2.jpg only

but file1.pdf exceed 500Mb

C. File3.tiff only

Image dimensions: Must be between 50 x 50 pixels and 10,000 x 10,000 pixels

D. File2.jpg and File3.tiff only

Jpg,bmp,Tiff, Pdf, png where pdf is primarily support

E. File1.pdf, File2.jpg, and File3.tiff

Answer: D

Explanation:

Question: 279

You have an Azure OpenAI model named All.

You are building a web app named App1 by using the Azure OpenAI SDK. You need to configure App1 to connect to All. What information must you provide?

A. the endpoint, key, and model name

Endpoint: Tells the SDK where to connect.

B. the deployment name, endpoint, and key

API Key: Authenticates the SDK to access your resource.

C. the endpoint, key, and model type

Deployment Name: Specifies which deployed model to use within your resource.

- D. the deployment name, key, and model name

Answer: B

Explanation:

Question: 280

You have an Azure subscription that contains an Azure AI Content Safety resource named CS1. You plan to build an app that will analyze user-generated documents and identify obscure offensive terms. You need to create a dictionary that will contain the offensive terms. The solution must minimize development effort. What should you use?

since minimal dev effort and dictionary so blacklist else Text Classifier

A. a text classifier

Text Classifiers are machine learning models specifically trained to categorize text into predefined classes. In this case, the classes would be "offensive" and "not offensive."

B. text moderation

C. language detection

Text moderation: While text moderation is a broader field encompassing methods for managing online content, it often relies on tools like text classifiers to automate the identification of offensive material. Text moderation itself is not a direct tool for scanning documents.

D. a blacklist

Language detection identifies the language of a text but doesn't classify whether it's offensive or not.

Answer: D

Explanation:

Question: 281

A blacklist: Blacklists are lists of predefined offensive words or phrases. While they can be helpful, they are easily bypassed by variations, misspellings, or contextually offensive language. Text classifiers are more sophisticated and can handle these challenges

You have an Azure subscription that contains an Azure OpenAI resource named All and a user named User1. You need to ensure that User1 can perform the following actions in Azure OpenAI Studio;

- Identify resource endpoints.
- View models that are available for deployment.
- Generate text and images by using the deployed models

Cognitive Services OpenAI Contributor

it also grants the ability to create and manage model deployments, create custom fine-tuned models, and upload datasets for fine-tuning

The solution must follow the principle of least privilege Which role should you assign to User1?

Cognitive Services OpenAI User

- A. Cognitive Services OpenAI User
- B. Cognitive Services Contributor
- C. Contributor
- D. Cognitive Services OpenAI Contributor

Cognitive Services Contributor(highest access)
ability to create and manage Azure
OpenAI resources

- 1. View the resource in Azure portal:
- 2. View the resource endpoint under Keys and Endpoint:
- 3. Ability to view the resource and associated model deployments in Azure AI Foundry portal:
Grants access to view deployed models
- 4. Ability to view what models are available for deployment in Azure AI Foundry portal

Answer: A

Explanation:

- 5. Use the Chat, Completions, and DALL-E (preview) playground experiences to generate text and images with any models that have already been deployed to this Azure OpenAI resource
- 6. Make inference API calls with Microsoft Entra ID:
Allows User1 to make API calls using Azure AD authentication

Question: 282

Contributor
excessive permissions at the resource group or subscription level, far exceeding what User1 needs

You have an Azure subscription that contains an Azure OpenAI resource named AH and an Azure AI Content Safety resource named CS1.

You build a chatbot that uses AH to provide generative answers to specific questions and CS1 to check input and output for objectionable content.

You need to optimize the content filter configurations by running tests on sample questions.

Solution: From Content Safety Studio, you use the Protected material detection feature to run the tests.

Does this meet the requirement?

- A. Yes
- B. No

The user needs to optimize the content filter configurations of their chatbot to ensure that both input (user prompts) and output (chatbot responses) are free of objectionable material. They intend to do this by running tests on sample questions. However, using the "Protected material detection" feature is incorrect for this specific optimization task

Answer: B

Explanation:

Question: 283

You have an Azure subscription that contains an Azure OpenAI resource named All and an Azure AI Content Safety resource named CS1.

You build a chatbot that uses All to provide generative answers to specific questions and CS1 to check input and output for objectionable content.

You need to optimize the content filter configurations by running tests on sample questions.

Solution: From Content Safety Studio, you use the **Safety metaprompt feature** to run the tests

Does this meet the requirement?

The "Safety metaprompt" feature in Azure AI Content Safety is designed to guide the AI model's responses to be safer and more aligned with desired behavior. This feature is not used to test the effectiveness of content filter configurations using sample questions

A. Yes

Purpose of this Feature: This feature focuses on influencing the behavior of the AI model itself by providing instructions and guidelines on generating safe responses, according to Microsoft documentation.

Answer: B

Explanation:

Question: 284

You have an Azure subscription that contains an Azure OpenAI resource named All and an Azure AI Content Safety resource named CS1.

You build a chatbot that uses All to provide generative answers to specific questions and CS1 to check input and output for objectionable content.

You need to optimize the content filter configurations by running tests on sample questions.

Solution: From Content Safety Studio, you use the Monitor online activity feature to run the tests

Does this meet the requirement?

- A. Yes "Monitor online activity" feature in Azure AI Content Safety is not the correct tool for optimizing content filter configurations by directly testing sample questions. While it provides valuable insights into the overall performance of the filters, it doesn't allow you to perform targeted testing with specific inputs.
- B. No

Moderate text content is Required to optimize content filter

Answer: B

Explanation:

To optimize content filter configurations by running tests on sample questions, the correct feature to use is "Moderate text content" in Azure AI Content Safety Studio.

Question: 285

You have an Azure subscription that contains an Azure AI Document Intelligence resource named Aldoc1 in the S0 tier.

You have the files shown in the following table.

Name	Format	Password-locked	Size (MB)
File1	JPG	No	400
File2	PDF	No	250
File3	PNG	Yes	180
File4	XLSX	No	900
File5	PDF	Yes	160

You need to train a custom extraction model by using Aldoc1.

Which files can you upload to Document Intelligence Studio?

- A. File1, and File2 only
- B. File2, File4, and File5 only
- C. File1, File2, and File4 only
- D. File1, and File5 only
- E. File1, File2, File3, File4, and File5

File1 (JPG, No Password, 400 KB): JPG is a supported image format for training, and the file size is within the limit (less than 500 MB for the paid S0 tier).

File2 (PDF, No Password, 250 KB): PDF is a supported file format for training, and the file size is within the limit.

File3 (PNG, Password Protected, 180 KB): Password-protected files are not supported and cannot be used for training.

File4 (XLSX, No Password, 900 KB): While XLSX is a supported file format for analysis, it is not a supported format for training custom extraction models.

Answer: C

Explanation:

File5 (PDF, Password Protected, 160 KB): Password-protected files are not supported

Question: 286

You have a custom Azure OpenAI model.

You have the files shown in the following table.

Name	Size
File1.tsv	80 MB
File2.xml	25 MB
File3.pdf	50 MB
File4.xlsx	200 MB

You need to prepare training data for the model by using the OpenAI CLI data preparation tool.

Which files can you upload to the tool?

Supported format : JSONL,CSV,TSV,XLSX,JSON , max file size 512MB ,direct upload <200MB

- A. File1.tsv only File1.tsv (80MB): TSV is a supported format, and the size is within acceptable limits.
- B. File2.xml only File2.xml (25MB): XML is not a supported file format for the OpenAI CLI data preparation tool.
- C. File3.pdf only File3.pdf (50MB): PDF is not a supported file format for this tool.
- D. File4.xlsx only File4.xlsx (200MB): While XLSX is a supported format, the size exceeds the typical limit for direct uploads. For larger files, using Azure Blob Storage and referencing it from the CLI tool might be necessary.
- E. File1.tsv and File4.xlsx only
- F. File1.tsv,File2.xml and File4.xlsx only
- G. File1.tsv, File2.xml, File3.pdf and File4.xlsx

Answer: A

Explanation:

Question: 287

You have an Azure subscription and 10,000 ASCII files.

You need to identify files that contain specific phrases. The solution must use cosine similarity.

Which Azure OpenAI model should you use?

- A. text-embedding-ada-002 - specific for generating text embedding. Cosine similarity is used to measure the similarity between these embeddings
- B. GPT-4
- C. GPT-35 Turbo
- D. GPT-4-32k

Answer: A

Explanation:

Question: 288

HOTSPOT

You have an Azure subscription that contains an Azure AI Content Safety resource.

You are building a social media app that will enable users to share images.

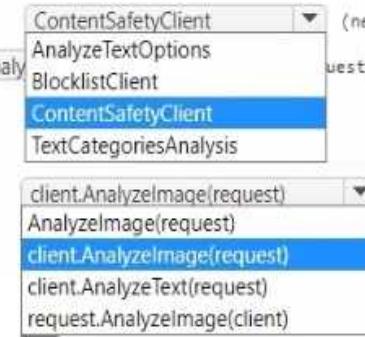
You need to configure the app to moderate inappropriate content uploaded by the users.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
public static void Analyze(AnalyzeImageOptions request)
{
    var endpoint = Environment.GetEnvironmentVariable("ENDPOINT");
    var key = Environment.GetEnvironmentVariable("KEY");
    var client = new ContentSafetyClient(new Uri(endpoint), new AzureKeyCredential(key));
    return client.AnalyzeImage(request);
}
return client.AnalyzeImage(request);
```



Type text here

Answer:

Explanation:

Answer Area

```
public static void Analyze(AnalyzeImageOptions request)
{
    var endpoint = Environment.GetEnvironmentVariable("ENDPOINT");
    var key = Environment.GetEnvironmentVariable("KEY");
    var client = new ContentSafetyClient(new Uri(endpoint), new AzureKeyCredential(key));
    return client.AnalyzeImage(request);
}
```

StreamingChatCompletions is a feature of the Azure OpenAI .NET SDK that allows you to receive chat completion responses in a streaming manner.

This means that you receive the response from the AI model incrementally as the tokens are generated, rather than waiting for the entire response to be complete. This can be useful for applications where you want to display the response to the user in real-time.

Question: 289

HOTSPOT

You have a chatbot that uses Azure OpenAI to generate responses.

You need to upload company data by using Chat playground. The solution must ensure that the chatbot uses the data to answer user questions.

How should you complete the code? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

```
// 3. Configure AzureChatExtensionsOptions
var extensions = new AzureChatExtensionsOptions();

// 4. Configure AzureCognitiveSearchChatExtensionConfiguration (assuming you are
using Azure Cognitive Search)
var searchExtension = new AzureCognitiveSearchChatExtensionConfiguration()
{
    SearchEndpoint = new Uri(searchEndpoint), // Replace with your search endpoint
    SearchKey = searchKey, // Replace with your search key
    IndexName = indexName, // Replace with your index name
};

// 5. Add the extension to the extensions list
extensions.Extensions.Add(searchExtension);
```

Answer Area

```
var options = new ChatCompletionsOptions()
{
    Messages =
    {
        new ChatMessage(ChatRole.User, "What are the differences between Azure Machine Learning and Azure AI services?"),
    },
}
```

Type text here

```
Extensions =
{
    new AzureChatExtensionConfiguration
    AzureChatExtensionsOptions
    SearchKey = new AzureKeyCredential(searchKey),
    IndexName = searchIndex,
}
```

ChatCompletionsOptions()

This sets the parameters for the chat completion API call, including the messages and extensions.

AzureChatExtensionsOptions()

Explanation:

Answer:

```
{  
    new AzureChatExtensionsOptions  
}
```

Question: 290

You have an Azure Subscription that contains an Azure OpenAI resource named AI1 and a user named User1.

You need to ensure that User1 can add custom data sources to AI1. The solution must follow the principle of least privilege.

Which role should you assign to User1?

A. Search Service Contributor

B. Cognitive Services OpenAI Contributor

C. Cognitive Services Contributor

D. Search index Data Contributor

B. Cognitive Services OpenAI Contributor: This role provides the necessary permissions to manage and configure Azure OpenAI resources, including the ability to add custom data sources (like Azure Cognitive Search indexes) to be used with Azure OpenAI models. This aligns perfectly with the stated requirement, according to Microsoft Learn

- 1.Create custom fine-tuned models
- 2.Upload datasets for fine-tuning
- 3.View, query, filter Stored completions data
- 4.Create new model deployments or edit existing model deployments
- 5.Add datasource to Azure openai data.

A. Search Service Contributor: This role grants permissions to manage Azure Cognitive Search resources, but it doesn't provide the necessary permissions to manage Azure OpenAI resources or configure data sources for them. It is more than is required.

C. Cognitive Services Contributor: This role grants broad access to Cognitive Services resources, which is more than what is required for this scenario. This role has permissions beyond just Azure OpenAI, violating the principle of least privilege, according to Microsoft Learn.

ccc Extracts structured data (fields and values) from documents.

Answer: C

Explanation:

Search Index Data Contributor: This role allows managing data within a search index, but it doesn't provide the necessary permissions to link this index as a data source to an Azure OpenAI resource.

Question: 291

You have an Azure DevOps pipeline named Pipeline1 that is used to deploy an app. Pipeline1 includes a step that will create an Azure AI services account.

You need to add a step to Pipeline1 that will identify the created Azure AI services account. The solution must minimize development effort.

Which Azure Command-Line interface (CLI) command should you run?

- A. Az resource link This command is used for managing resource links between Azure resources, such as linking a virtual network to a private endpoint. It's not relevant for identifying an Azure AI services account.
- B. Az account list This command lists Azure subscriptions associated with a user account. It does not provide information about cognitive services accounts.
- C. Az cognitiveservices account network-rule
- D. Az cognitiveservices account show

This command is designed to retrieve detailed information about an existing Azure Cognitive Services account (which includes Azure AI services). It will return the properties of the created account, including its name, resource group, endpoint, and other relevant details. This directly addresses the requirement of identifying the created account with minimal effort.

Answer: D

Explanation:

account network-rule This command is used to manage network rules (e.g., IP filtering, virtual network integration) for a Cognitive Services account. It's not the right tool for account details.

Question: 292

You are building an app that will use the Azure AI Speech service.

You need to ensure that the app can authenticate to the service by using a Microsoft Entra ID token.

Which two actions should you perform? Each answer part of the solution.

NOTE: Each correct selection is worth one point.

- A. Create a Conditional Access
- B. Create a private endpoint

private Endpoint: It is required because the Speech service uses custom subdomains with private endpoints only for Microsoft Entra ID authentication. Private endpoints enable secure communication between your application and the Speech service, ensuring that traffic does not traverse the public internet.

- C. Request an X.509 certificate
- D. Certificate a custom subdomain.
- E. Enable a virtual network service endpoint.

E. Enable a virtual network service endpoint: Virtual network service endpoints are used to secure access to Azure services from within a virtual network. While service endpoints can enhance security, they are not a direct requirement for Microsoft Entra ID authentication in the Speech service.

Certificate a custom subdomain: Azure AI services, including the Speech service, require a custom subdomain to be configured for Microsoft Entra ID authentication. This is because the regional endpoints do not support Microsoft Entra authentication

Answer: B, C

Explanation:

Microsoft Entra ID authentication provides a more secure way to access Azure resources compared to using API keys. By using Microsoft Entra ID, you can leverage role-based access control (RBAC) to manage permissions and access to the Speech service

Question: 293

You are designing a content management system.

You need to ensure that the reading experience is optimized for users who have reduced comprehension and learning differences, such as dyslexia.

Which Azure service should you include in the solution?

- A. Azure AI Translator
- B. Azure AI Document Intelligence
- C. Azure AI Immersive Reader
- D. Azure AI Language

Azure AI Immersive Reader is specifically designed to enhance reading comprehension and accessibility for users with various learning differences, including dyslexia

Answer: C

Explanation:

Question: 294

You are building an image sharing app that will use Azure AI to prevent users from sharing sexually explicit images.

You need to ensure that inappropriate images are identified correctly. The solution must minimize development effort.

What should you use?

- A. Visual Studio
 - B. Vision Studio in Azure AI Vision
 - C. Azure AI Content Safety Studio
 - D. Azure AI Studio
- D. Azure AI Studio: This is a platform for building and deploying AI solutions. It does not provide the specific content moderation tools and pre-trained models found in Content Safety Studio

Azure AI Content Safety Studio provides a user-friendly interface and pre-built models for content moderation, including the detection of sexually explicit content in images

B. Vision Studio in Azure AI Vision: Vision Studio is great for exploring various image analysis features, but it doesn't offer the same pre-built content moderation workflows and customization options as Content Safety Studio. While it can detect adult content, it's not the primary focus.

Explanation:

Answer: B

Question: 295

You are building a solution in Azure that will use Azure AI Language service to process sensitive customer data.

You need to ensure that only specific Azure processes can access the Language service. The solution must minimize administrative.

What should you include in the solution?

- A. Azure Application Gateway
- B. IPsec rules
- C. A virtual network gateway
- D. Virtual network rules

While Application Gateway can provide web traffic management and security features, it is not the primary solution for restricting access

IPsec rules can be complex to manage and are typically used for securing communication between networks, not for controlling access to individual services within Azure.

A virtual network gateway is used for connecting on-premises networks to Azure virtual networks. It's not directly relevant for restricting access

Virtual network rules allow you to restrict access to your Azure AI Language service to specific subnets within your Azure virtual networks. This approach leverages Azure's networking infrastructure to control access based on network origin, which is a more secure and manageable method than relying on IP addresses or complex network configurations.

Answer: D

Explanation:

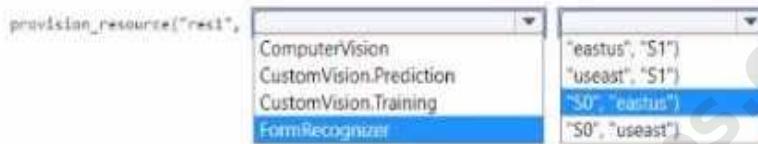
Question: 296

HOTSPOT

You plan to provision Azure AI service resources by using the following method.

```
static void provision_resource(CognitiveServicesManagementClient client, string name, string kind, string tier,
    string location)
{
    CognitiveServicesAccount parameters =
        new CognitiveServicesAccount(null, null, kind, location, name,
            new CognitiveServicesAccountProperties(), new Sku(tier));
    result = client.Accounts.Create(resource_group_name, tier, parameters);
}
```

Answer Area



FormReconizer

SO, eastus

Answer:

Explanation:

Answer Area

eastus (East US): Popular choice for general-purpose workloads, large capacity.
eastus2 (East US 2): Another popular choice for general-purpose workloads, large capacity.
westus (West US): Good for users in the Western US, often used in conjunction with eastus.
westus2 (West US 2)
centralus (Central US)
southcentralus (South Central US)

Question: 297**HOTSPOT**

S0- entry level small document

s1- higher for large scale computer vision applicaitons

You have an app that uses the AI Language custom question answering service.

You need to ad alternatives for the word testing by using the Authoring API.

How should you complete the JSON payload? To answer, select the appropriate options in the answer area.

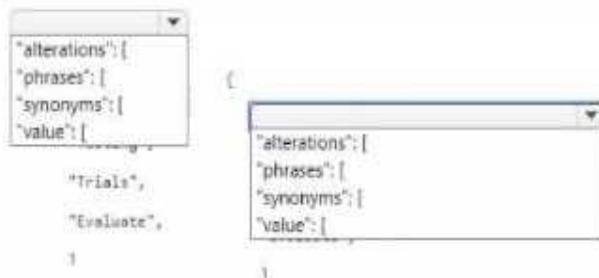
NOTE: Each correct selection is worth one point.

```
{  
    "synonyms": [  
        {  
            "value": "testing",  
            "alterations": [  
                "validation",  
                "verification",  
                "evaluation",  
                "quality assurance"  
            ]  
        }  
    ]  
}
```

Answer Area

Synonyms

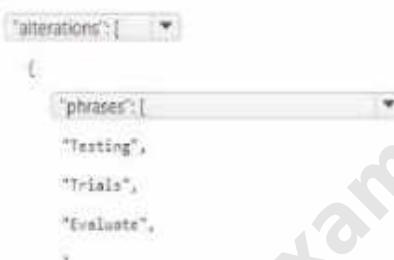
Alteratinos



Answer:

Explanation:

Answer Area



Question: 298

You have the following C# function.

```
static void MyFunction(TextAnalyticsClient textAnalyticsClient, string text)
{
    var response = textAnalyticsClient.ExtractKeyPhrases(text);
    Console.WriteLine("Key phrases:");
    foreach (string keyphrase in response.Value)
    {
        Console.WriteLine($"{keyphrase}");
    }
}
```

You call the function by using the following code.

```
MyFunction(textAnalyticsClient, "the quick brown fox jumps over the lazy dog");
```

Following 'key phrases' what output will you receive?

A. Jumps over the

B. The quick brown fox jumps over the lazy dog

C. Quick brown fox lazy dog

D. The quick

Key Phrase Extraction: This Azure AI Language feature identifies the main points in a text. It analyzes the text and returns phrases that represent the core concepts.

Noun Phrases: The most important elements of the sentence are the subject "quick brown fox" and the object "lazy dog".

Explanation:

Ignoring Minor Words: The key phrase extractor tends to ignore common words like "the" and "over" if they are not crucial to the meaning of the extracted phrase

Question: 299

DRAG DROP

You have a web app that uses Azure AI search.

When reviewing activity, you see greater than expected search query volumes. You suspect that the query key is compromised.

You need to prevent unauthorized access to the search endpoint and ensure that users only have read only access to the documents collection. The solution must minimize app downtime.

Which three action should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

- | Actions |
|--|
| >Delete the compromised key. |
| Regenerate the secondary admin key. |
| Regenerate the primary admin key. |
| Change the app to use the secondary admin key. |
| Change the app to use the new key. |
| Add a new query key. |

Answer Area



Answer:

Explanation:

Answer Area

- 1 :: Add a new query key.
- 2 :: Change the app to use the new key.
- 3 :: Delete the compromised key.

1. Regenerate the secondary admin key
2. update the app to use secondary admin key
3. regenerate the primary key
4. update the application to use new primary key

Question: 300

HOTSPOT

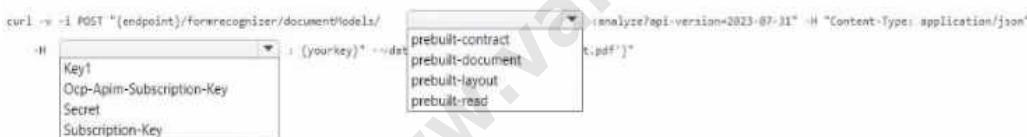
You have an Azure subscription that contains an Azure AI Document intelligence resource named D1.

You create a PDF document named test.pdf that contain tabular data.

You need to analyze Test.pdf by using D1.

How should you complete the command? To answer, select the appropriate option in the answer area.

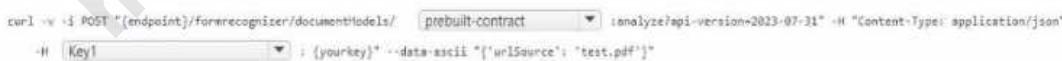
NOTE: Each correct selection is worth one point.

Answer Area

prebuilt-read: This model is designed for extracting text from documents, without understanding layout or structure. It would not reliably identify and extract tables as separate entities.

Answer:

Explanation: prebuilt-layout is optimized for understanding the visual layout of documents. If your main goal is table extraction, prebuilt-layout provides higher accuracy. If you don't need the semantic content, using prebuilt-layout is more efficient in terms of processing time and cost.



prebuilt-layout: designed for extracting text from documents, without understanding layout or structure.

prebuilt-document: This model is more comprehensive and can extract both text and layout, but it also extracts other document elements (paragraphs, headers, etc.). If you only need the tables, prebuilt-layout is more efficient. can detect table but not primary purpose

- 1) prebuilt-layout
- 2) ocp-apim-subscription-key

Question: 301

You have an Azure subscription that contain an Azure OpenAI resource named AI1.

You build a chatbot that uses AI1 to provide generation answers to specific questions.

You need to ensure that the chatbot checks all input output for objectionable content.

Which types of resource should you create first?

Azure AI Content Safety is specifically designed for content moderation tasks. It offers pre-built models and customizable workflows to detect harmful content, such as hate speech, sexual content, and violence, in both text and images.

- A. Azure Machine Learning
- B. Log Analytics
- C. Azure AI Content Safety
- D. Microsoft Defender Threat intelligence (Defender TI)

Answer: C

Explanation:

Question: 302

HOTSPOT

You are developing an app that will use the Azure AI vision API to analyze an image.

You need configure the request that will be used by the app to identity whether an image is clipart or a line drawing.

How should you complete the request? To answer. Select the appropriate options in the answer area.

NOTE: Each correct select is worth one point.

imageType Visual Feature: The imageType visual feature analyzes the image and determines if it is a photograph, a clipart image, or a line drawing. This is exactly what you need for your scenario

Answer Area

GET	https://cognitiveservices.azure.com/vision/v3.2/analyze?visualFeatures=	description	details={string}&language=e
PATCH		imageType	
POST		objects	
post		tags	

imageType

Answer:

Explanation:

description Visual Feature: This feature generates a textual description of the image content.

Answer Area

GET https://cognitiveservices.azure.com/vision/v3.2/analyze?visualFeatures= objects details={string}&language=e

objects Visual Feature: This feature detects and classifies objects present in the image. It's not relevant to identifying the image type

tags Visual Feature: This feature generates tags associated with the image content

Question: 303

DRAG DROP

You have an Azure subscription that contains an Azure OpenAI resource named AH.

You need to analyze an image to obtain a text description.

Which four actions should you perform in sequence from Azure OpenAI Studio? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions	Answer Area
Open Completions playground and select the deployed model.	
Create a new deployment and select a DALL-E model.	
Open Chat playground and select the deployed model.	
In the System message field, enter "You are an AI assistant that describes images."	1.Create a new deployment, select GPT-4 model, and set model version to vision-preview (correct model for vision tasks)
Create a new deployment, select a GPT-4 model, and set Model version to vision-preview.	2.Open chat playground and select deployment model. - takes you to the chat interface where vision models can be used.
Create a new deployment, select a text-embedding-ada-002 model, and set Model version to 2.0.	3.In System message field, enter "You are an AI Assistant that describes images." - tell the model its role.
In the Chat session pane, enter a text prompt of Describe this image, and upload an image by using the attachment button.	4.In chat session pane, enter a text prompt of "describe this image" and upload an image by using attachments in button. - performs the actual image analysis and text generation.

Not Correct

text-embedding-ada-002: not for image

The DALL-E models are for image generation, not image analysis and description.

Completions playground: The completions playground is designed for text-based tasks and doesn't support image input.

Answer:

Explanation:

Actions

- :: Open Completions playground and select the deployed model.
- :: Create a new deployment and select a DALL-E model.
- :: Open Chat playground and select the deployed model.

Answer Area

- 1 :: In the System message field, enter You are an AI assistant that describes images.
- 2 :: Create a new deployment, select a GPT-4 model, and set Model version to vision-preview.
- 3 :: Create a new deployment, select a text-embedding-ada-002 model, and set Model version to 2.0.
- 4 :: In the Chat session pane, enter a text prompt of **Describe this image**, and upload an image by using the attachment button.

Question: 304

You have a local folder that contains the files shown in the following table.

Name	Format	Length (mins)	Size (MB)
File1	WMV	34	400
File2	AVI	90	1,200
File3	MOV	300	980
File4	MP4	80	1,800

You need to analyze the files by using Azure Ai Video Indexer. Which files can you upload to the Video Indexer website?

- A. File1, File2 and File4 only
- B. File1, and File2 only
- C. File1, File2, and File3 only
- D. File1, File2, File3 and File4
- E. File1, and File3 only

Supported Formats
 MP4, MOV, MKV, MPG, MPEG
 WMV
 AVI, ASF
 3GP

Type text here

Use direct upload for files up to 2 GB.

Use a URL for files up to 30 GB.

Based on the provided table:

File1.WMV: Supported.
 File2.AVI: Supported.
 File3.MOV: Supported.
 File4.MP4: Supported.

File Duration Limit: The maximum duration of a video file is 6 hours for most presets. However, for the Basic Audio preset, the limit is 12 hours.

Direct Upload Limitations: When uploading via direct upload (as a byte array), there is a higher dependency on network, service reliability, connectivity, upload speed, and lost packets which might affect your performance and ability to scale.

Answer: E

Explanation:

Number of Files: When uploading and indexing through the Video Indexer website, you can upload up to 10 videos per request.

API Request Limit: When submitting indexing requests through the API, there's a limit of 10 requests per second and up to 120 requests per minute.

Question: 305

You have an app named App1 that uses a custom Azure AI Document Intelligence model to recognize contract documents. You need to ensure that the model supports an additional contract format. The solution must minimize development effort. What should you do?

- A. Lower the confidence score threshold of App1.
- B. Lower the accuracy threshold of App1.
- C. Add the additional contract format to the existing training set. Retrain the model.**
- D. Create a new training set and add the additional contract format to the new training set.
- E. Create and train a new custom model.

Answer: C

Explanation:

Question: 306

HOTSPOT

You have an Azure subscription.

You plan to build a solution that will analyze scanned documents and export relevant fields to a database.

You need to recommend which Azure AI service to deploy for the following types of documents:

- Internal expenditure request authorization forms
- Supplier invoices

The solution must minimize development effort.

What should you recommend for each document type? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point

Answer Area Internal expenditure request authorization forms typically have unique layouts and contain specific fields that are not found in standardized document types like invoices or receipts

Internal expenditure request authorization forms:	<input type="text" value="An Azure AI Document Intelligence custom model"/> An Azure AI Document Intelligence custom model An Azure AI Document Intelligence pre-built model Azure AI Custom Vision Azure AI Immersive Reader Azure AI Vision
Supplier invoices:	<input type="text" value="Azure AI Immersive Reader"/> An Azure AI Document Intelligence custom model An Azure AI Document Intelligence pre-built model Azure AI Custom Vision Azure AI Immersive Reader Azure AI Vision

If a pre-built model is available for internal expenditure authorization forms, that would be the quickest path. Otherwise, you would need to create and train a custom model, which is the second best option.

Answer:

Explanation:

Document Intelligence is designed for form processing. It can automatically extract key-value pairs and tables from documents. If you cannot find a pre-built model, you would need to train a custom model to achieve the task

Answer Area

Internal expenditure request authorization forms:	<input type="text" value="An Azure AI Document Intelligence custom model"/>
Supplier invoices:	<input type="text" value="Azure AI Immersive Reader"/>

Internal Expenditure

An Azure AI Document Intelligence custom model

Supplier Invoices

An Azure AI Document Intelligence pre-built model

Question: 307

You have an Azure subscription that contains an Azure OpenAI resource named All and ari Azure AI Content Safety resource named CS1.

You build a chatbot that uses All to provide generative answers to specific questions and CS1 to check input and output for objectionable content

You need to optimize the content filter configurations by running tests on sample questions.

Solution: From Content Safety Studio, you use the Moderate text content feature to run the tests.

Does this meet the requirement?

A. Yes

B. No

Custom Extraction Model:Extracts structured data (fields and values) from documents.
Explanation:

Type text here

Question: 308

You have an Azure subscription that contains an Azure OpenAI resource named OpenAI1 and a user named User1.

You need to ensure that User1 can upload datasets to OpenAI1 and finetune the existing models.
The solution must follow the principle of least privilege.

Which role should you assign to User1?

Duplicate

- A. Cognitive Services Contributor
- B. Contributor
- C. Cognitive Services OpenAI User
- D. Cognitive Services OpenAI Contributor

Answer: C

Explanation:

Question: 309

DRAG DROP

You have an Azure subscription.

You are building a chatbot that will use an Azure OpenAI model.

You need to deploy the model.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions

⋮ Apply for access to Azure OpenAI.

⋮ Deploy the embeddings model.

⋮ Provision Azure API Management.

⋮ Provision an Azure OpenAI resource.

⋮ Deploy the GPT model.

⋮ Deploy the DALL-E model.

Answer Area

1. Gain access to Azure OpenAI.

2. Create the necessary resources.

3. Deploy a suitable model for a chatbot.

Not correct

Deploy the embeddings model: Embedding models are used for semantic search and other tasks, not for text generation, which is required for a chatbot.

Provision Azure API Management: While API Management is valuable for managing and securing your API endpoints, it's not a prerequisite for deploying an Azure OpenAI model.

Deploy DALL-E model: DALL-E models are for image generation, not for text generation in a chatbot.

Answer:**Explanation:****Actions**

⋮ Apply for access to Azure OpenAI.

⋮ Deploy the embeddings model.

⋮ Provision Azure API Management.

Answer Area

1 ⋮ Provision an Azure OpenAI resource.

2 ⋮ Deploy the GPT model.

3 ⋮ Deploy the DALL-E model.

Flow

Access: You must have access to the Azure OpenAI service before you can provision a resource or deploy any models.

Resource Creation: You need to provision the Azure OpenAI resource before you can deploy the GPT model.

Model Deployment: The GPT model needs to be deployed before your application can use it.

Question: 310**HOTSPOT**

You have 1,000 scanned images of hand-written survey responses. The surveys do NOT have a consistent layout.

You have an Azure subscription that contains an Azure AI Document Intelligence resource named Aldoc1.

You open Document Intelligence Studio and create a new project.

You need to extract data from the survey responses. The solution must minimize development effort.

To where should you upload the images, and which type of model should you use? To answer, select the appropriate options in the answer area.

Not correct

NOTE: Each correct selection is worth one point.

Identity Document (ID): This pre-built model is designed for extracting data from identity documents (like passports or driver's licenses), not from survey responses.

Answer Area

Upload to:

- An Azure Storage account
- An Azure Cosmos DB account
- An Azure Files share
- An Azure Storage account**

Model type:

- Identity document (ID)
- Custom neural
- Custom template
- Identity document (ID)**

Custom Template: Custom template models require consistent layouts and may not be suitable for the handwritten documents in this case.

Identity Document (ID): This is the same as the first option and is unsuitable for extracting data from survey responses

Answer:

Explanation:

Answer Area

Upload to: An Azure Storage account

Azure Storage Account.

Model type: Identity document (ID)

Custom Neural model

Custom Neural Models: Custom Neural models are well-suited for scenarios where the document layout is variable and requires advanced learning capabilities. You can train a custom model using your specific dataset of handwritten survey responses.

Question: 311

You are developing an app that will use the text-to-speech capability of the Azure AI Speech service.

The app will be used in motor vehicles.

You need to optimize the quality of the synthesized voice output.

Which Speech Synthesis Markup Language (SSML) attribute should you configure?

A. the style attribute of the msstts: express -as element

The msstts:express-as element is used to control the speaking style and emotional tone of the voice,

B. the level attribute of the emphasis element

The emphasis element adjusts the emphasis (loudness and speaking rate) of certain words or phrases

C. the pitch attribute of the prosody element

The prosody element is used to control the pitch, rate, and volume of the speech,

D. the effect attribute of the voice element

The effect attribute within the voice element in SSML is specifically designed to optimize the synthesized speech for different environments

For example, setting effect="eq_car" will optimize the speech for playback in vehicles, accounting for factors like speaker response and background noise

Answer: D

Explanation:

Question: 312

HOTSPOT

You are building a text-to-speech app that will use a custom neural voice.

You need to create an SSML file for the app. The solution must ensure that the voice profile meets the following requirements:

- Expresses a calm tone
- Imitates the voice of a young adult female

How should you complete the code? To answer, select the appropriate options in the answer area


a. NOTE: Each correct selection is worth one point.

Answer Area

```
<mstts:express-as
    role="YoungAdultFemale"
    style="gentle">
```

How can I assist you?

```
</mstts:express-as>
```



Answer:

Explanation: role for Custom Voices: The role attribute can select the base voice profile when using custom neural voices. For example, choose a custom voice trained to sound like a "youngAdultFemale."

Answer Area

style for Refinements: The style attribute can further adjust the voice within the <voice> tag after the voice profile has been selected, such as "gentle" or "calm."

```
<mstts:express-as role="YoungAdultFemale" style="gentle">  
    How can I assist you?  
</mstts:express-as>  
...  


| Role      | Style |
|-----------|-------|
| Confirmed |       |


```

Question: 313

You develop a custom question answering project in Azure AI Language. The project will be used by a chatbot.

You need to configure the project to engage in multi-turn conversations. What should you do?

- A. Add alternate questions. confirmed
 - B. Enable chit-chat.
 - C. Add follow-up prompts. Follow-up prompts are specifically designed to enable multi-turn conversations in Azure AI Language's custom question answering projects. duplicate
 - D. Enable active learning. enable the chatbot to proactively guide the conversation. Alternate questions are about matching the same answer, not enabling a multi-turn dialog.

Enable chit-chat: Chit-chat adds general conversational capabilities to a bot, allowing it to handle greetings, small talk, and other basic interactions. While this makes a chatbot more natural, it doesn't specifically enable structured, multi-turn question-and-answer flows.

Answer: C

Active learning is a feedback mechanism where the system learns from user interactions to improve its accuracy over time. While helpful for model improvement, it doesn't directly enable the multi-turn conversational flow you need for a chatbot.

Explanation:

Question: 314

HOTSPOT

1. Enable Active learning
2. open the review suggestion pane
3. update the question and answers pane
4. Retrain

You are building a language learning solution.

You need to recommend which Azure services can be used to perform the following tasks:

- Analyze lesson plans submitted by teachers and extract key fields, such as lesson times and required texts.
- Analyze learning content and provide students with pictures that represent commonly used words or phrases in the text

The solution must minimize development effort.

Which Azure service should you recommend for each task? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Analyze lesson plans:	Azure AI Document Intelligence
	Azure AI Search
	Azure AI Custom Vision
	Azure AI Document Intelligence
	Immersive Reader

Analyze learning content:	Immersive Reader
	Azure AI Search
	Azure AI Custom Vision
	Azure AI Document Intelligence
	Immersive Reader

Reasoning: Azure AI Custom Vision allows training custom image classification models with minimal coding. The model can be trained to associate words or phrases with representative images. This is ideal for creating visual aids for language learners. The API can then be used to analyze the learning content and retrieve the corresponding pictures.

Answer:

Explanation:

Answer Area

Azure AI Document Intelligence

Analyze lesson plans:	Azure AI Document Intelligence
-----------------------	--------------------------------

Analyze learning content:	Immersive Reader
---------------------------	------------------

Azure AI Custom Vision

Reasoning: Azure AI Custom Vision allows training custom image classification models with minimal coding. The model can be trained to associate words or phrases with representative images. This is ideal for creating visual aids for language learners. The API can then be used to analyze the learning content and retrieve the corresponding pictures.

Custom Extraction Model

Question: 315

You have a product knowledgebase that contains multiple PDF documents.

You need to build a chatbot that will provide responses based on data in the knowledgebase. The solution must minimize development effort and costs.

What should you include in the solution?

designed for building more sophisticated natural language understanding models that can handle complex conversational flows and intent detection

A. Azure AI Language conversational language understanding (CLU)

This service simply identifies the language of the input text. While potentially useful, it's not directly relevant to building a question-answering chatbot. It doesn't help with content extraction or answering user questions.

C. Azure AI Language custom question answering

access to large language models that can perform various tasks, including question answering.

D. Azure OpenAI

However, it requires more effort to integrate with your custom knowledge base, build prompts

Azure AI Language custom question answering (formerly QnA Maker): This service is specifically designed for building question answering systems using existing content. You can upload your PDF documents to the service, and it will automatically extract the content and create a question-answer knowledge base. Then, you can easily query this knowledge base to find the answers relevant to user questions. It significantly reduces development time and cost compared to building a custom solution from scratch. It handles the extraction, indexing, and querying for you

Answer: C

Explanation:

Question: 316

You need to ensure that the chatbot can classify user input into separate categories. The categories must be dynamic and defined at the time of inference.

Which service should you use to classify the input?

A. Azure OpenAI text summarization

Azure AI Language custom text classification.

B. Azure OpenAI text classification

This service allows you to build custom models for classifying text into categories that you define and train. The key advantage here is that you can define your categories dynamically during inference, which aligns with the requirements.

C. Azure AI Language custom named entity recognition (NER)
name,date,location only

Not correct

D. Azure AI Language custom text classification

Azure OpenAI text summarization: This service summarizes large text documents. It is not designed for classifying text based on categories. It does not allow dynamic categories.

Answer: D

Azure OpenAI text classification: This service can classify text using pre-trained models, but it does not support custom models with dynamic category definitions. While useful, it doesn't meet the requirements of the task.

Explanation:

Question: 317

In Azure AI Studio, you use Completions playground with the GPT-35 Turbo model.

You have a prompt that contains the following code.

```
function F(n)
{
    var f = [0, 1];
    for (var i = 2; i < n; i++) f[i] = f[i-1] + f[i-2];
    return f;
}
```

You need the model to create an explanation of the code. The solution must minimize costs. What should you do?

- A. Change the model to GPT-4-32Ic
- B. Add// what does function F do? to the prompt.**
- C. Add function F(explanation) to the prompt.
- D. Set the temperature parameter to 1.

The goal is to get a good explanation while keeping costs low. Using GPT-35 Turbo is already more cost-effective than GPT-4.

Answer: B

Explanation:

Question: 318

You have an Azure subscription that contains an Azure OpenAI resource named AM.

You build a chatbot that uses AI to provide generative answers to specific questions.

You need to ensure that questions intended to circumvent built-in safety features are blocked.

Which Azure AI Content Safety feature should you implement?

- A. Protected material text detection
- B. Jailbreak risk detection**
- C. Monitor online activity
- D. Moderate text content

This feature is specifically designed to identify and block user prompts (questions) that are intentionally crafted to bypass the safety mechanisms built into AI models like Azure OpenAI

Answer: B

Explanation:

Question: 319

HOTSPOT

You have an Azure subscription that contains an Azure OpenAI resource. Multiple different models are deployed to the resource.

You are building a chatbot by using Chat playground in Azure AI Studio.

You need to ensure that the chatbot generates text in concise formal business language. The solution must meet the following requirements:

- Reduce the cost of running the language model.
- Maintain the size of the chatbot history window.

Which two settings should you configure? To answer, select the appropriate settings in the answer are

a. NOTE: Each correct selection is worth one point.

The screenshot shows the 'Chat playground' setup interface. On the left, a sidebar lists various sections: Home, Get started, Model catalog, Playgrounds (with Chat selected), Assistants (PREVIEW), Real-time audio (PREVIEW), Images, Completions, Tools (Fine-tuning, Stored completions, PREVIEW, Batch jobs), Shared resources (Deployments, Quota, Content filters, Data files, Vector stores PREVIEW), and a 'Start with a sample prompt' section. The main area is titled 'Chat playground' and contains a 'Setup' panel with a deployment dropdown set to 'gpt-4o (version:2024-05-13)', a 'Give the model instructions and context' text area containing 'You are an AI assistant that helps people find information.', and buttons for 'Apply changes' and 'Generate prompt'. Below this is a 'Add your data' section with a red border, containing text about asking questions about own data and a link to learn more about data protection, along with a '+ Add a data source' button. To the right is a 'Chat history' panel showing a blue speech bubble icon and the text 'Start with a sample prompt'. Three sample prompts are listed: 'Marketing Slogan' (Create a catchy marketing slogan for a new eco-friendly product.), 'Creative Storytelling' (Write a short story about a time traveler who accidentally changes a major historical event.), and 'Historical Fiction' (Write a scene set in ancient Rome, focusing on the daily life of a common citizen.). A large watermark 'https://www.validexamdump.com' is diagonally across the page.

Answer:

Explanation:

Answer Area

Chat playground

Regenerate the seconda

Export View Code Prompt flow Evaluate Deploy to a web app Import

Deployment * + Create new deployment ✓
gpt-35-turbo (version:0301)

System message Add your data Parameters

Apply changes Reset to default

System message ⓘ ✓
You are an AI assistant that helps people find information.

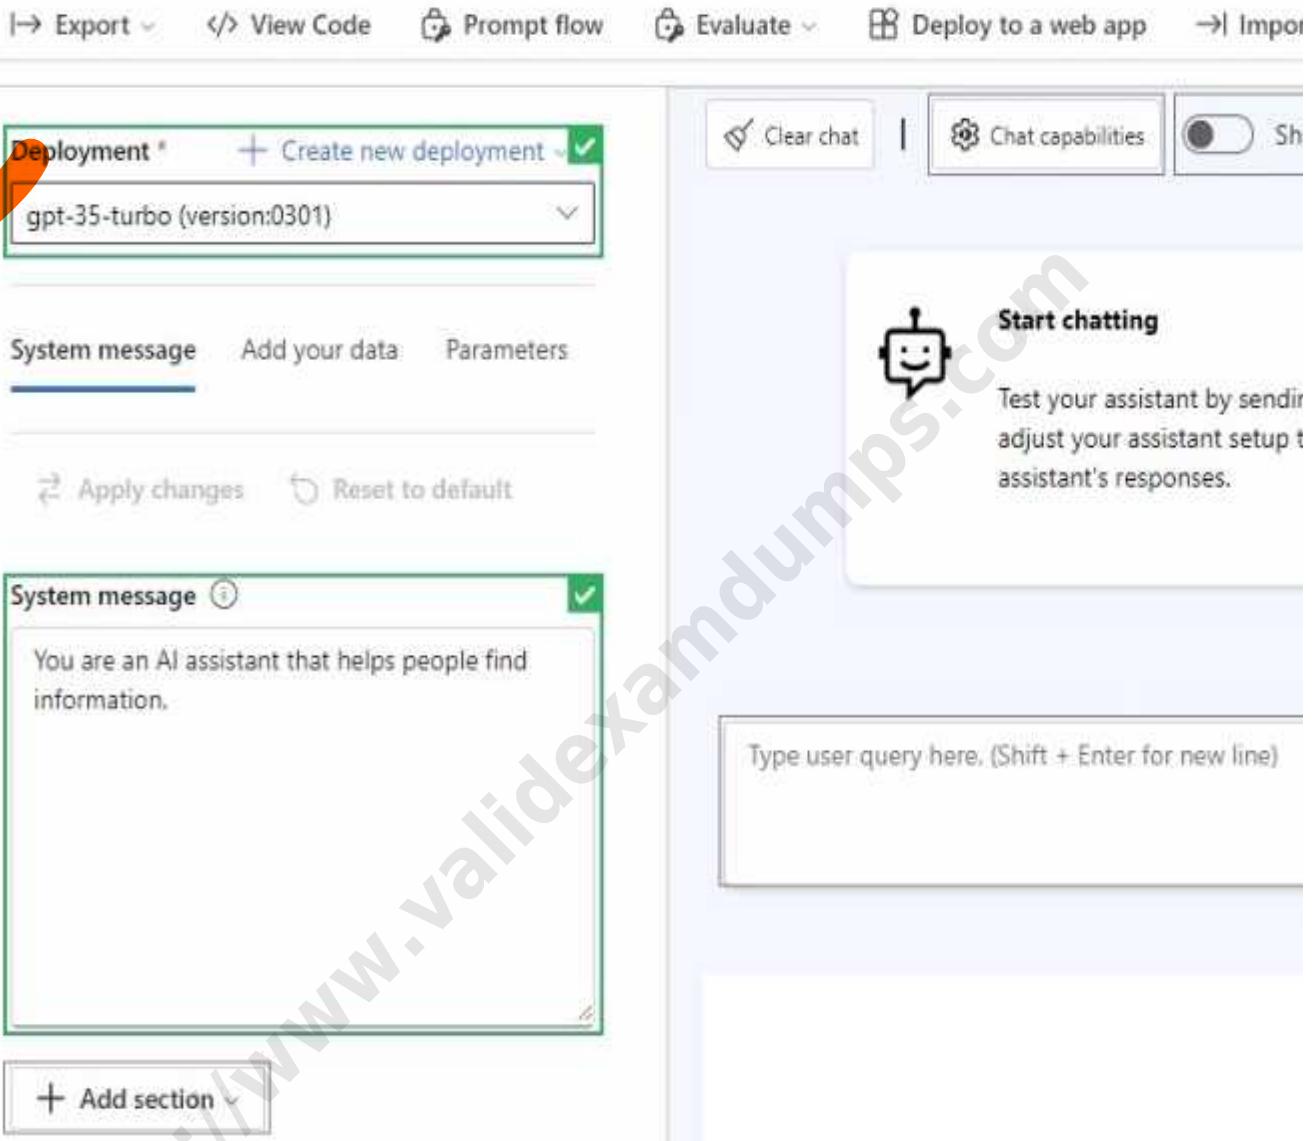
+ Add section

Note: A good system message ensures the model follows the instructions and also minimizes the number of tokens generated, thus reducing costs

Clear chat | Chat capabilities | Sh

Start chatting
Test your assistant by sending a message. You can adjust your assistant setup to change its responses.

Type user query here. (Shift + Enter for new line)

**Question: 320****HOTSPOT**

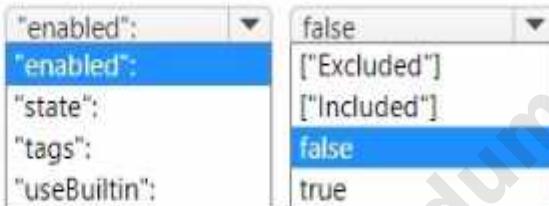
You have an Azure subscription that contains an Azure AI Video Indexer account.

You need to add a custom brand and logo to the indexer and configure an exclusion for the custom brand. How should you complete the REST API call? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

```
{  
    "referenceUrl": "https://www.contoso.com/Contoso",  
    "id": 97974,  
    "name": "Contoso",  
    "accountId": "ContosoAccountId",  
    "lastModifierUserName": "SampleUserName",  
    "created": "2023-04-25T14:59:52.7433333",  
    "lastModified": "2023-04-25T14:59:52.7433333",  
    "enabled": false  
}
```



You can only update the category of a brand. You can't change the name of a brand because names of brands are unique. If you need to change the brand name, delete the entire brand (see next section) and create a new brand with the new name.

Answer:

Explanation:

Select the pencil icon next to the brand that you want to edit.
Select the Update button to update the brand with the new inform

Answer Area

```
{  
    "referenceUrl": "https://www.contoso.com/Contoso",  
    "id": 97974,  
    "name": "Contoso",  
    "accountId": "ContosoAccountId",  
    "lastModifierUserName": "SampleUserName",  
    "created": "2023-04-25T14:59:52.7433333",  
    "lastModified": "2023-04-25T14:59:52.7433333",  
    "enabled": false  
}
```

Add a brand to exclude list

Select + Create new brand.

Provide a name (required), category (optional).

Select Save. The brand is added to the Exclude brands list.

Question: 321

HOTSPOT

You are developing an application that will use the Azure AI Vision client library. The application has the following code.

```
def analyze_image(local_image):
    with open(local_image, "rb") as image_stream:
        image_analysis = client.analyze_image_in_stream(
            image=image_stream,
            visual_features=[
                VisualFeatureTypes.tags,
                VisualFeatureTypes.description
            ]
        )
        for caption in image_analysis.description.captions:
            print(f"\n{caption.text} with confidence {caption.confidence}")
        for tag in image_analysis.tags:
            print(f"\n{tag.name} with confidence {tag.confidence}")
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Answer Area

Statements	Yes	No
The code will perform face recognition.	<input type="radio"/>	<input type="radio"/>
The code will list tags and their associated confidence.	<input type="radio"/>	<input type="radio"/>
The code will read an image file from the local file system.	<input type="radio"/>	<input type="radio"/>

Answer:

Explanation:

The code will perform face recognition. No, the code does not include any features related to face recognition. It is only analyzing the image for tags and descriptions.

The code will list tags and their associated confidence. Yes, the code includes a loop that prints each tag and its associated confidence level.

The code will read an image file from the local file system. Yes, the code opens an image file from the local file system for analysis.

Answer Area
Type text here



Statements	Yes	No
The code will perform face recognition.	<input type="radio"/>	<input checked="" type="radio"/>
The code will list tags and their associated confidence.	<input checked="" type="radio"/>	<input type="radio"/>
The code will read an image file from the local file system.	<input checked="" type="radio"/>	<input type="radio"/>

Question: 322

You have an Azure subscription that contains an Azure AI Content Safety resource named CS1.

You create a test image that contains a circle.

You submit the test image to CS1 by using the curl command and the following command-line parameters.

```
--data-raw '{  
    "image": {  
        "content": "<base_64_string>"  
    },  
    "categories": [  
        "violence"  
    ],  
    "outputType": "EightSeverityLevels"  
}
```

What should you expect as the output?

- A. 0 is CORRECT because the image being submitted contains only a simple circle, which does not contain any violent content. When using Azure AI Content Safety with the "Violence" category and the outputType set to "EightSeverityLevels," the service evaluates the image for signs of violence. The "EightSeverityLevels" output ranges from 0 to 7, where 0 indicates "Non-offensive" and 7 represents "Very High" severity. Since a circle is considered harmless, the expected severity score is 0, indicating no detection of violent content.
- B. 0.0
- C. 7
- D. 100

Answer: A,C

Explanation:

Question: 323

You have an Azure subscription.

You are building a social media app that will enable users to share images.

You need to ensure that inappropriate content uploaded by the users is blocked. The solution must minimize development effort

What are two tools that you can use? Each correct answer presents a complete solution

NOTE: Each correct selection is worth one point.

- A. Microsoft Defender for Cloud Apps
- B. Azure AI Custom Vision
- C. Azure AI Vision
- D. Azure AI Content Safety
- E. Azure AI Document Intelligence

While powerful for building custom image recognition models, it would require significantly more development effort to train a model to detect inappropriate content, compared to using pre-built services.

Answer: B,D

Explanation:

Question: 324

You have an Azure subscription.

You need to build an app that will compare documents for semantic similarity. The solution must meet the following requirements:

- Return numeric vectors that represent the tokens of each document.
- Minimize development effort.

Which Azure OpenAI model should you use?

- A. GPT-3.5
- B. embeddings
- C. DALL-E
- D. GPT-4

Answer: B

Explanation:

Question: 325

You have an Azure subscription that contains an Azure OpenAI resource.

You deploy the GPT-4 model to the resource.

You need to ensure that you can upload files that will be used as grounding data for the model.

Which two types of resources should you create? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Azure AI Bot Service
- B. Azure SQL
- C. Azure AI Document Intelligence
- D. Azure Blob Storage
- E. Azure AI Search

Purpose: Blob storage is the primary service for storing unstructured data in Azure, including the files that will be used to ground the GPT-4 model.

Azure AI Search (formerly known as Azure Cognitive Search) is used to create a searchable index of the content in your blob storage.

Answer: D, E

Explanation:

Question: 326

HOTSPOT

You have 100,000 images.

You need to build an app that will perform the following actions:

- Identify road signs in the images and extract the text on the signs.
- Analyze the text to identify well-known locations.

The solution must minimize development effort.

What should you use for each action? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Answer Area

Azure AI Vision

Extract the text:

Azure AI Vision
Azure AI Document Intelligence
Azure AI Language
Azure AI Search
Azure AI Vision

Azure AI Language (Named Entity Recognition)

Identify well-known locations:

Azure AI Search
Azure AI Document Intelligence
Azure AI Language
Azure AI Search
Azure AI Vision

Explanation:

Answer:

Answer Area

Extract the text:

Identify well-known locations:

Question: 327

HOTSPOT

You have an Azure subscription that contains an Azure AI Content Safety resource named Resource1.

You create the following cURL command.

```
curl -X POST "https://resource1.cognitiveservices.azure.com/contentsafety/text:detectProtectedMaterial?api-version=2024-09-01" \
-H "Content-Type: application/json" \
-H "Ocp-Apim-Subscription-Key: <your_subscription_key>" \
-d '{ "text": "<your_content>" }'
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth point.

1. The text value is a string representing the text, and the overall structure is JSON, not that the content itself must be json.

Answer Area

Statements	Yes	No
The <code>text</code> value must use JSON-formatted text.	<input type="radio"/>	<input checked="" type="radio"/>
The command will analyze inputted text and identify whether the text contains published song lyrics.	<input checked="" type="radio"/>	<input type="radio"/>
The <code>Ocp-Apim-Subscription-Key</code> value must contain the ID of the Azure subscription that hosts Resource1.	<input type="radio"/>	<input type="radio"/>

2. detectProtectedMaterial feature of the API is specifically designed to analyze text and determine if it contains potentially sensitive material, including text similar to copyrighted content. This might be content such as lyrics.

Explanation:

3. It should be the Content Safety Resource API key, not the Azure subscription ID

Answer Area

Statements	Yes	No
The <code>text</code> value must use JSON-formatted text.	<input checked="" type="radio"/>	<input type="radio"/>
The command will analyze inputted text and identify whether the text contains published song lyrics.	<input type="radio"/>	<input checked="" type="radio"/>
The <code>Ocp-Apim-Subscription-Key</code> value must contain the ID of the Azure subscription that hosts Resource1.	<input type="radio"/>	<input checked="" type="radio"/>

Question: 328

HOTSPOT

You have an Azure subscription that contains an Azure AI Document Intelligence resource named DM.

You create a PDF document named Test.pdf that contains tabular data.

You need to analyze Test.pdf by using DM.

How should you complete the command? To answer, select the appropriate options in the answer area

NOTE: Each correct selection is worth one point.

Answer Area

```
curl -v -i POST "{endpoint}/formrecognizer/documentModels/" :analyze?api-version=2023-07-31" -H "Content-Type: application/json" -H "Ocp-Apim-Subscription-Key: {yourkey}" --data-ascii "{\"urlSource": \"test.pdf\"}"
```

1. prebuild-layout- good for tables, text and selection marks from pdf

2. ocp-apim-subscription-key

Explanation:

Answer:

Answer Area

```
curl -v -i POST "{endpoint}/formrecognizer/documentModels/" :analyze?api-version=2023-07-31" -H "Content-Type: application/json" -H "Ocp-Apim-Subscription-Key: {yourkey}" --data-ascii "{\"urlSource\": \"test.pdf\"}"
```

Type text here

Question: 329

You are building a social media messaging app.

You need to identify in real time the language used in messages.

Which SDK package should you install?

- A. Azure.AI.Translation.Text
- B. Microsoft.CognitiveServices.Speech
- C. Azure.AI.Translation.Document
- D. Azure.AI.Translation.Speech

B. Microsoft.CognitiveServices.Speech: This SDK is used for speech-to-text, text-to-speech, and speech translation functionalities, which are not needed for the current task.

C. Azure.AI.Translation.Document: This SDK is designed for translating documents, not individual text messages.

A.Azure.AI.Translation.Text is the SDK specifically designed for text-based translation and language detection using Azure AI Translator. It offers features like language detection, translation, transliteration, dictionary lookup, and dictionary examples.

Answer: A

Explanation:

D.Azure.AI.Translation.Speech: This SDK focuses on translating speech, not text messages.

Question: 330

You have a training dataset that contains 10,000 PDF documents. The documents contain scanned books, comics, and magazines.

You are building a solution that will use Azure AI and a custom model.

You need to train the model by using Language Studio. The solution must meet the following requirements:

Multi-label Classification:

- Tag each item as a book, comic, or magazine.
- Minimize development effort.

Purpose: This type of project is designed to assign one or more labels (categories) to a single document or text input.

What should you use?

Matching Requirements: It directly addresses the requirement to tag each PDF document as either a "book", "comic", or "magazine." One document can only belong to one of the three categories.

- A. a custom extraction model
- B. a multi label classification project
- C. a custom named entity recognition (NER) project
- D. a multi label image classification model

Custom Extraction Model:**Answer: B**

Extracts structured data (fields and values) from documents.

Explanation:

Question: 331

DRAG DROP

You are building a phone call handling solution that will use the Azure AI Speech service and a custom neural voice.

You need to create a custom speech model.

Which five actions should you perform in sequence from Speech Studio? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Actions

- :: Upload a consent statement for the voice talent as a WAV file.
- :: Upload speech samples as WMA files.
- :: Create a custom voice project.
- :: Upload a consent statement for the voice talent as a signed PDF file.
- :: Analyze the quality of the audio data and resolve identified issues.
- :: Upload speech samples as MP3 files.
- :: Train the model by using a neural training method.

Answer Area

1. Create a custom voice project:
2. Upload a consent statement for voice talent as a signed PDF file:
for voice voice essential for compliance.
3. Upload speech samples as WAV files:
4. Analyze the quality of audio data and resolve identified
5. Train the model by using neural training method:

Explanation:

Actions

- :: Upload a consent statement for the voice talent as a WAV file.
- :: Upload speech samples as WMA files.

Answer Area

- 1 :: Create a custom voice project.
- 2 :: Upload a consent statement for the voice talent as a signed PDF file.
- 3 :: Analyze the quality of the audio data and resolve identified issues.
- 4 :: Upload speech samples as MP3 files.
- 5 :: Train the model by using a neural training method.

Question: 332

You are building an app that will analyze documents by using the Azure AI Language service.

You need to identify industry-specific technical terms in the documents. The solution must minimize development effort.

What should you use?

A. key phrase extraction

B. custom named entity recognition (NER)

C. conversational language understanding (CLU)

D. language detection detect the language of text , not for technical terms extraction.

Custom NER is specifically designed to identify and extract custom entities from text, which can be tailored to your needs by creating a custom entity type for industry-specific terminology

Key Phrase Extraction-Extracts important phrases from text, not necessarily specific industry terms.

CLU - not identify industry terms. focus on natural language intent
not specialized for terminology extraction.

Answer: B

Explanation:

Question: 333

You have an Azure subscription that contains an Azure AI Language custom question answering project named QA1.

You need to import question and answer pairs to QA1.

Which two file formats can you use? Each correct answer presents a complete solution

NOTE; Each correct selection is worth one point.

A. Excel

B. TSV

C. JSON

D. LU

E. CSV

These are the two file formats that can be used to import question and answer pairs into an Azure AI Language custom question answering project (QA1). Both formats represent data in a tabular structure, which can be used to import multiple question and answer pairs at once, along with optional metadata.

Excel is a common format for managing data, it is not a supported format for importing directly into a Custom Question Answering project. To import Excel data you'd first need to save it as a TSV or CSV.

Answer: B, E

Explanation:

Question: 334

DRAG DROP

You Have a chatbot that uses the Azure AI Language custom question answering service. The model used by the service was trained by using an internal support FAQ document.

You discover that the chatbot fails to provide correct answers to common questions.

You need to increase the accuracy of the responses provided by the chatbot. The solution must minimize development effort.

Which three actions should you perform in sequence from Language Studio? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Enable Active Learning First: This sets up the process of collecting feedback and identifying areas where the model is struggling.

Review Suggestions: This is the next step where you leverage active learning data, identify errors, and pinpoint what needs to be improved.

Update Q&A Pairs: Based on what you find in review, this is the step where you take corrective action in the data.

Retrain and Republish: Lastly, to make the changes take effect, a retrain and republish is needed

Actions	Answer Area
Type text here	
Update the question and answer pairs.	1. Enable Active learning
Review and accept the alternative phrases.	2. open the review suggestion pane
Open the Edit knowledge base pane.	3. update the question and answers pane
Open the Review suggestions pane.	4. Retrain and republish model
Enable active learning.	
Retrain and republish the model.	
Modify the FAQ document, and then reload it.	

Answer:

Explanation:

Actions	Answer Area
Open the Edit knowledge base pane.	
Update the question and answer pairs.	1. <input type="checkbox"/> Enable active learning.
Retrain and republish the model.	2. <input type="checkbox"/> Open the Review suggestions pane.
Modify the FAQ document, and then reload it.	3. <input checked="" type="checkbox"/> Review and accept the alternative phrases.

Not correct

Review and accept the algorithm phrases: This action relates to a different part of Language Studio - conversational language understanding, not question answering

Question: 335

You have an Azure subscription that contains an Azure AI Document Intelligence resource named Aldoc1.

You have an app named App1 that uses Aldoc1. App1 analyzes business cards by calling business card model v2.1.

You need to update App1 to ensure that the app can interpret QR codes. The solution must minimize administrative effort.

What should you do first?

- A. Deploy a custom model. **Requires training a new model, labeling data, and managing its lifecycle. This increases effort and complexity.**
- B. Implement the read model.
- C. Upgrade the business card model to v3.0
- D. Implement the contract model

Business Card Model v3.0 and Barcode/QR Code Support: The Document Intelligence service evolves, with newer versions of prebuilt models incorporating more features. The business card model, from version 3.0 onwards, supports barcode and QR code interpretation as an add-on feature.

Answer: C

Explanation:

Question: 336

Implementing the Read Model (B): The Read model primarily focuses on text extraction, not specialized data extraction like barcodes and QR codes. While it can extract text from a QR code, it doesn't interpret the encoded information itself.

DRAG DROP

You have an Azure subscription that contains an Azure AI Search resource named AS1.

You implement a custom skill in AS1 that performs language and sentiment analysis of documents.

You are evaluating the use of AS1 as part of an enrichment pipeline.

In which order will AS1 index the documents? To answer, move all indexing stages from the list of stages to the answer area and arrange them in the correct order.

Stages
document cracking
skillset execution
push to index
output field mappings
field mappings

1. document cracking:
Azure AI Search takes raw documents from a data source (like Blob Storage) and extracts their content.

2. Field Mapping:
how specific fields from the data source are mapped to fields in the search index.
e.g 'title' field from your document should be mapped to the 'title' field in the search index schema.

3. Skillset Selection:
This is the step where the enrichment pipeline is applied.

4. Output Field Mapping:

Explanation: This step defines how the outputs of the skillset are mapped to fields in the index.

Answer:

Stages
1. Document Cracking:
2. Field Mapping:
3. Skillset Selection:
4. Output Field Mapping:
5. Push to Index

Answer Area
1. document cracking
2. skillset execution
3. output field mappings
4. field mappings
5. push to index

5. Push to Index

enhanced with the results of the skillset and field mapping, are pushed into the search index.

Question: 337

You have a product support manual.

You need to build a product support chatbot based on the manual. The solution must minimize development effort and costs.

What should you use?

- A. Azure AI Phi-3-medium with fine-tuning
- B. Azure A1 Language Custom question answering
- C. Azure OpenA1 GPI-4 with grounding data that uses Azure AI Search
- D. Azure AI Document Intelligence

B. Purpose: This service is specifically designed to build conversational question answering systems from structured content, like FAQs or manuals.

You can import documents directly, and it automatically extracts question-answer pairs.

Answer: B

Explanation:

Minimized Effort: You can import your product manual, and the service automatically extracts question-answer pairs. You can then refine these pairs and train the model with minimal coding.

A. Azure AI Phi-3-medium with fine-tuning: While Phi-3 is a powerful model, fine-tuning it for a specific task like product support would require more development effort and potentially higher costs compared to using a dedicated question-answering service.

Thank you for your visit.
To try more exams, please visit below link
<https://www.validexamdumps.com/AI-102.html>

```
curl -i POST https://<your-endpoint>.cognitiveservices.azure.com/formrecognizer/documentModels  
/prebuilt-layout:analyze?api-version=2024-11-30 \  
-H 'Content-Type: application/pdf' \  
-H 'Ocp-Apim-Subscription-Key: <your-key>' \  
--data-binary @/path/to/Test.pdf
```

The key in the context of the Azure AI Document Intelligence curl command refers to the Ocp-Apim-Subscription-Key, which is your subscription key for accessing the service.

Explanation:

Subscription Key: This key is unique to your Azure subscription and is used to authenticate your requests to the Document Intelligence service.

Ocp-Apim-Subscription-Key: This is the name of the header in the HTTP request where you need to provide the subscription key value.

Incorrect options:

key1, secret: These are placeholders and are not the standard terms for the Document Intelligence API key.

subscription-key: While this is a commonly used term for a subscription key, it's not the specific name of the header required in the curl command for Document Intelligence.